

CS 5600

Computer Systems

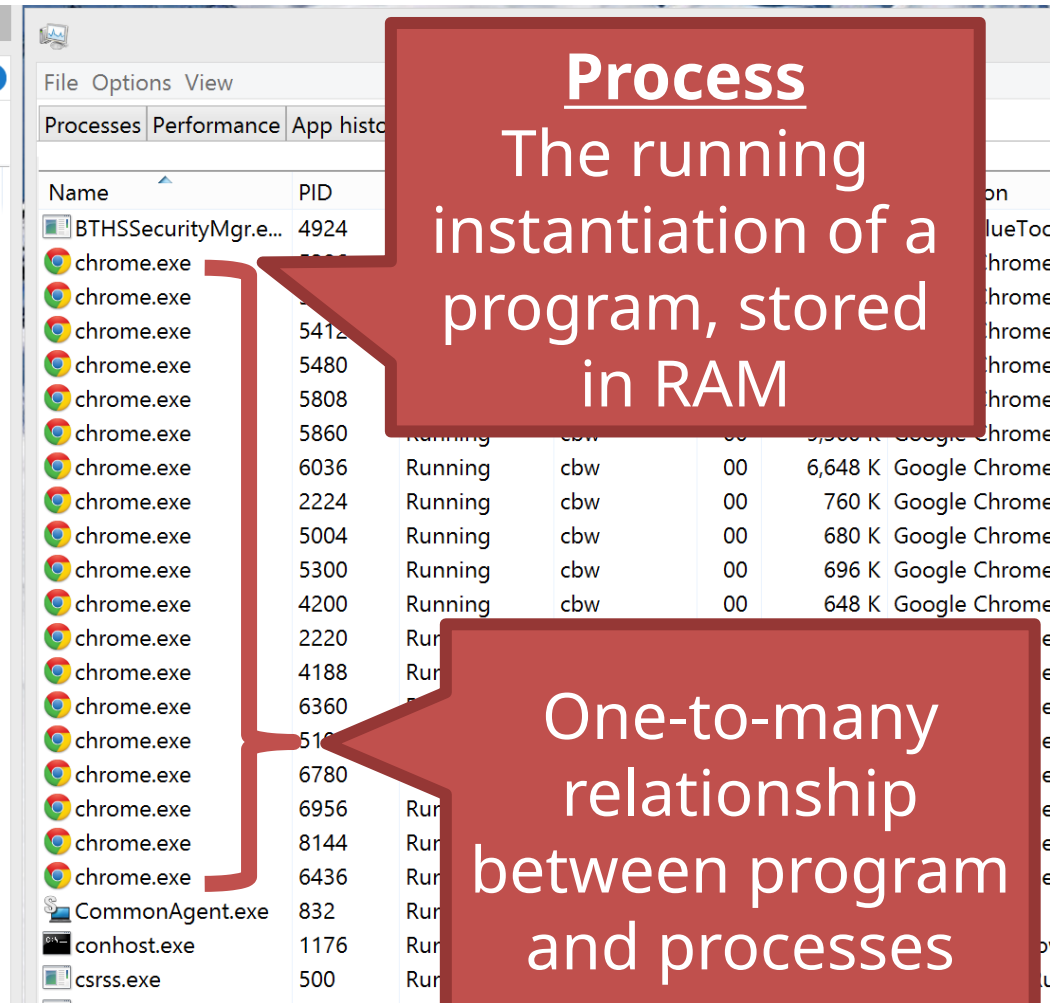
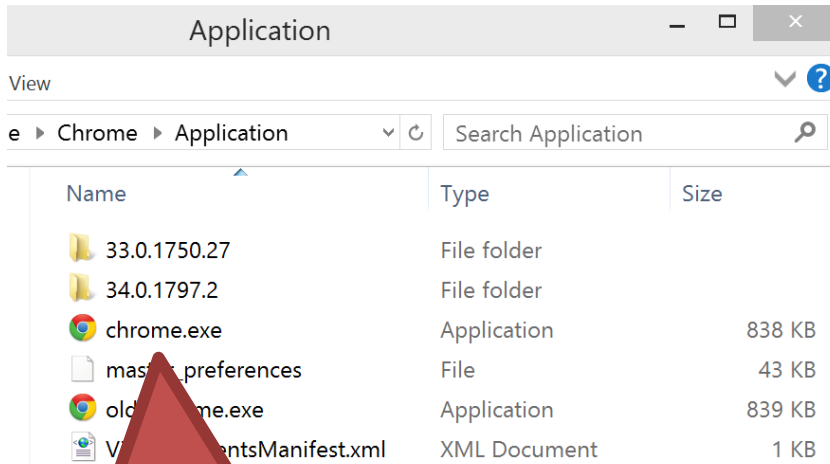
Lecture 4: Programs, Processes, and Threads

- Programs
- Processes
- Context Switching
- Protected Mode Execution
- Inter-process
Communication
- Threads

Running Dynamic Code

- One basic function of an OS is to execute and manage code dynamically, e.g.:
 - A command issued at a command line terminal
 - An icon double clicked from the desktop
 - Jobs/tasks run as part of a batch system (MapReduce)
- A **process** is the basic unit of a program in execution

Programs and Processes



Process

The running instantiation of a program, stored in RAM

Program

An executable file in long-term storage

One-to-many relationship between program and processes

How to Run a Program?

- When you double-click on an .exe, how does the OS turn the file on disk into a process?
- What information must the .exe file contain in order to run as a program?

Program Formats

- Programs obey specific file formats
 - CP/M and DOS: COM executables (*.com)
 - DOS: MZ executables (*.exe)
 - Named after Mark Zbikowski, a DOS developer
 - Windows Portable Executable (PE, PE32+) (*.exe)
 - Modified version of Unix COFF executable format
 - PE files start with an MZ header. Why?
 - Unix/Linux: Executable and Linkable Format (ELF)

test.c

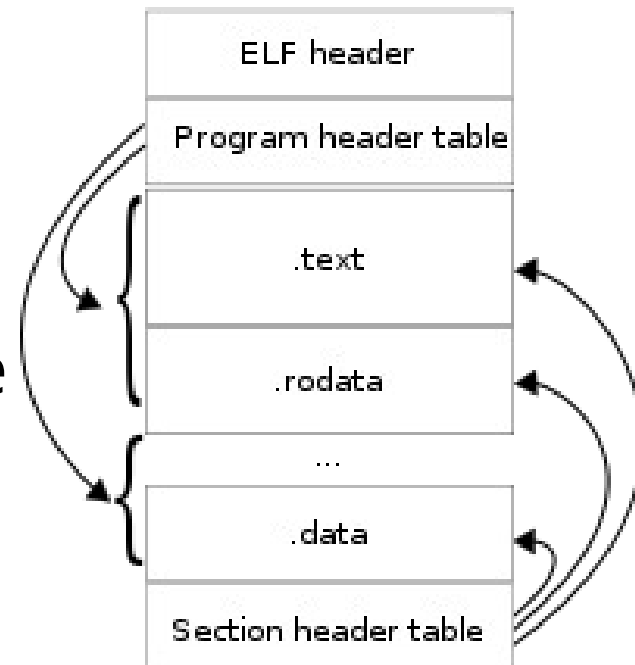
```
#include <stdio.h>
```

```
int big_big_array[10 * 1024 * 1024];  
char *a_string = "Hello, World!";  
int a_var_with_value = 100;
```

```
int main(void) {  
    big_big_array[0] = 100;  
    printf("%s\n", a_string);  
    a_var_with_value += 20;  
  
    printf("main is : %p\n", &main);  
    return 0;  
}
```

ELF File Format

- ELF Header
 - Contains compatibility info
 - **Entry point** of the executable code
- Program header table
 - Lists all the segments in the file
 - Used to load and execute the program
- Section header table
 - Used by the linker



ELF Header Format

```
typedef struct {
```

```
1   unsigned char e_ident[EI_NIDENT];
```

```
   Elf32_Half e_type;
```

```
5   Elf32_Half e_machine,
```

```
   Elf32_Word e_version;
```

```
   Elf32_Addr e_entry,
```

```
   Elf32_Off e_phoff,
```

```
   Elf32_Off e_shoff;
```

```
10  Elf32_Word e_flags;
```

```
   Elf32_Half e_ehsize;
```

```
   Elf32_Half e_phentsize;
```

```
   Elf32_Half e_phnum,
```

```
   Elf32_Half e_shentsize;
```

```
15  Elf32_Half e_shnum,
```

```
   Elf32_Half e_shstrndx;
```

```
} Elf32_Ehdr;
```

ISA of executable

- Entry point of executable code
- What should EIP be set to initially?

of program headers

of section headers

ELF Header Example

```
$ gcc -g -o test test.c
```

```
$ readelf --header test
```

ELF Header:

```
Magic:      7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
Class:      ELF64
Data:      2's complement, little endian
Version:    1 (current)
OS/ABI:     UNIX - System V
ABI Version: 0
Type:      EXEC (Executable file)
Machine:    Advanced Micro Devices X86-64
Version:    0x1
Entry point address: 0x400460
Start of program headers: 64 (bytes into file)
Start of section headers: 5216 (bytes into file)
Flags:      0x0
Size of this header: 64 (bytes)
Size of program headers: 56 (bytes)
Number of program headers: 9
Size of section headers: 64 (bytes)
Number of section headers: 36
Section header string table index: 33
```

Investigating the Entry Point

```
int main(void) {  
    ...  
    printf("main is : %p\n", &main);  
    return 0;  
}
```

```
$ gcc -g -o test test.c  
$ readelf --headers ./test | grep Entry  
  Entry point address:          0x400460  
$ ./test  
Hello World!  
main is : 0x400544
```

Entry point != &main

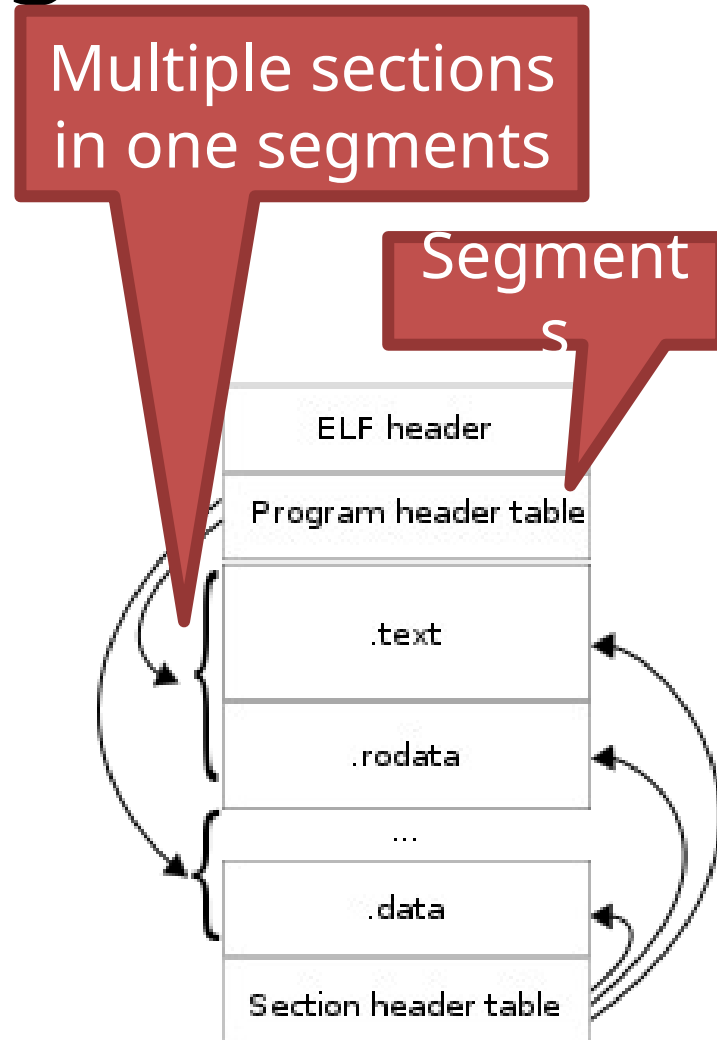
```
$ ./test
Hello World!
main is : 0x400544
$ readelf --headers ./test | grep Entry
Entry point address: 0x400460
$ objdump --disassemble -M intel ./test
```

```
...
0000000000400460 < start>:
400460: 31 ed                xor    ebp,ebp
400462: 49 89 d1             mov    r9,rdx
400465: 5e                  pop    rsi
400466: 48 89 e2             mov    rdx,rsp
400469: 48 83 e4 f0          and    rsp,0xfffffffffffffff0
40046d: 50                  push   rax
40046e: 54                  push   rsp
40046f: 49 c7 c0 20 06 40 00 mov    r8,0x400620
400476: 48 c7 c1 90 05 40 00 mov    rcx,0x400590
40047d: 48 c7 c7 44 05 40 00 mov    rdi,0x400544
400484: e8 c7 ff ff ff      call  400450 <__libc_start_main@plt>
...
```

- Most compilers insert extra code into compiled programs
- This code typically runs before and after main()

Sections and Segments

- **Sections** are the various pieces of code and data that get linked together by the compiler
- Each **segment** contains one or more sections
 - Each segment contains sections that are related
 - E.g. all code sections
 - Segments are the basic units for the **loader**



Common Sections

- Sections are the various pieces of code and data that compose a program
- Key sections:
 - `.text` – Executable code
 - `.bss` – Global variables initialized to zero
 - `.data`, `.rodata` – Initialized data and strings
 - `.strtab` – Names of functions and variables
 - `.symtab` – Debug symbols

String variable → .data

Example

Empty 10 MB
array → .bss

```
int big_big_array[10*1024*1024];  
char *a_string = "Hello, World!";  
int a_var_with_value = 0x100;
```

Initialized
global variable
→ .data

```
int main(void) {  
    big_big_array[0] = 100;  
    printf("%s\n", a_string);  
    a_var_with_value += 20;  
    ...  
}
```

String constant
→ .rodata

Code → .text

```
$ readelf --headers ./test
```

...
Section to Segment mapping:
Segment Sections...

00	
01	.interp
02	.interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt .init .plt .text .fini .rodata .eh_frame_hdr .eh_frame
03	.ctors .dtors .jcr .dynamic .got .got.plt .data .bss
04	.dynamic
05	.note.ABI-tag .note.gnu.build-id
06	.eh_frame_hdr
07	
08	.ctors .dtors .jcr .dynamic .got

...
There are 36 section headers, starting at offset 0x1460:

Section Headers:

[Nr]	Name	Type	Address	Offset	Size	ES	Flags	Link	Info	Align
[0]		NULL	00000000		00000000		00000000	00		0
0 0										
[1]	.interp	PROGBITS	00400238		00000238		0000001c	00	A	0
0 1										
[2]	.note.ABI-tag	NOTE	00400254		00000254		00000020	00	A	0
0 4										
[3]	.note.gnu.build-I	NOTE	00400274		00000274		00000024	00		
A 0		0	4							
[4]	.gnu.hash	GNU_HASH	00400298		00000298		0000001c	00	A	5

.text Example Header

```
typedef struct {
```

```
    Elf32_Word p_type;
```

```
    Elf32_Off  p_offset;
```

```
5    Elf32_Addr p_vaddr;
```

```
    Elf32_Addr p_paddr;
```

```
    Elf32_Word p_filesz;
```

```
    Elf32_Word p_memsz;
```

```
    Elf32_Word p_flags;
```

```
10   Elf32_Word p_align;
```

Address to load
section in
memory

Offset of data in the
file

How many bytes (in
hex) are in the
section

Executable

Data for
the
program

```
Sections ./test
```

```
...
```

```
Section Headers:
```

```
...
```

[Nr]	Name	Type	Address	Offset	Size	ES	Flags	Link
[13]	.text	PROGBITS	00400460	00000460	00000218	00	AX	

```
0 0 16
```

.bss Example Header

```
int big_big_array[10*1024*1024];
```

```
typedef struct {
```

```
Elf32_Word p_type;  
Elf32_Off p_offset;  
Elf32_Addr p_vaddr;  
Elf32_Addr p_paddr;  
Elf32_Word p_filesz;  
Elf32_Word p_memsz;  
Elf32_Word p_flags;  
Elf32_Word p_align;
```

Offset of data in the
file
(Notice the length =
0)

5

Address to load
section in
memory

Contains
no
data

10

$\text{hex}(4 \cdot 10 \cdot 1024 \cdot 1024)$
= 0x2800020

Writabl
e

```
$ readelf -sections test
```

```
...  
Section Headers:
```

```
...  
[Nr] Name      Type      Address      Offset      Size      ES      Flags      Link  
Info Align  
[25] .bss      NOBITS    00601040     00001034    02800020    00      WA      0      0  
32
```

```
[26] comment
```

```
PROGBITS
```

```
00000000
```

```
00001034
```

Segments

- Each segment contains one or more sections
 - All of the sections in a segment are related, e.g.:
 - All sections contain compiled code
 - Or, all sections contain initialized data
 - Or, all sections contain debug information
 - ... etc...
- Segments are used by the loader to:
 - Place data and code in memory
 - Determine memory permissions (read/write/execute)

Segment Header

typedef struct {

Elf32_Word p_type;

Elf32_Off p_offset;

Elf32_Addr p_vaddr;

Elf32_Addr p_paddr;

Elf32_Word p_filesz;

Elf32_Word p_memsz;

Elf32_Word p_flags;

Elf32_Word p_align;

}

Type of segment

Offset within the ELF
Location to load the
segment into
memory

Size of the segment in
memory

- Flags describing the section data
- Examples: executable, read-only

\$ readelf --segments ./test

Elf file type is EXEC (Executable file)

Entry point 0x400460

There are 9 program headers, starting at offset 64

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flags	Align
PHDR	0x00000040	0x00400040		0x00400040		0x000001f8	
0x000001f8	R E	8					
INTERP	0x00000238	0x00400238		0x00400238		0x0000001c	
0x0000001c	R	1					
LOAD	0x00000000	0x00400000		0x00400000		0x0000077c	
0x0000077c	R E	200000					
LOAD	0x00000e28	0x00600e28		0x00600e28		0x0000020c	
0x02800238	RW	200000					
DYNAMIC	0x00000e50	0x00600e50		0x00600e50		0x00000190	
0x00000190	RW	8					
NOTE	0x00000254	0x00400254		0x00400254		0x00000044	
0x00000044	R	4					
GNU_EH_FRAME	0x000006a8	0x004006a8		0x004006a8		0x0000002c	
0x0000002c	R	4					
GNU_STACK	0x00000000	0x00000000		0x00000000		0x00000000	
0x00000000	RW	8					
GNU_RELRO	0x00000e28	0x00600e28		0x00600e28		0x000001d8	
0x000001d8	R	1					



Section to Segment mapping:

Segment Sections...

What About Static Data?

```
#include <stdio.h>
```

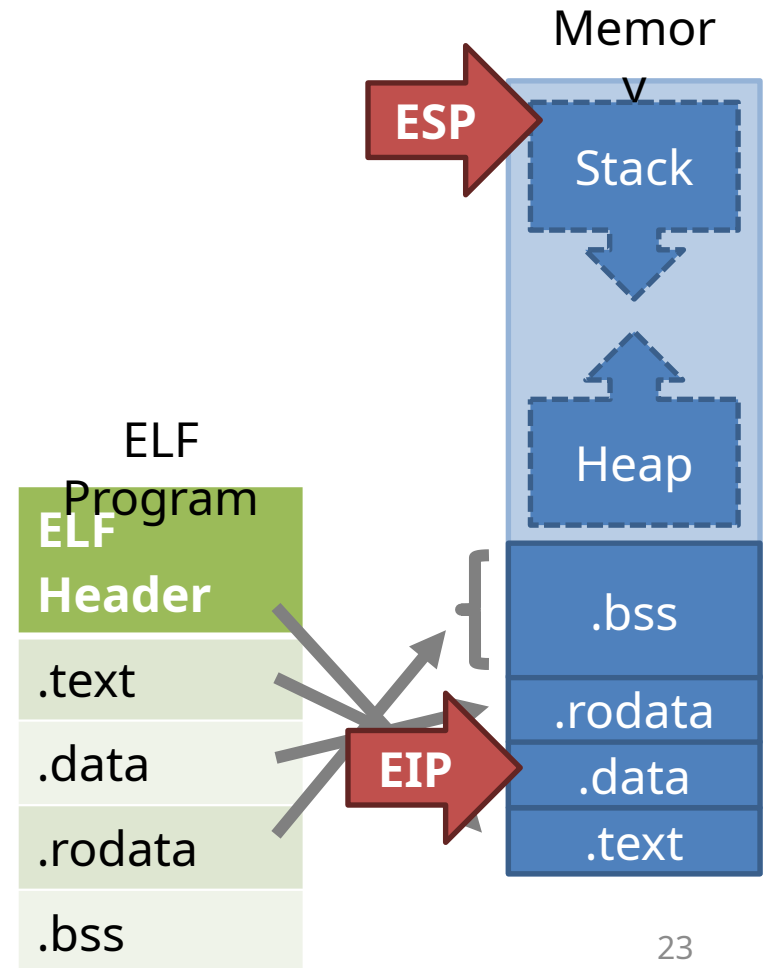
```
int big_big_array[10 * 1024 *  
1024];  
char *a_string = "Hello, World!";  
int a_var_with_value = 100;
```

```
int main(void) {  
    big_big_array[0] = 100;  
    printf("%s\n", a_string);  
    a_var_with_value += 20;  
  
    printf("main is : %p\n", &main);  
    return 0;  
}
```

```
$ strings -t d ./test  
568  /lib64/ld-linux-x86-  
64.so.2  
817  __gmon_start_  
832  libc.so.6  
842  puts  
847  printf  
854  __libc_start_main  
872  GLIBC_2.2.5  
1300 fff.  
1314 =  
1559 |$ L  
1564 t$(L  
1569 |$0H  
1676 Hello, World!  
1690 main is : %p  
1807 ;*3$"
```

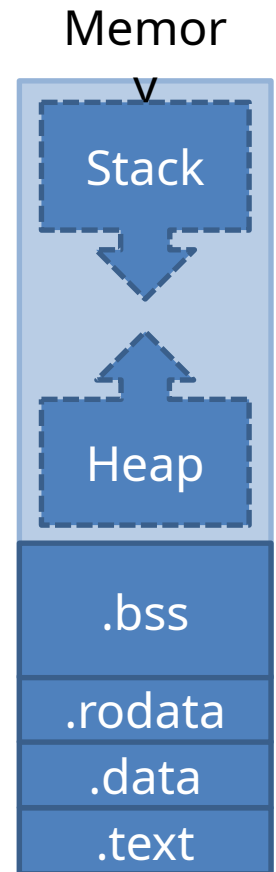
The Program Loader

- OS functionality that loads programs into memory, creates processes
 - Places segments into memory
 - Expands segments like .bss
 - Loads necessary dynamic libraries
 - Performs relocation
 - Allocates the initial stack frame
 - Sets EIP to the programs entry point



Single-Process Address Space

- The stack is used for local variables and function calls
 - Grows downwards
- Heap is allocated dynamically (malloc/new)
 - Grows upwards
- When the stack and heap meet, there is no more memory left in the process
 - Process will probably crash
- Static data and global variables are fixed at compile time



Problem: Pointers in Programs

- Consider the following code:

```
int foo(int a, int b) { return a * b - a / b; }  
int main(void) { return foo(10, 12); }
```

- Compiled, it might look like this:

000FE4D8 <foo>:

000FE4D8: mov eax, [esp+4]

000FE4DB: mov ebx, [esp+8]

000FE4DF: mul eax, ebx

...

000FE21A: push eax

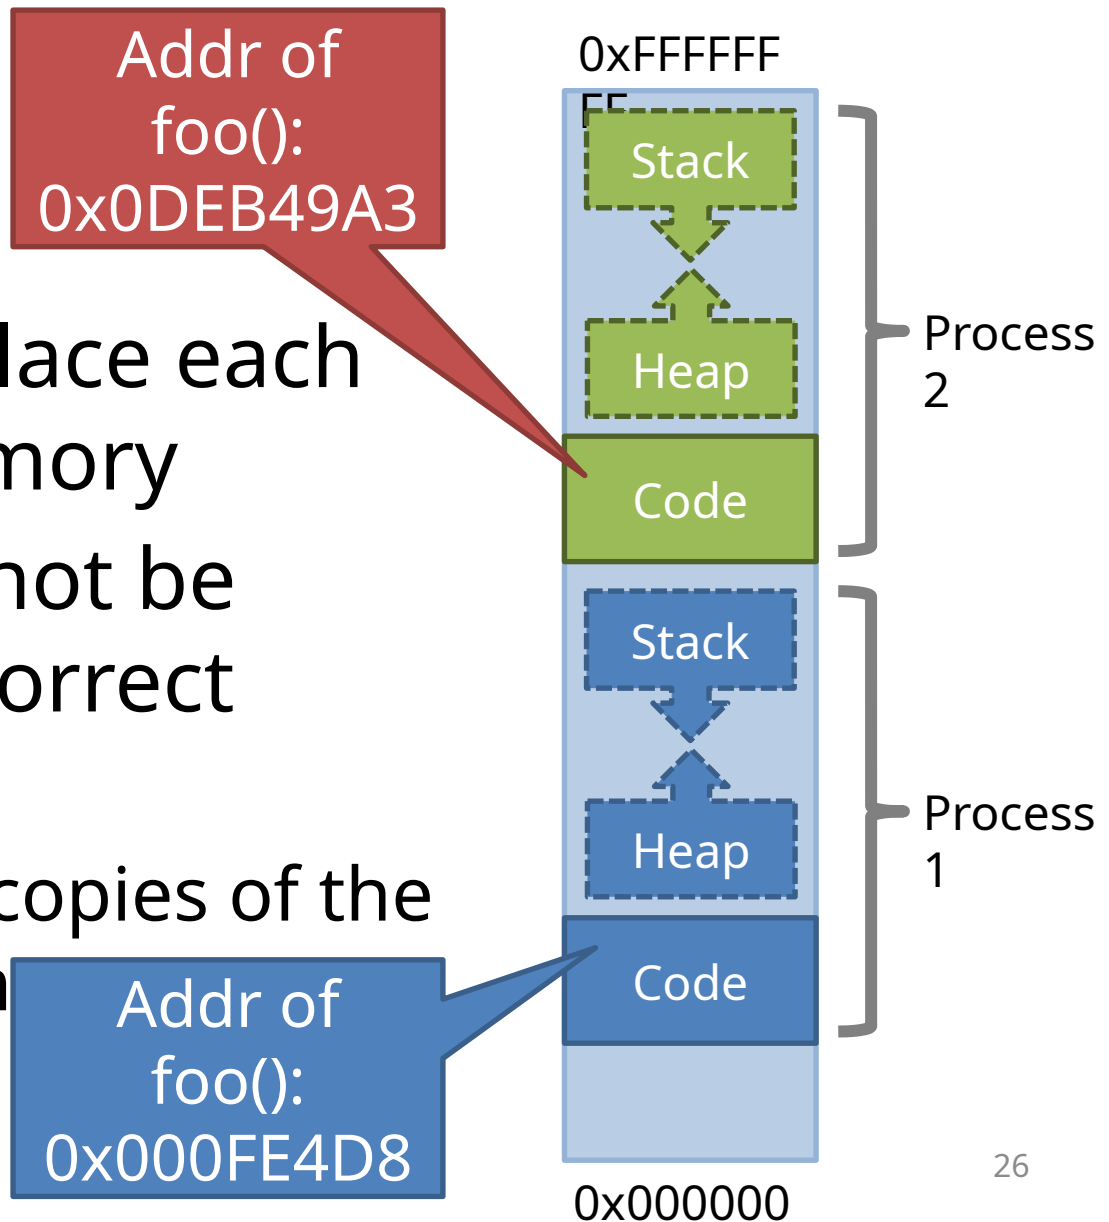
000FE21D: push ebx

000FE21F: call 0x000FE4D8

- ... but this assembly assumes foo() is at address 0x000FE4D8

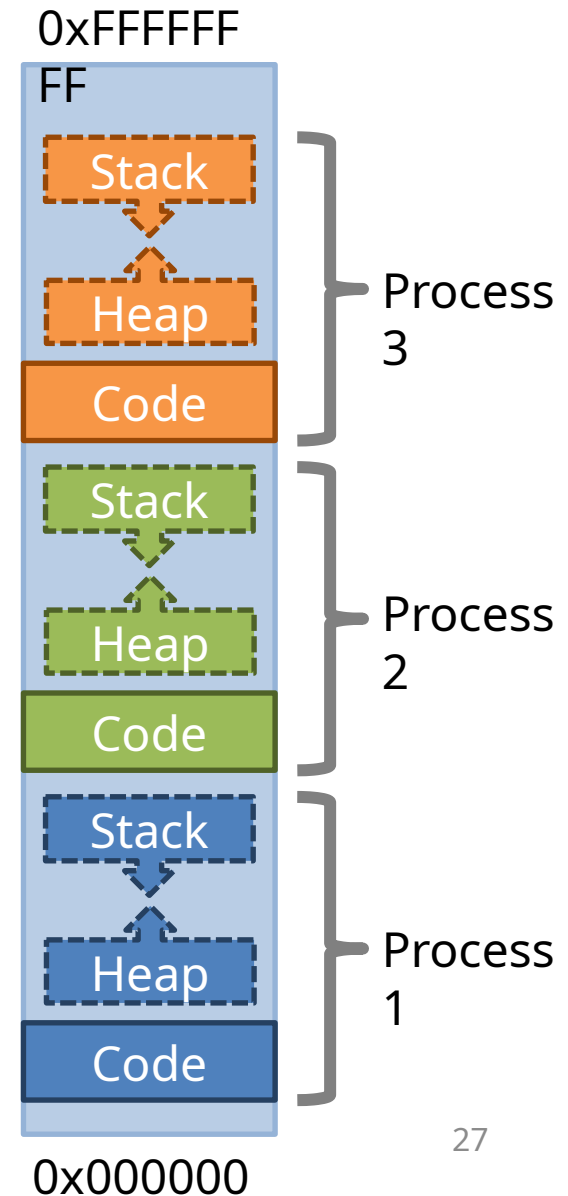
Program Load Addresses

- Loader must place each process in memory
- Program may not be placed at the correct location!
 - Example: two copies of the same program



Address Spaces for Multiple Processes

- Many features of processes depend on pointers
 - Addresses of functions
 - Addresses of strings, data
 - Etc.
- For multiple processes to run together, they all have to fit into memory together
- However, a process may not always be loaded into the same memory location



Address Spaces for Multiple Processes

- There are several methods for configuring address spaces for multiple processes
 1. Fixed address compilation
 2. Load-time fixup
 3. Position independent code
 4. Hardware support

Fixed-Address Compilation

Single Copy of Each Program

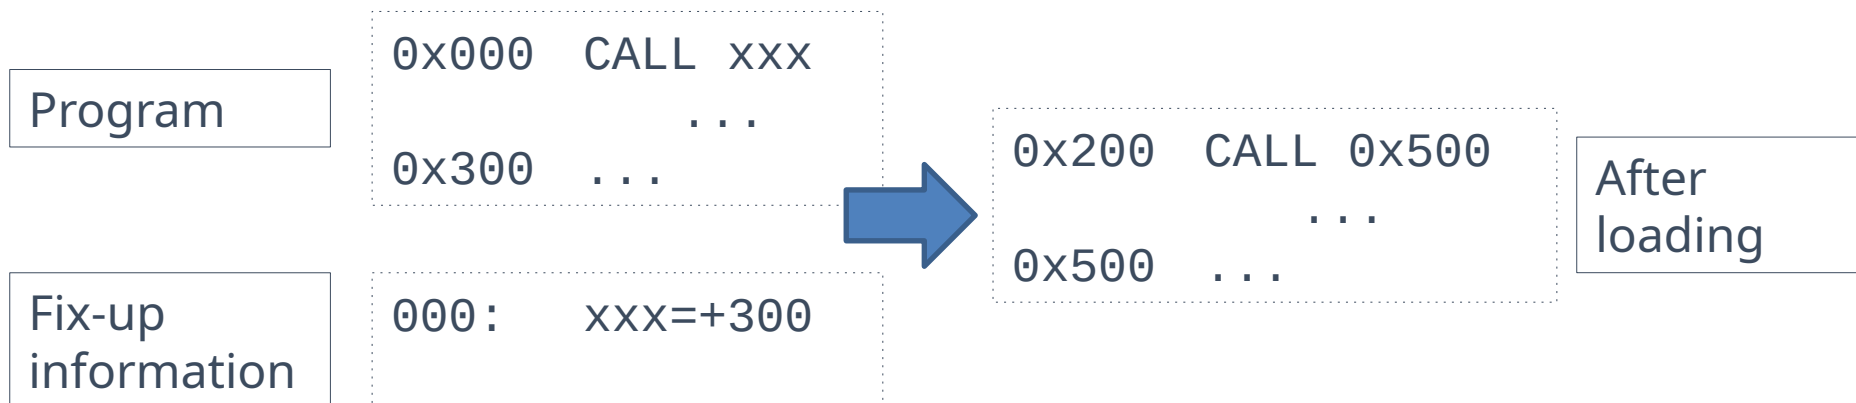
- Compile each program once, with fixed addresses
- OS may only load program at the specified offset in memory
- Typically, only one process may be run at any time
- Example: MS-DOS 1.0

Multiple Copies of Each Program

- Compile each program multiple times
- Once for each possible starting address
- Load the appropriate compiled program when the user starts the program
- Bad idea
 - Multiple copies of the same program

Load-Time Fixup

- Calculate addresses at load-time instead of compile-time
- The program contains a list of locations that must be modified at startup
 - All relative to some starting address
- Used in some OSes that run on low-end microcontrollers without virtual memory hardware



Position-Independent Code

- Compiles programs in a way that is independent of their starting address
 - PC-relative address
- Slightly less efficient than absolute addresses
- Commonly used today for security reasons

Absolute
addressing

```
0x200  CALL 0x500
...
0x500  ...
```

PC-relative
addressing

```
0x200  CALL PC+0x300
...
0x500  ...
```

Hardware Support

- Hardware address translation
- Most popular way of sharing memory between multiple processes
 - Linux
 - OS X
 - Windows
- Program is compiled to run at a fixed location in **virtual memory**
- The OS uses the **MMU** to map these locations to physical memory

MMU and Virtual Memory

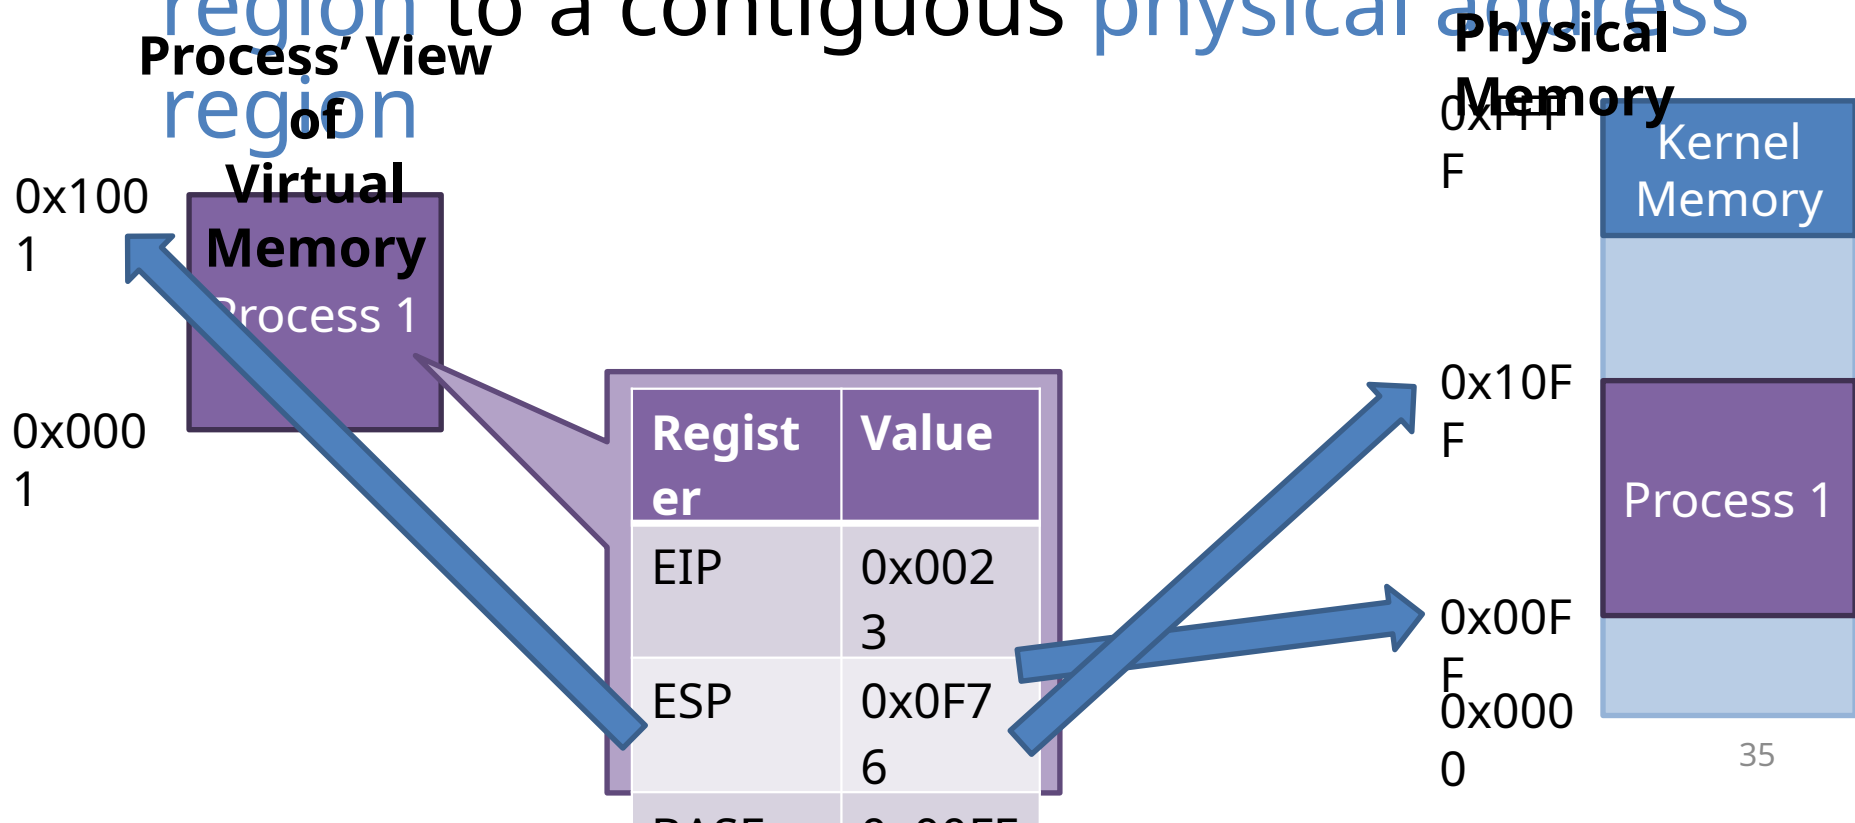
- The Memory Management Unit (MMU) translates between virtual addresses and physical addresses
 - Process uses **virtual address** for calls and data load/store
 - MMU translates virtual addresses to **physical addresses**
 - The physical addresses are the true locations of code and data in RAM

Advantages of Virtual Memory

- Flexible memory sharing
 - Simplifies the OS's job of allocating memory to different programs
- Simplifies program writing and compilations
 - Each program gets access to 4GB of RAM (on a 32-bit CPU)
- Security
 - Can be used to prevent one process from accessing the address of another process
- Robustness
 - Can be used to prevent writing to addresses belonging to the OS (which may cause the OS to crash)

Base and Bounds Registers

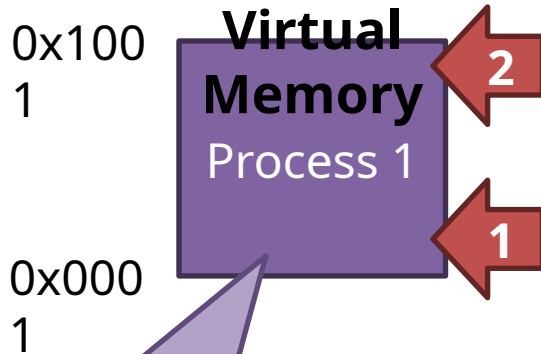
- A simple mechanism for address translation
- Maps a contiguous **virtual address region** to a contiguous **physical address region**



Base and Bounds Example

Process' View of

Virtual Memory
Process 1



`0x0023 mov eax, [esp]`

1) Fetch instruction

$$0x0023 + 0x00FF = 0x0122$$

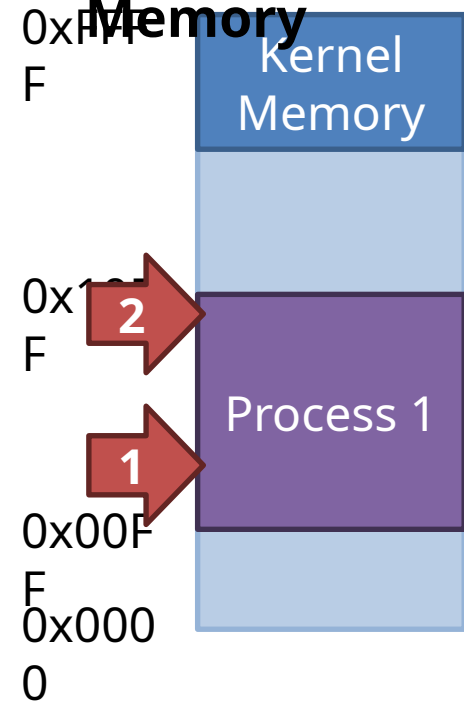
2) Translate memory access

$$0x0F76 + 0x00FF = 0x1075$$

3) Move value to register
[0x1075] → eax

Register	Value
EIP	0x0023
ESP	0x0F76
BASE	0x00FF
BOUND	0x1000

Physical Memory



Confused About Virtual Memory?

- That's okay :)
- We will discuss virtual memory at great length later in the semester
- In project 3, you will implement virtual memory in Pintos

- Programs
- Processes
- Context Switching
- Protected Mode Execution
- Inter-process Communication
- Threads

From the Loader to the Kernel

- Once a program is loaded, the kernel must manage this new process
- Program Control Block (PCB): kernel data structure representing a process
 - Has at least one thread (possibly more...)
 - Keeps track of the memory used by the process
 - Code segments
 - Data segments (stack and heap)
 - Keeps runtime state of the process
 - CPU register values
 - EIP

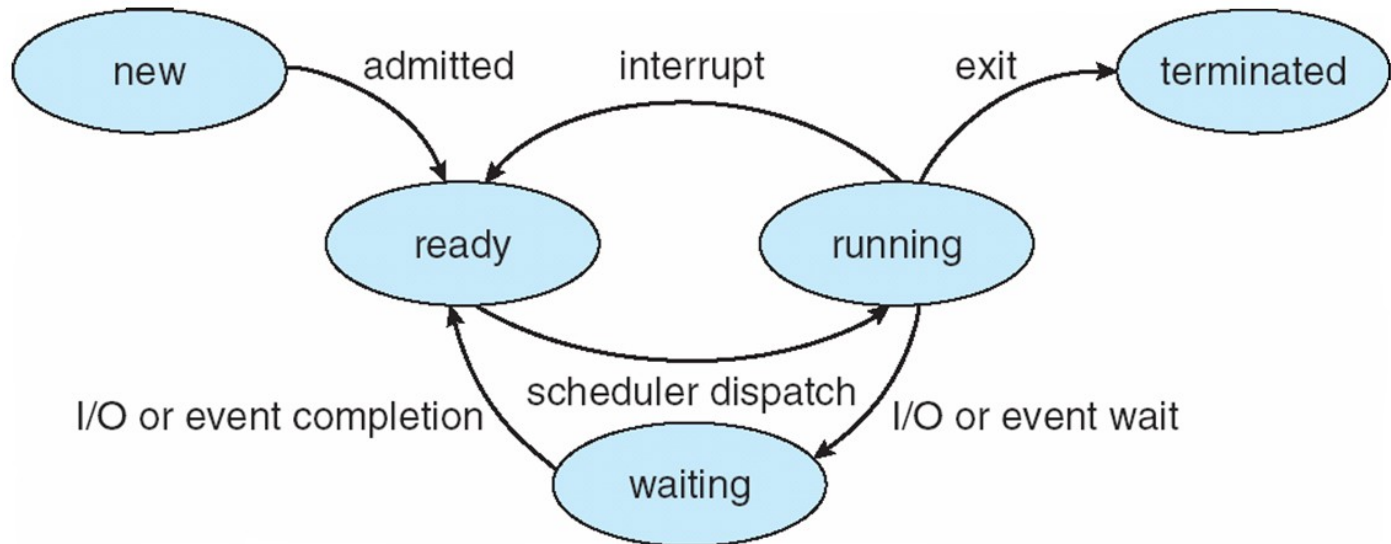
Program Control Block (PCB)

- OS structure that represents a process in memory
- Created for each process by the loader
- Managed by the kernel

```
struct task_struct {  
    // Typical Unix PCB  
    pid_t pid;  
    // process identifier  
    long state;  
    // state of the process  
    unsigned int time_slice;  
    // scheduling information  
    struct task_struct *parent;           // this  
    // process's parent  
    struct list_head children;           // this  
    // process's children  
    struct files_struct *files;           // list of  
    // open files  
    struct mm_struct *mm; // address space of this process  
};
```


Process States

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution



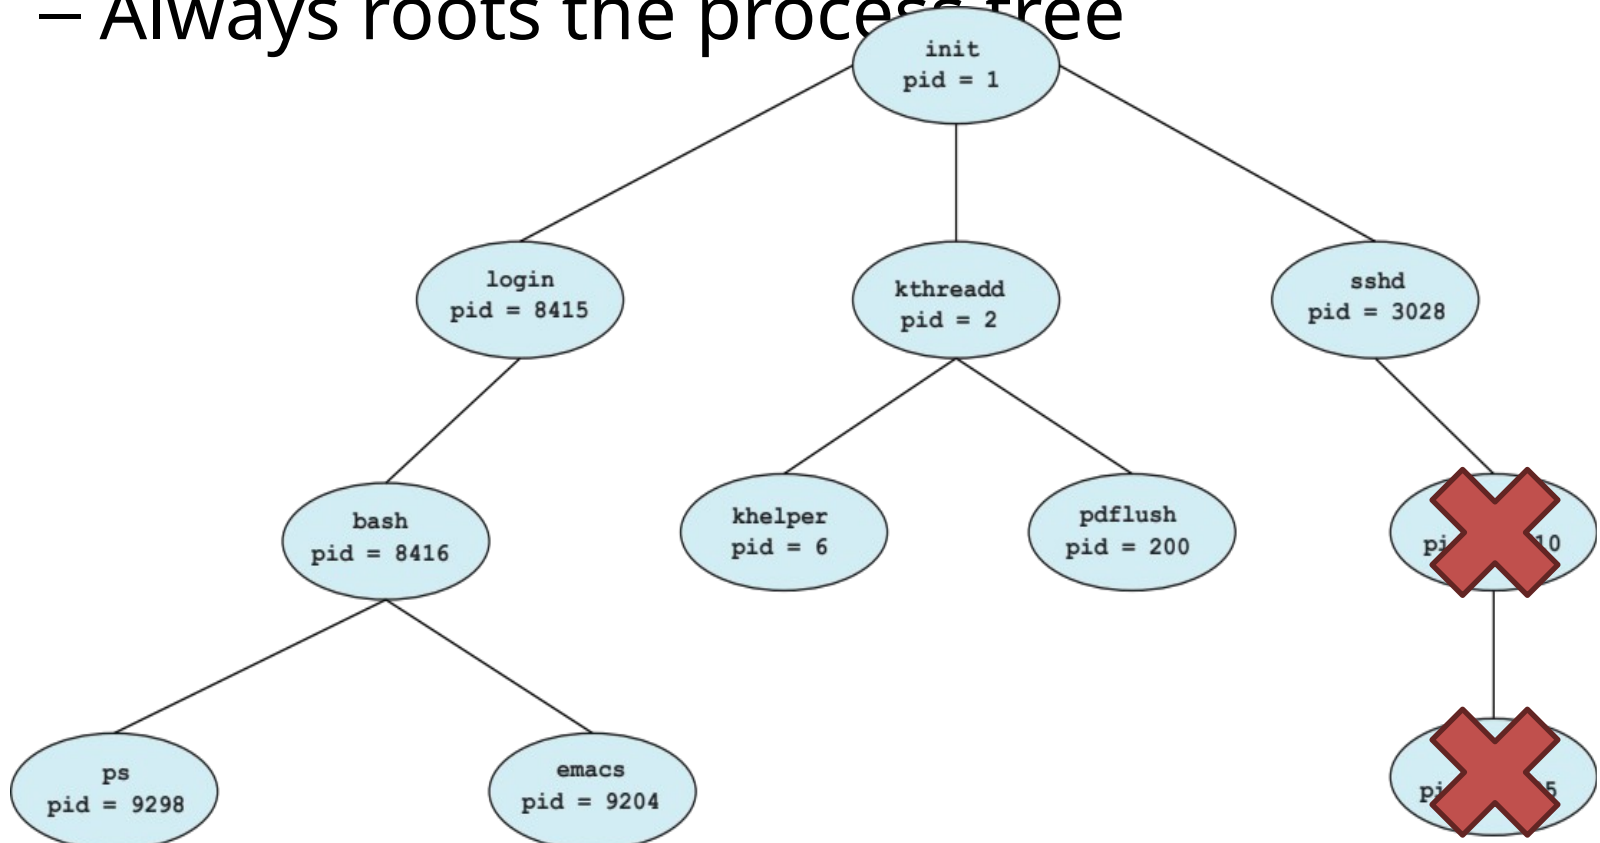
Parents and Children

- On Unix/Linux, all processes have **parents**
 - i.e. which process executed this new process?
- If a process spawns other processes, they become it's **children**
 - This creates a tree of processes
- If a parent exits before its children, the children become **orphans**
- If a child exits before the parent calls `wait()`, the child becomes a **zombie**



Process Tree

- **init** is a special process started by the kernel
 - Always roots the process tree



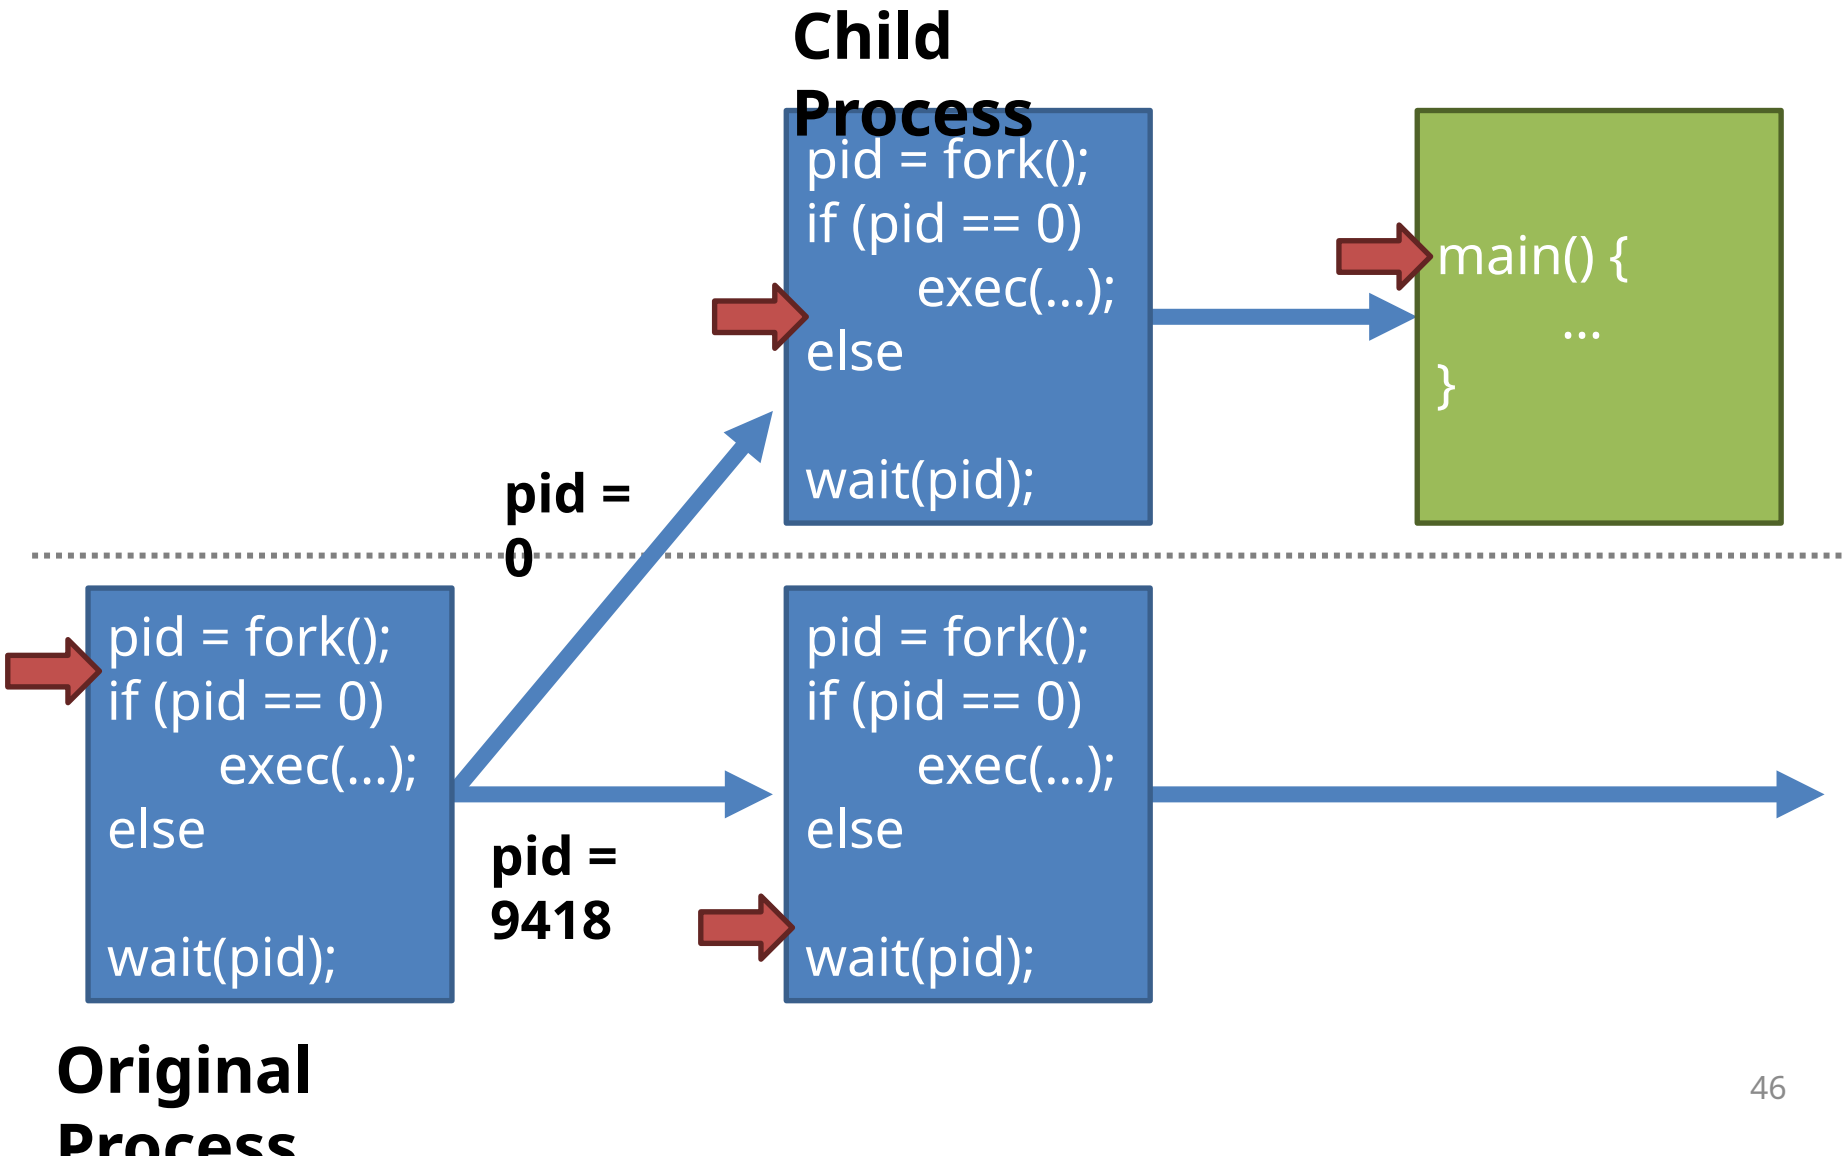
Additional Execution Context

- File descriptors
 - stdin, stdout, stderr
 - Files on disk
 - Sockets
 - Pipes
- Permissions
 - User and group
 - Access to specific APIs
 - Memory protection
- Environment
 - \$PATH
- Shared Resources
 - Locks
 - Mutexes
 - Shared Memory

UNIX Process Management

- `fork()` – system call to create a copy of the current process, and start it running
 - No arguments!
- `exec()` – system call to change the program being run by the current process
- `wait()` – system call to wait for a process to finish
- `signal()` – system call to send a notification to another process

UNIX Process Management



Question: What does this code print?

```
int child_pid = fork();
if (child_pid == 0) {           // I'm the child
    process
    printf("I am process # %d\n", getpid());
    return 0;
} else {                       // I'm the parent process
    printf("I am parent of process # %d\n",
        child_pid);
    return 0;
}
```

Questions

- Can UNIX `fork()` return an error? Why?
- Can UNIX `exec()` return an error? Why?
- Can UNIX `wait()` ever return immediately? Why?

Implementing UNIX fork()

- Steps to implement UNIX fork()
 1. Create and initialize the process control block (PCB) in the kernel
 2. Create a new address space
 3. Initialize the address space with a copy of the entire contents of the address space of the parent
 4. Inherit the execution context of the parent (e.g., any open files)
 5. Inform the scheduler that the new process is ready to run

Implementing UNIX exec()

- Steps to implement UNIX exec()
 1. Load the new program into the current address space
 2. Copy command line arguments into memory in the new address space
 3. Initialize the hardware context to start execution
 - EIP = Entry point in the ELF header
 - ESP = A newly allocated stack

Process Termination

- Typically, a process will `wait(pid)` until its child process(es) complete
- `abort(pid)` can be used to immediately end a child process

- Programs
- Processes
- Context Switching
- Protected Mode Execution
- Inter-process Communication
- Threads

The Story So Far...

- At this point, we have gone over how the OS:
 - Turns programs into processes
 - Represents and manages running process
- Next step: context switching
 - How does a process access OS APIs?
 - i.e. System calls
 - How does the OS share the CPU between several programs?
 - Multiprocessing

Context Switching

- Context switching
 - Saves state of a process before a switching to another process
 - Restores original process state when switching back
- Simple concept, but:
 - How do you save the state of a process?
 - How do you stop execution of a process?
 - How do you restart the execution of process that has been switched out?

The Process Stack

- Each process has a stack in memory that stores:
 - Local variables
 - Arguments to functions
 - Return addresses from functions
- On x86:
 - The stack grows downwards
 - ESP (Stack Pointer register) points to the bottom of the stack (i.e. the newest data)
 - EBP (Base Pointer) points to the base of the current frame
 - Instructions like `push`, `pop`, `call`, `ret`, `int`, and `iret` all modify the stack

```
int bar(int a, int b) {  
    int r = rand();  
    return a + b - r;  
}
```

stack_exam.c

```
int foo(int a) {  
    int x, y;  
    x = a * 2;  
    y = a - 7;  
    return bar(x, y);  
}
```

```
int main(void) {  
    ...  
    foo(12);  
    ...  
}
```



```
$ gcc -g -fno-stack-protector -m32 -o
stack_exam stack_exam.c
$ objdump --disassemble -M intel ./stack_exam
```

```
...
804842a: e8 c0 ff ff ff      call 80483ef
<foo>
```

```
804842f: b8 00 00 00 00      mov    eax,0x0
...
```

```
080483ef <foo>:
```

```
80483ef: 55                  push   ebp
```

```
80483f0: 89 e5              mov    ebp, esp
```

```
80483f2: 83 ec 28          sub    esp, 0x28
```

```
80483f5: 8b 45 08          mov    eax, [ebp+0x8]
```

```
80483f8: 01 c0             add    eax, eax
```

```
80483fa: 89 45 f4          mov    [ebp-0xc], eax
```

```
80483fd: 8b 45 08          mov    eax, [ebp+0x8]
```

```
8048400: 83 e8 07          sub    eax, 0x7
```

```
8048403: 89 45 f0          mov    [ebp-0x10], eax
```

```
8048406: 8b 45 f0          mov    eax, [ebp-0x10]
```

```
8048409: 89 44 24 04       mov    [esp+0x4], eax
```

```
804840d: 8b 45 f4          mov    eax, [ebp-0xc]
```

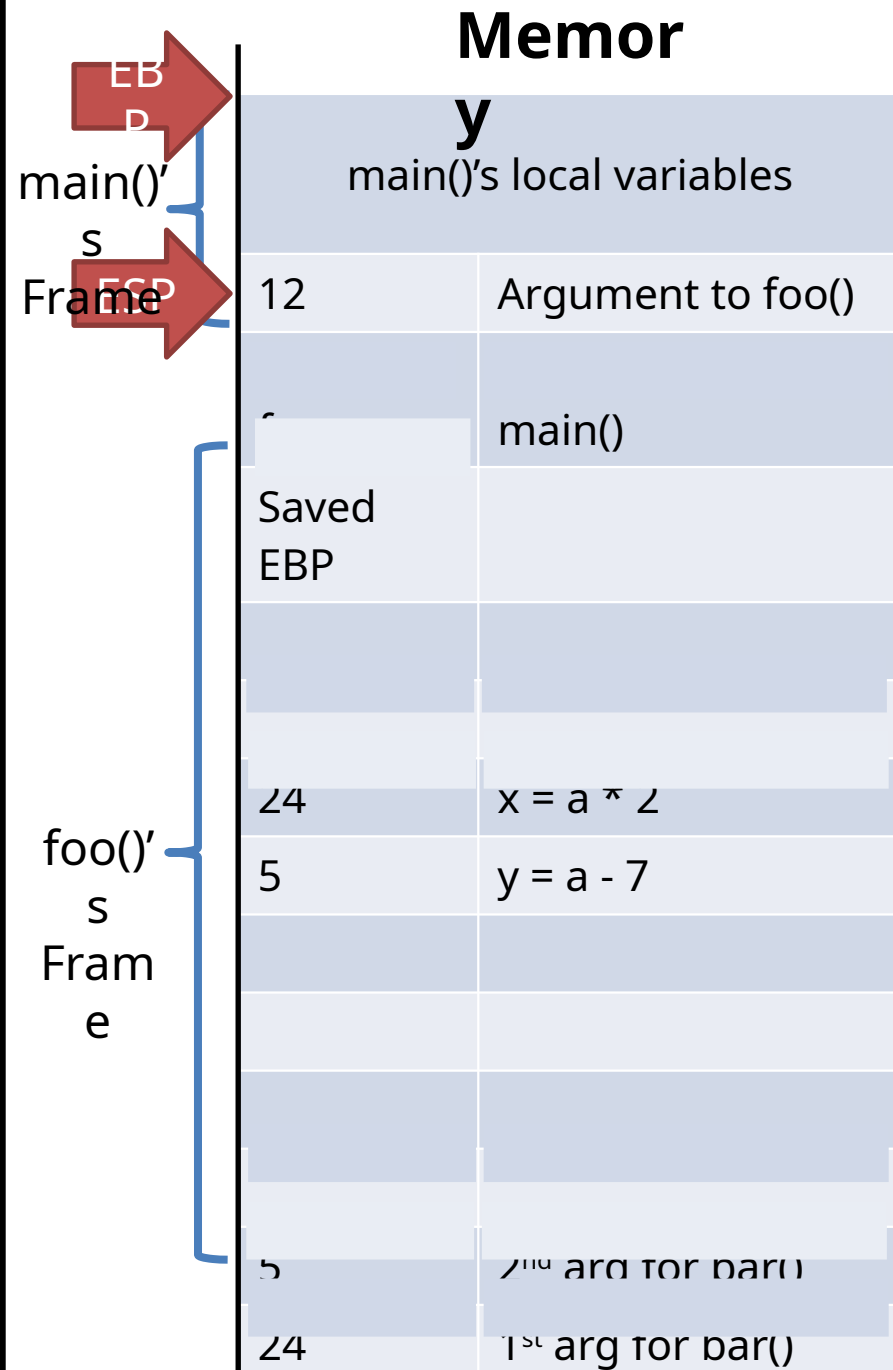
```
8048410: 89 04 24          mov    [esp], eax
```

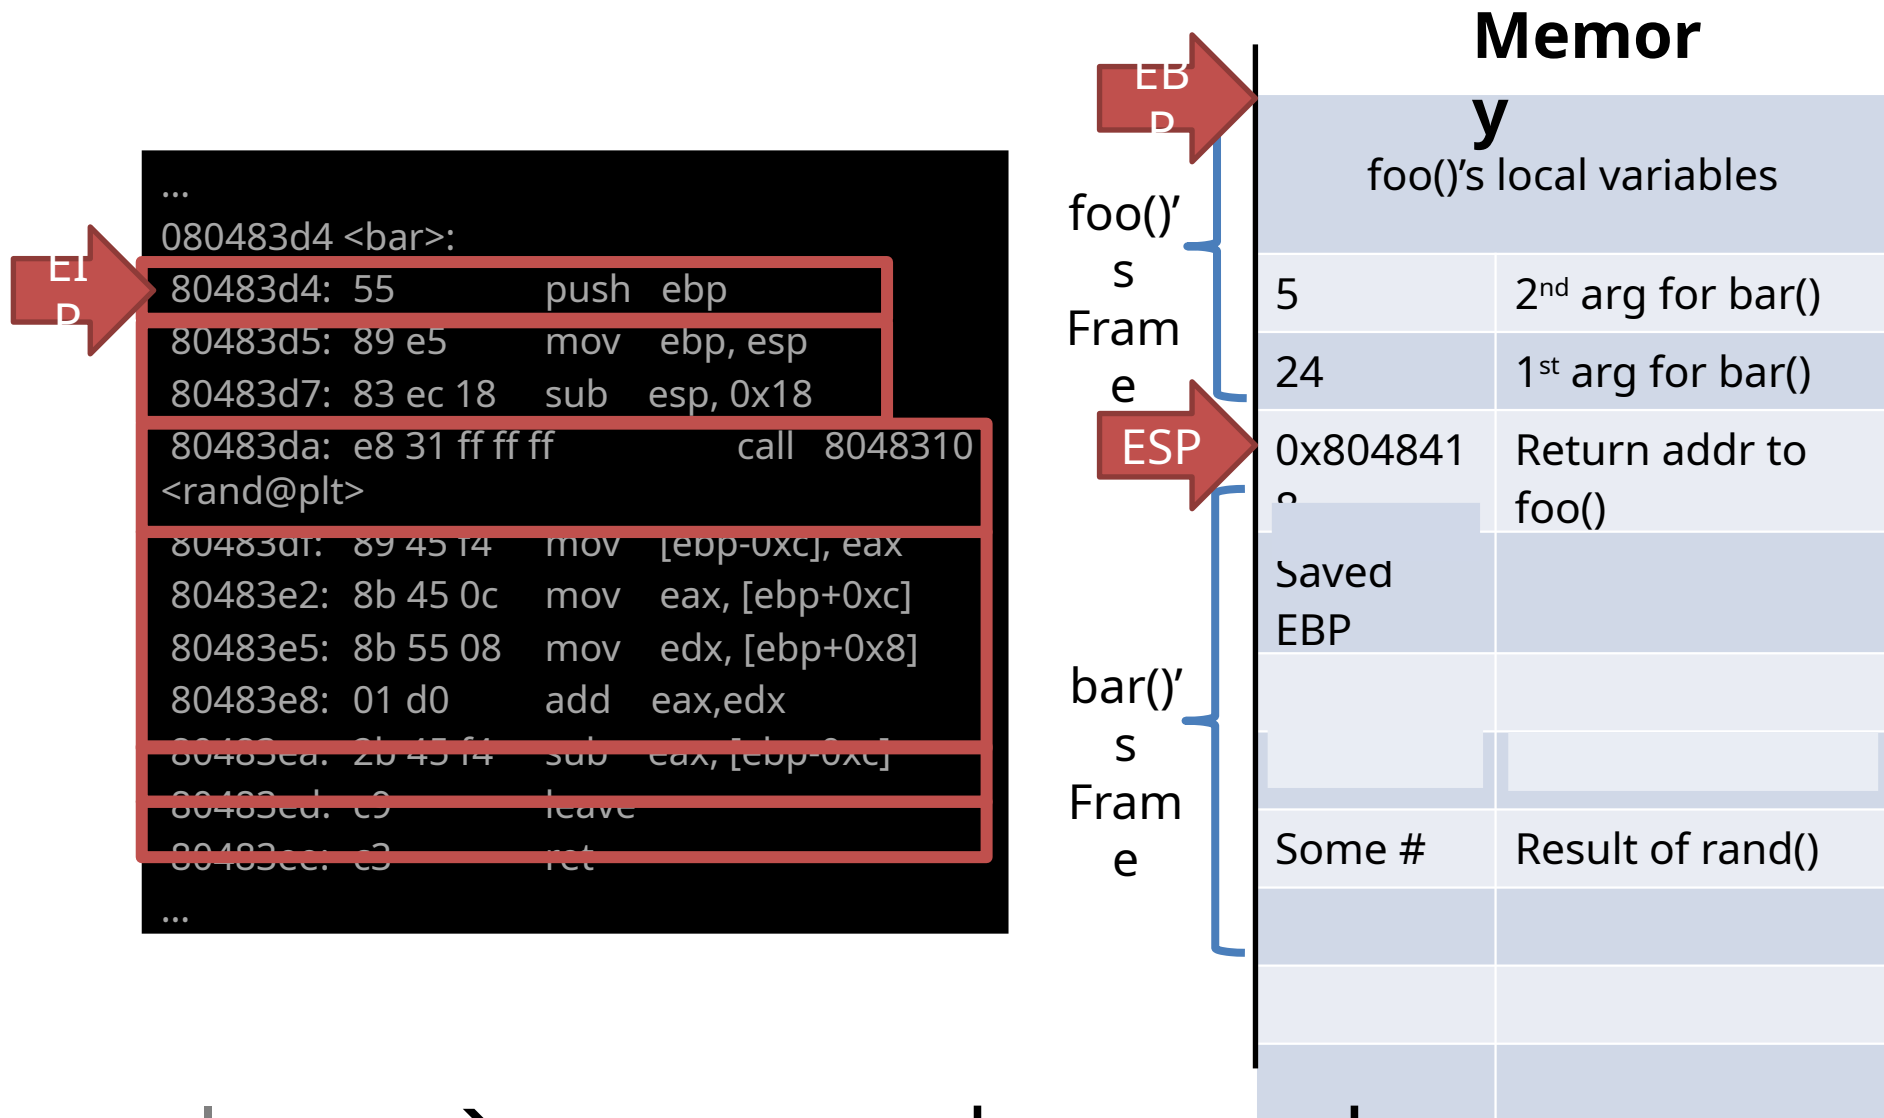
```
8048413: e8 bc ff ff ff      call 80483d4
```

```
<bar>
```

```
8048418: c9                leave
```

```
8048419: c3                ret
```





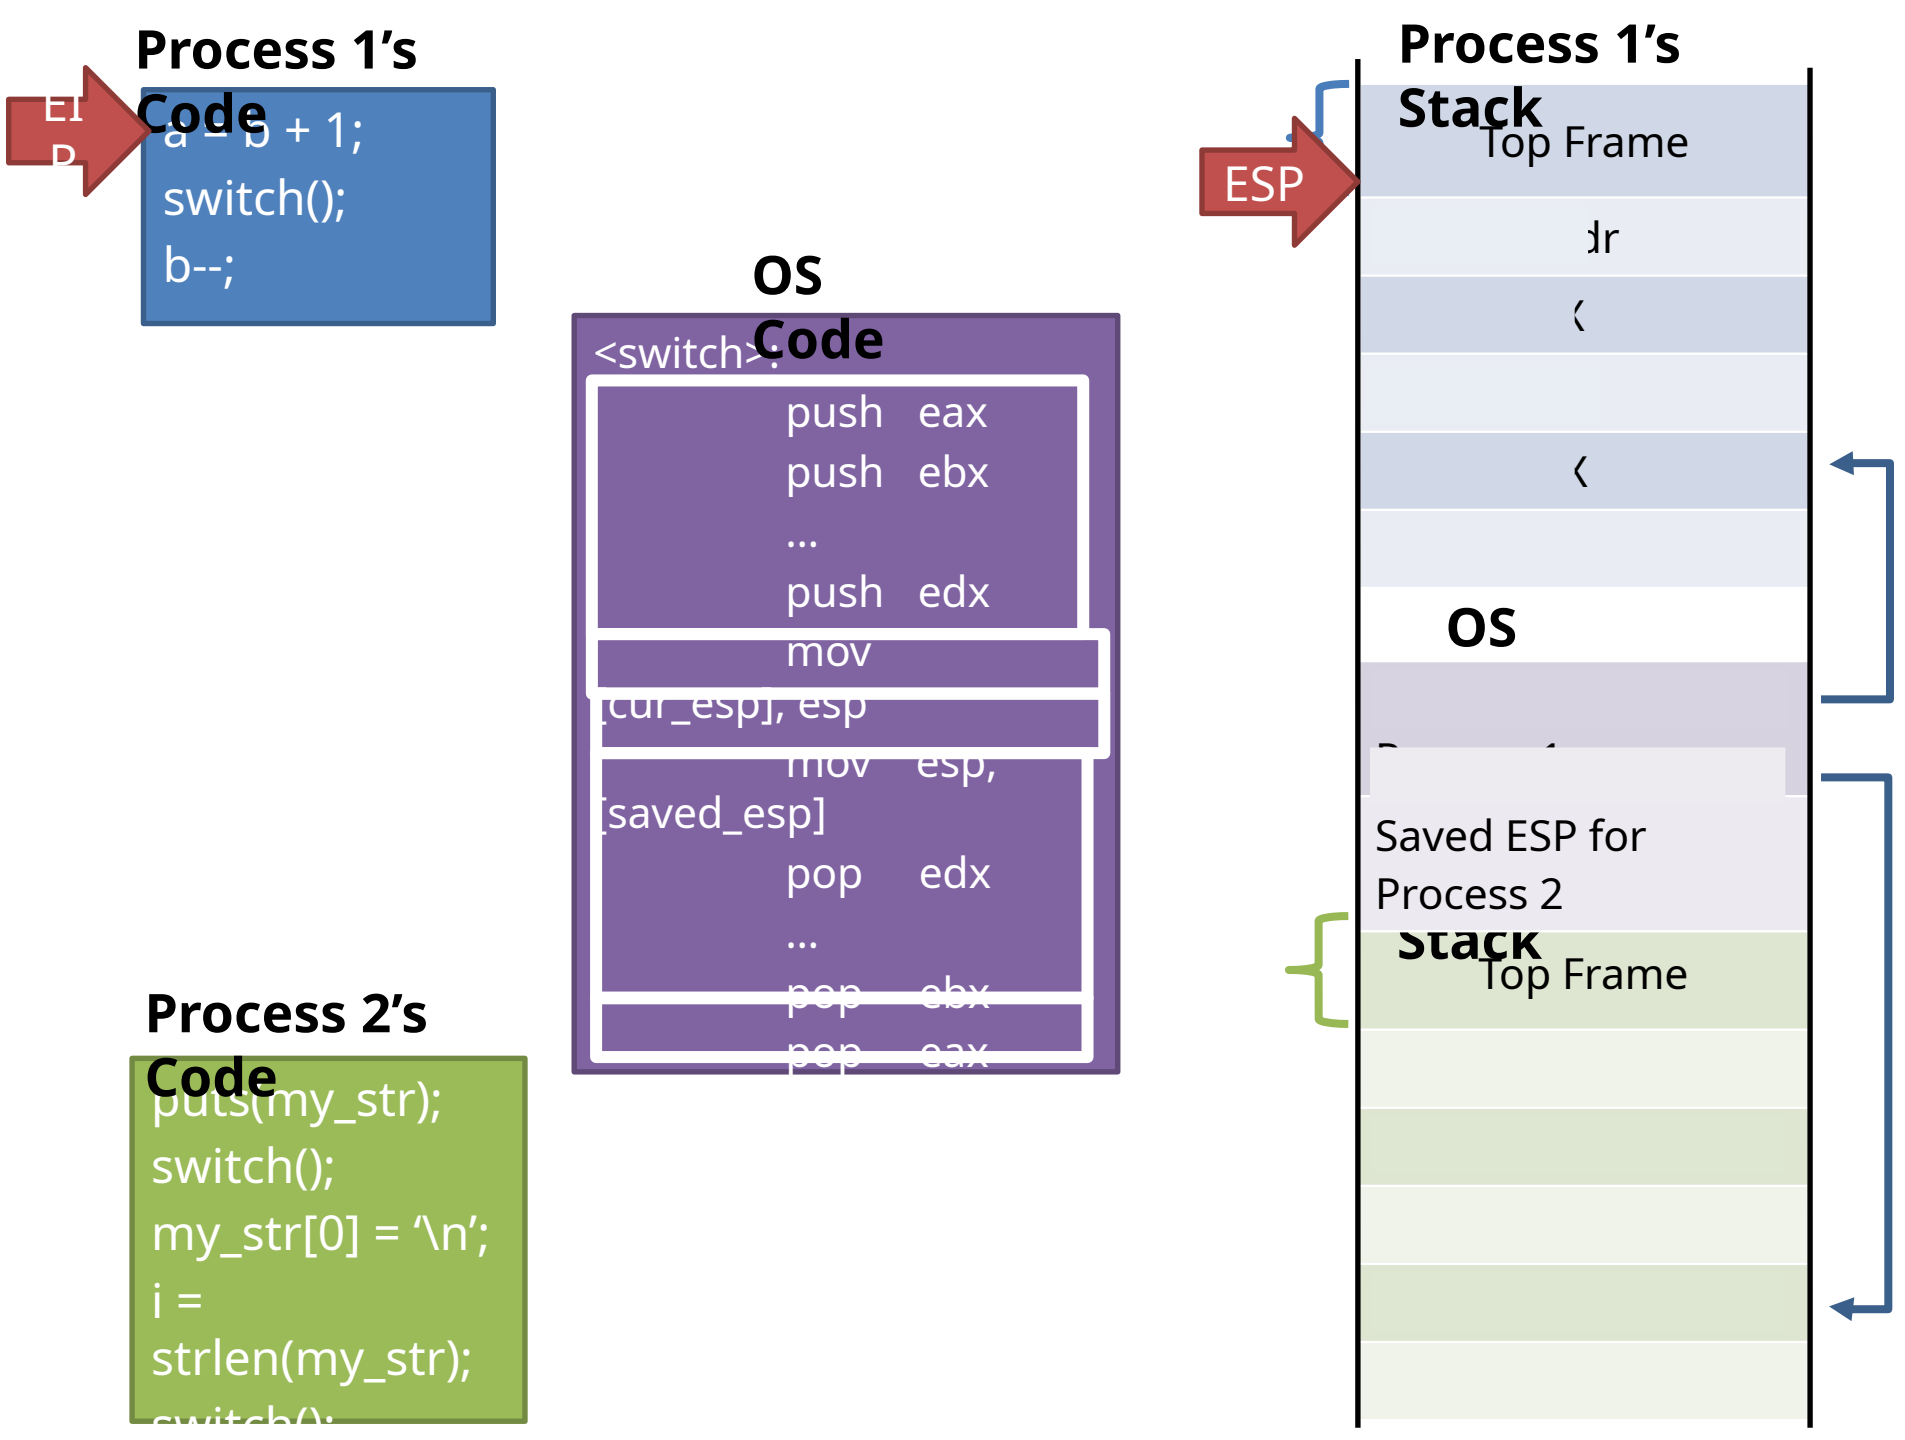
- `leave` → `mov esp, ebp; pop ebp;`
- Return value is placed in EAX

Stack Switching

- We've seen that the stack holds
 - Local variables
 - Arguments to functions
 - Return addresses
 - ... basically, the state of a running program
- Crucially, a process' **control flow** is stored on the stack
- If you modify the stack, you also modify control flow
 - Stack switching is effectively process switching

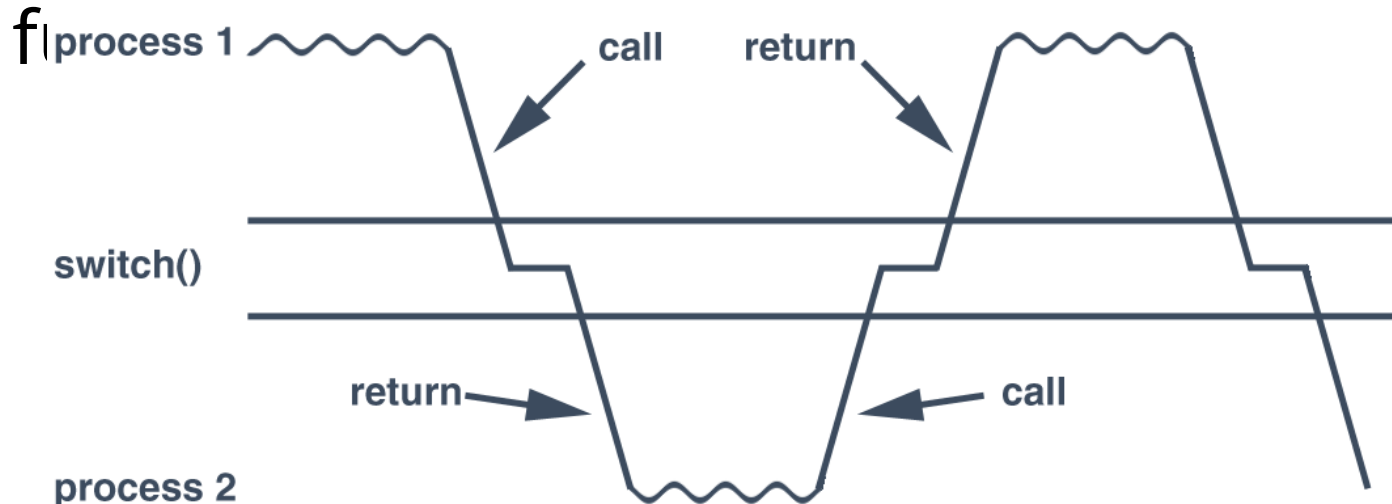
Switching Between Processes

1. Process 1 calls into `switch()` routine
2. CPU registers are pushed onto the stack
3. The stack pointer is saved into memory
4. The stack pointer for process 2 is loaded
5. CPU registers are restored
6. `switch()` returns back to process 2



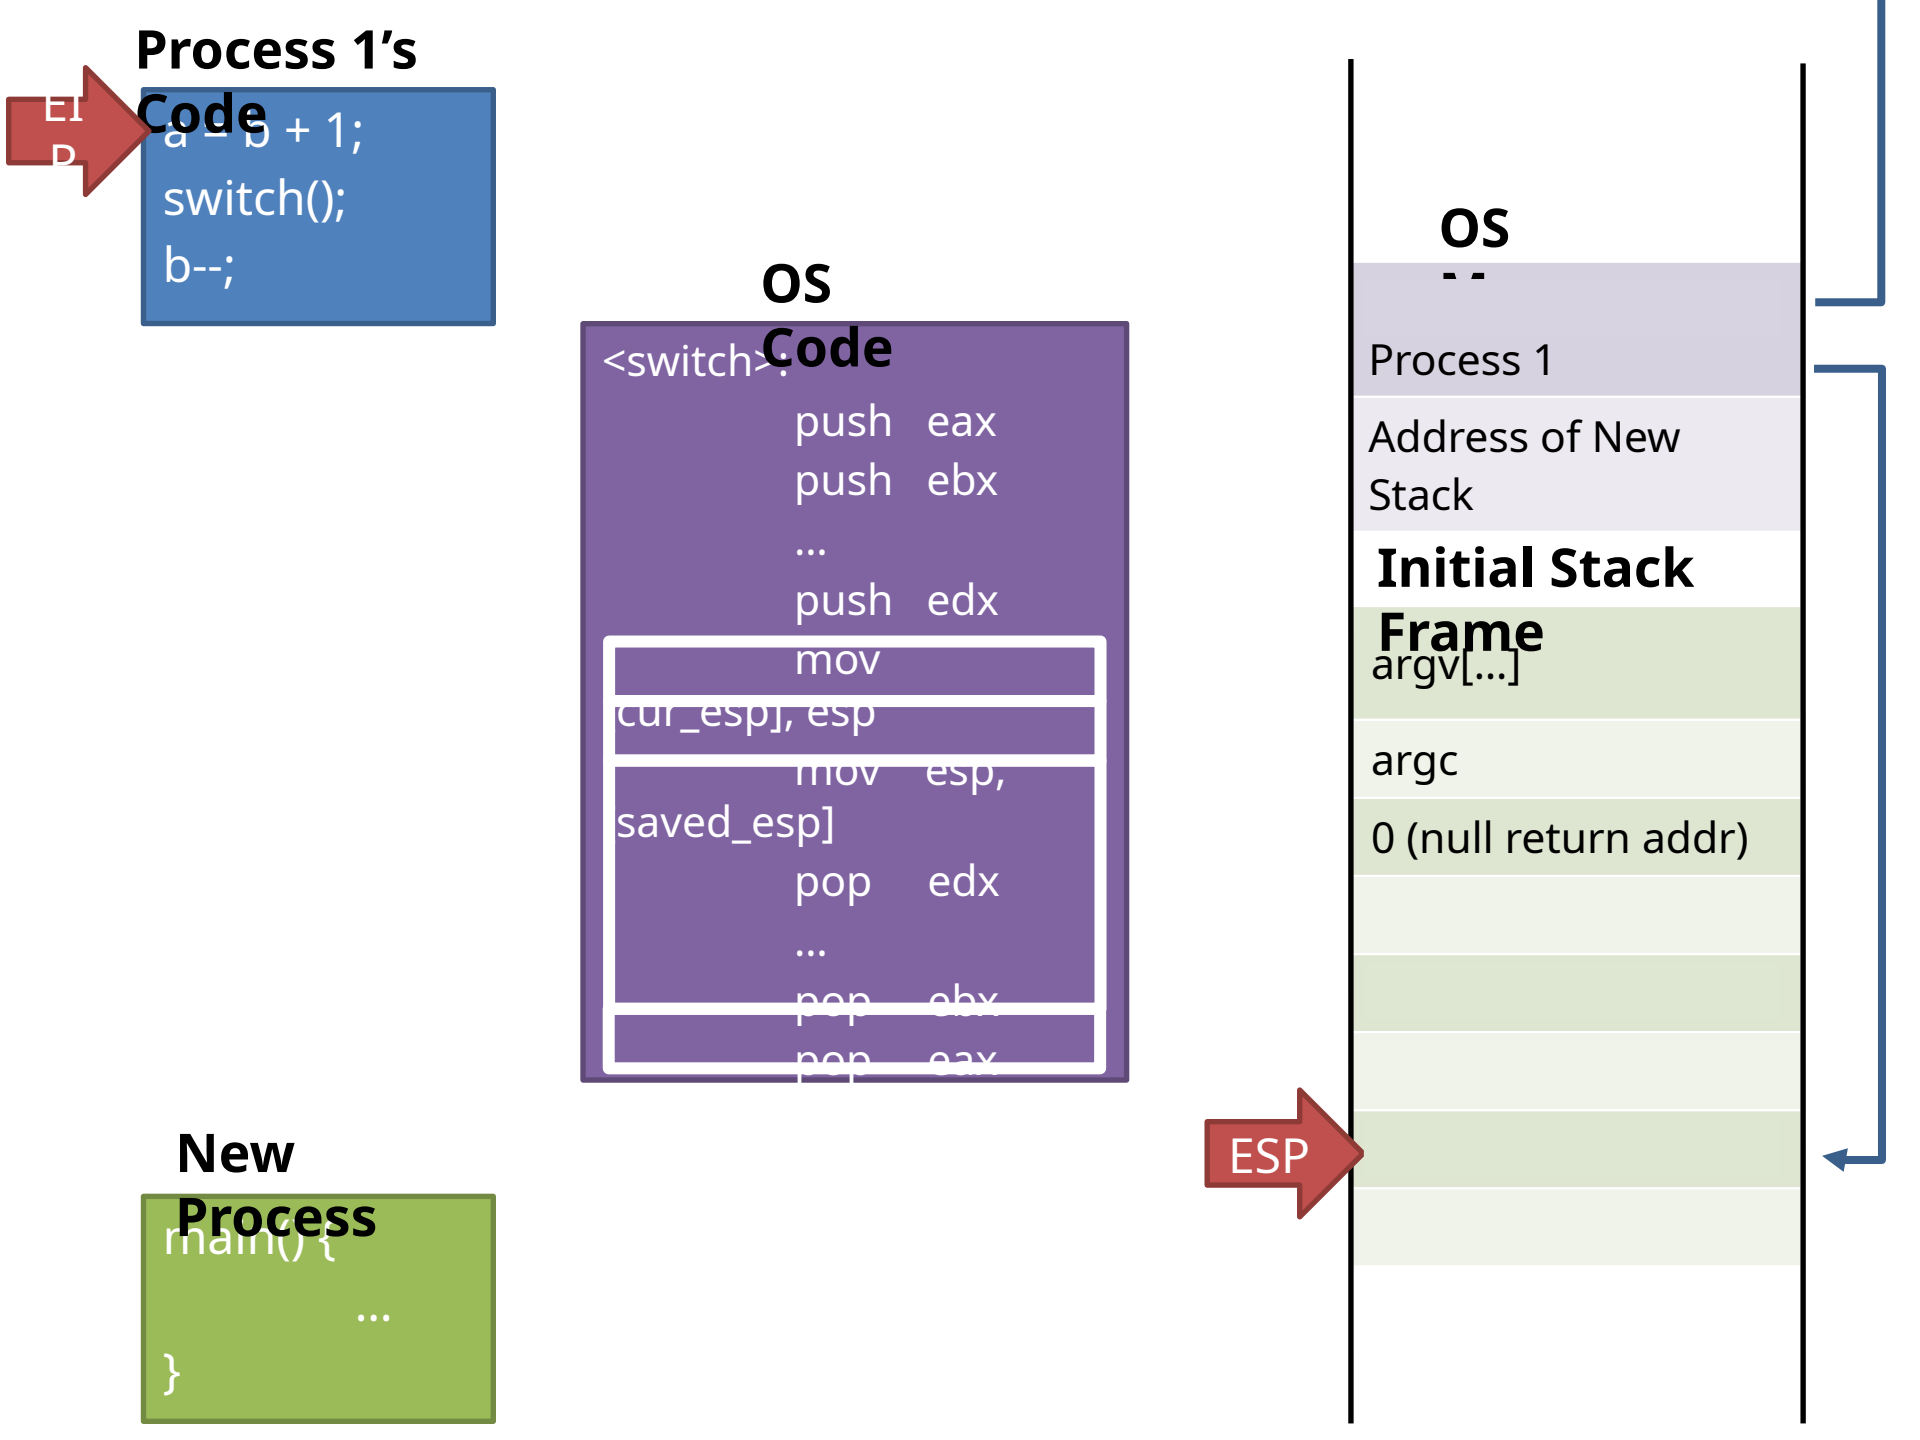
Abusing Call and Return

- Context switching uses function call and return mechanisms
 - Switches **into** a process by **returning** from a function
 - Switches **out** of a process by **calling** into a



What About New Processes?

- But how do you start a process in the first place?
 - A new process doesn't have a stack...
 - ... and it never called into switch()
- Pretend that there *was* a previous call
 - Build a fake initial stack frame
 - This frame looks exactly like the instruction just before main() called into switch()
 - When switch() returns, it'll allow main() to run from the beginning



When Do You Switch Processes?

- To share CPU between multiple processes, control must eventually return to the OS
 - When should this happen?
 - What mechanisms implements the switch from user process back to the OS?
- Four approaches:
 1. Voluntary yielding
 2. Switch during API calls to the OS
 3. Switch on I/O
 4. Switch based on a timer interrupt

Voluntary Yielding

- Idea: processes must voluntarily give up control by calling an OS API, e.g. `thread_yield()`
- Problems:
 - Misbehaving or buggy apps may never yield
 - No guarantee that apps will yield in a reasonable amount of time
 - Wasteful of CPU resources, i.e. what if a process is idle-waiting on I/O?

Interjection on OS APIs

- Idea: whenever a process calls an OS API, this gives the OS an opportunity to context switch
 - E.g. `printf()`, `fopen()`, `socket()`, etc...
- The original Apple Macintosh used this approach
 - Cooperative multi-tasking
- Problems:
 - Misbehaving or buggy apps may never yield
 - Some normal apps don't use OS APIs for long periods of time
 - E.g. a long, CPU intensive matrix calculation

I/O Context Switch Example

- What's happening here?

```
struct terminal {
    queue<char> keystrokes; /* buffered keystrokes - array or list */
    process *waiting;      /* process waiting for input */
    ...
};
process *current;          /* the currently running process */
queue<process *> active;    /* linked list of other processes ready to run */

char get_char(struct terminal *term) {
    if (term->keystrokes.empty()) {
        term->waiting = current; /* sleep waiting for input */
        switch_to(active.pop_head()); /* and switch to next active process */
    }
    return term->keystrokes.pop_head();
}

void interrupt(struct terminal *term, char key) {
    term->keystrokes.push_tail(key); /* add keystroke to buffer */
    if (term->waiting) {
        active.push_tail(term->waiting); /* and wake up sleeping process */
        term->waiting = NULL;
    }
}
```

Context Switching on I/O

- Idea: when one process is waiting on I/O, switch to another process
 - I/O APIs already go through the OS, so context switching is easy
- Problems:
 - Some apps don't have any I/O for long periods of time

Preemptive Context Switching

- So far, our processes will not switch to another process until some action is taken
 - e.g. an API call or an I/O interrupt
- Idea: use a timer interrupt to force context switching at set intervals
 - Interrupt handler runs at a fixed frequency to measure how long a process has been running
 - If it's been running for some max duration (scheduling quantum), the handler switches to the next process
- Problems:
 - Requires hardware support (a programmable timer)
 - Thankfully, this is built-in to most modern CPUs

- Programs
- Processes
- Context Switching
- Protected Mode Execution
- Inter-process Communication
- Threads

Process Isolation

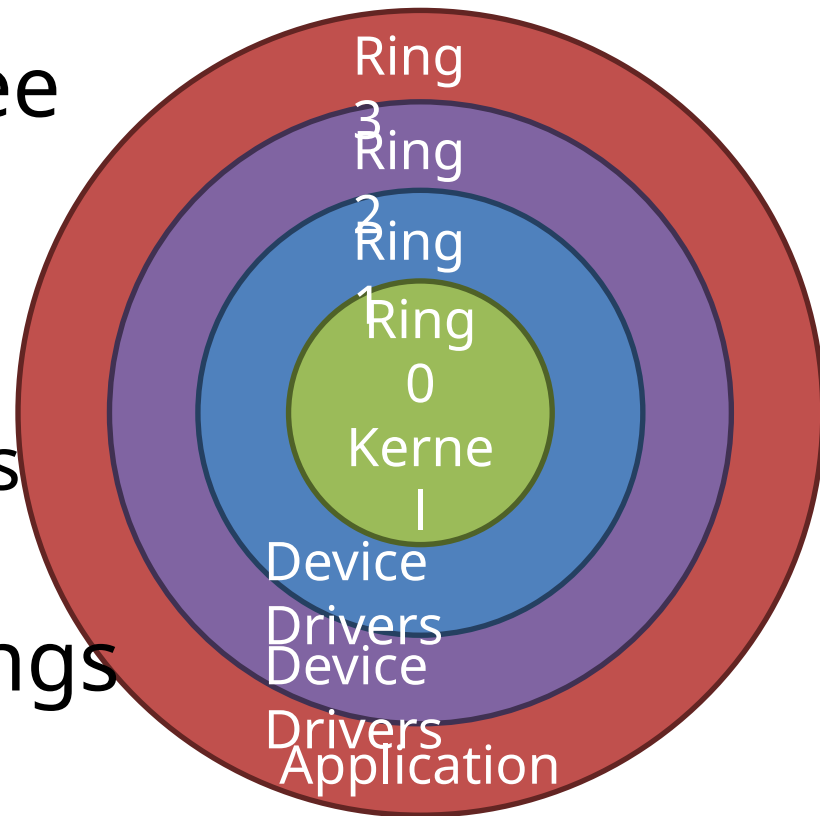
- At this point, we can execute multiple processes concurrently
- Problem: how do you stop processes from behaving badly?
 - Overwriting kernel memory
 - Reading/writing data from other processes
 - Disabling interrupts
 - Crashing the whole computer
 - Etc.

Thought Experiment

- How can we implement execution with limited privilege?
 - Use an interpreter or a simulator
 - Execute each program instruction in a simulator
 - If the instruction is permitted, do the instruction
 - Otherwise, stop the process
 - Basic model in Javascript, Java, ...
- However, interpreters and simulators are slow
- How do we go faster?
 - Run the **unprivileged** code directly on the CPU

Protected Mode

- Most modern CPUs support **protected mode**
- x86 CPUs support three rings with different privileges
 - Ring 0: OS kernel
 - Ring 1, 2: device drivers
 - Ring 3: userland
- Most OSes only use rings 0 and 3
- What about hypervisors?



Real vs. Protected

- On startup, the CPU starts in 16-bit **real** mode
 - Protected mode is disabled
 - Assumes segment:offset addressing
- Typically, bootloader switches CPU to protected mode

```
mov eax, cr0
```

```
or eax, 1 ; set bit 1 of CR0 to 1 to enable  
pmode
```

```
mov cr0, eax
```

Dual-Mode Operation

- Ring 0: kernel/supervisor mode
 - Execution with the full privileges of the hardware
 - Read/write to any memory, access any I/O device, read/write any disk sector, send/read any packet
- Ring 3: user mode or “userland”
 - Limited privileges
 - Only those granted by the operating system kernel

Protected Features

- What system features are impacted by protection?
 - Privileged instructions
 - Only available to the kernel
 - Limits on memory accesses
 - Prevents user code from overwriting the kernel
 - Access to hardware
 - Only the kernel may directly interact with peripherals
 - Programmable Timer Interrupt
 - May only be set by the kernel
 - Used to force context switches between processes

Privileged Instructions

- Examples?
 - sti/cli – Enable and disable interrupts
 - Any instruction that modifies the CR0 register
 - Controls whether protected mode is enabled
 - hlt – Halts the CPU
- What should happen if a user program attempts to execute a privileged instruction?
 - General protection (GP) exception gets thrown by the CPU
 - Control is transferred to the OSes exception handler

Changing Modes

- Applications often need to access the OS
 - i.e. system calls
 - Writing files, displaying on the screen, receiving data from the network, etc...
- But the OS is ring 0, and apps are ring 3
- How do apps get access to the OS?
 - Apps invoke system calls with an interrupt
 - E.g. int 0x80
 - int causes a mode transfer from ring 3 to ring 0

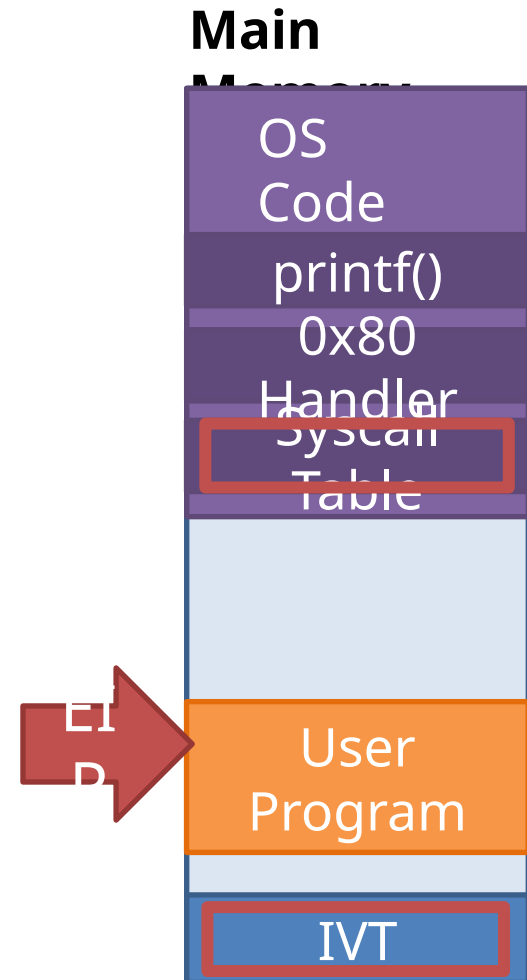
Mode Transfer

Userland
d
Kernel
Mode

1. Application executes trap (`int`) instruction
 - EIP, CS, and EFLAGS get pushed onto the stack
 - Mode switches from ring 3 to ring 0
2. Save the state of the current process
 - Push EAX, EBX, ..., etc. onto the stack
3. Locate and execute the correct syscall handler
4. Restore the state of process
 - Pop EAX, EBX, ... etc.
5. Place the return value in EAX
6. Use `iret` to return to the process
 - Switches back to the original mode (typically 3)

System Call Example

1. Software executes `int 0x80`
 - Pushes EIP, CS, and EFLAGS
2. CPU transfers execution to the OS handler
 - Look up the handler in the IVT
 - Switch from ring 3 to 0
3. OS executes the system call
 - Save the processes state
 - Use EAX to locate the system call
 - Execute the system call
 - Restore the processes state
 - Put the return value in EAX
4. Return to the process with `iret`
 - Pops EIP, CS, and EFLAGS
 - Switches from ring 0 to 3



Alternative Syscall Mechanisms

- Thus far, all examples have used `int/iret`
- However, there are other syscall mechanisms on x86
 - `sysenter/sysexit`
 - `syscall/sysret`
- The `sys*` instructions are much faster than `int/iret`
 - Jump directly to OS code, rather than looking up handlers in the IVT
 - Used by modern OSes, including the Linux kernel

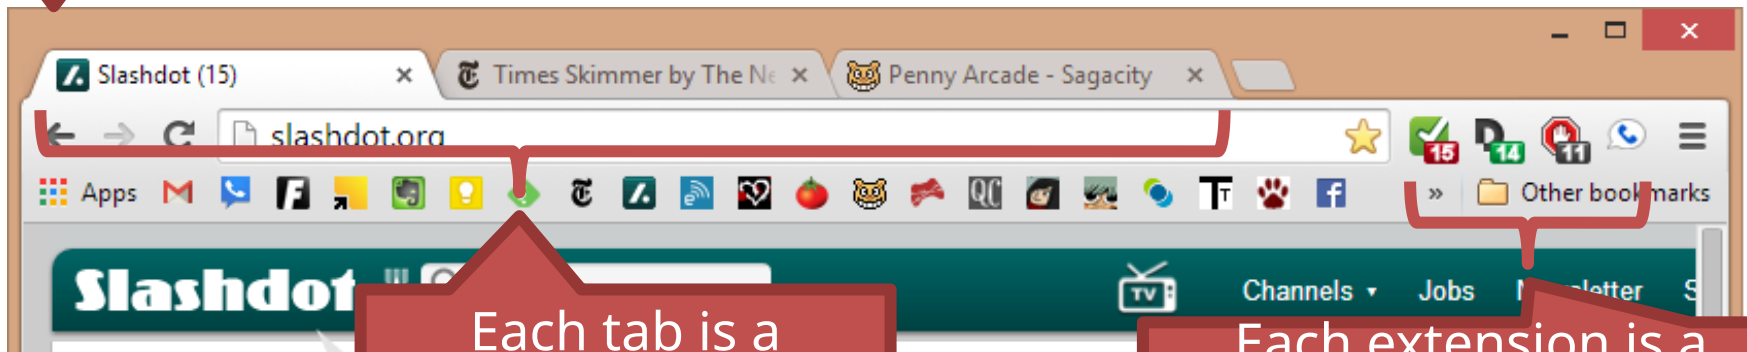
- Programs
- Processes
- Context Switching
- Protected Mode Execution
- **Inter-process
Communication**
- Threads

Processes are not Islands

- Thus far:
 - We can load programs as processes
 - We can context switch between processes
 - Processes are protected from each other

Browser core
is a process

if one or more processes want to
communicate with each other?



Each tab is a
process

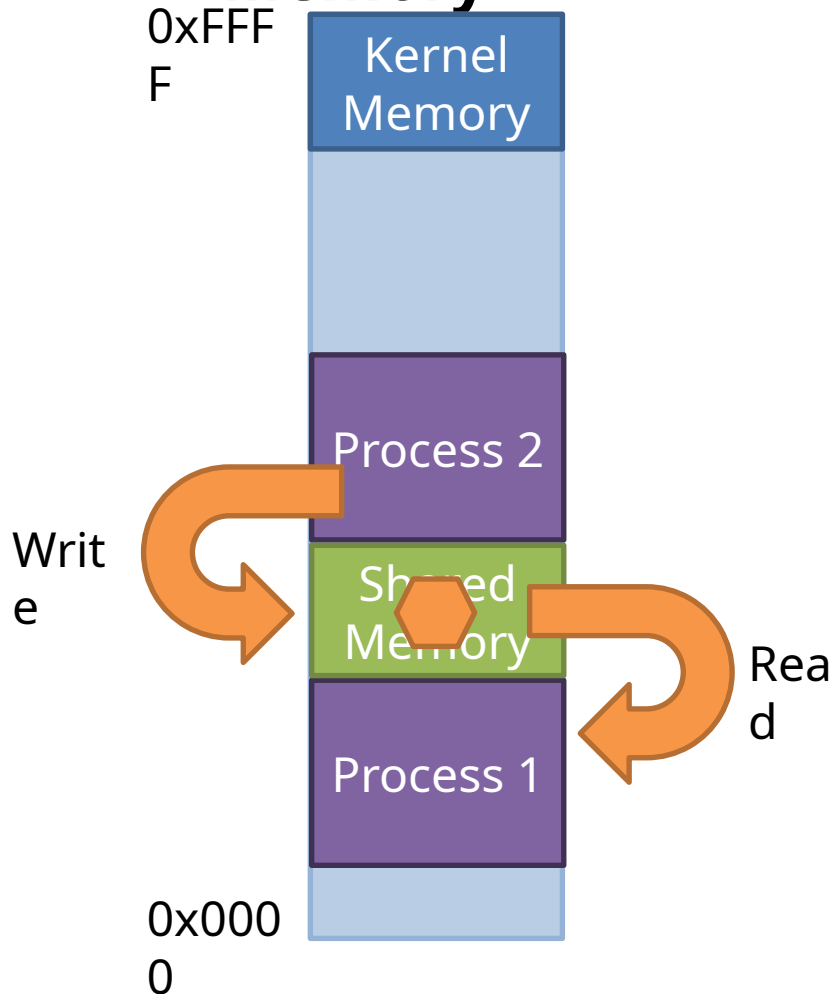
Each extension is a
process

Mechanisms for IPC

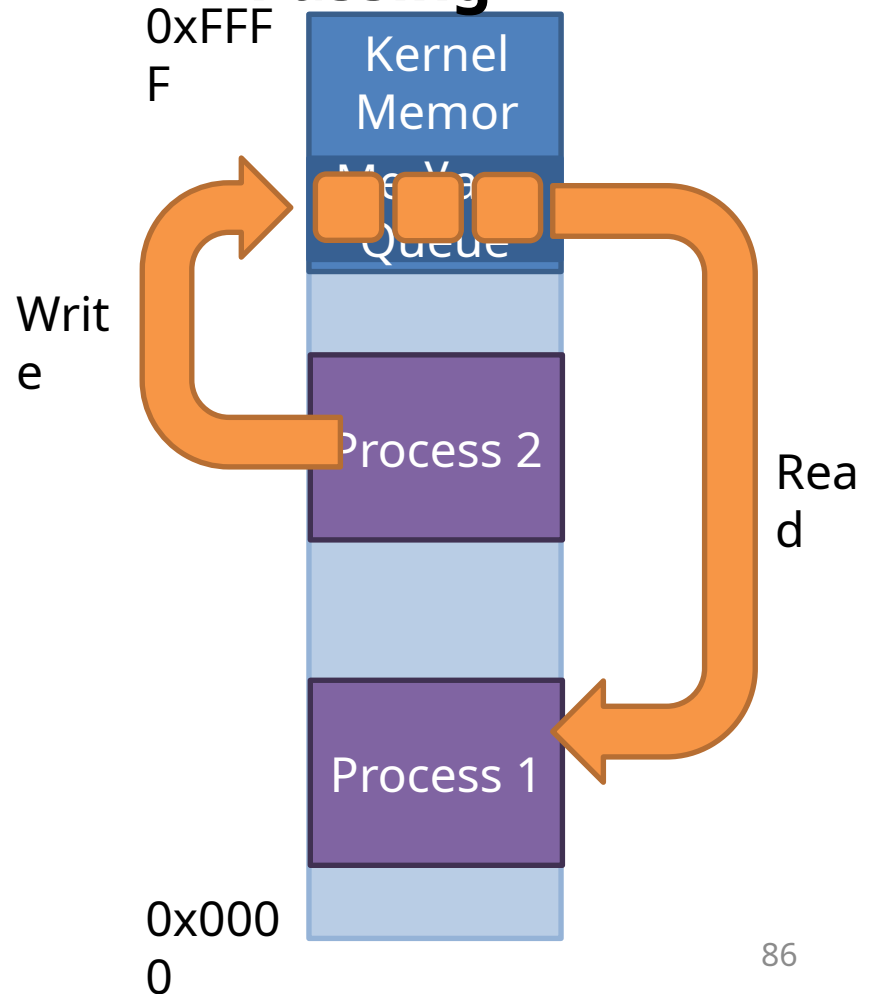
- Typically, two ways of implementing Inter-process communication (IPC)
 - Shared memory
 - A region of memory that many processes can all read/write
 - Message passing
 - Various OS-specific APIs
 - Pipes
 - Sockets
 - Signals

IPC Examples

Shared Memory



Message Passing



Posix Shared Memory API

- `shm_open()` – create and/or open a shared memory page
 - Returns a file descriptor for the shared page
- `ltruncate()` or `ftruncate()` – limit the size of the shared memory page
- `mmap()` – map the memory page into the processes address space
 - Now you can read/write the page using a pointer
- `close()` – close a file descriptor
- `shm_unlink()` – remove a shared page
 - Processes with open references may still access the page

```
/* Program to write some data in shared memory */
int main() {
    const int SIZE = 4096; /* size of the shared page */
    /* name of the shared page */
    const char * NAME = "MY_PAGE";
    const char * msg = "Hello World!";
    int shm_fd;
    char * ptr;

    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    ftruncate(shm_fd, SIZE);
    ptr = (char *) mmap(0, SIZE, PROT_WRITE,
        MAP_SHARED, shm_fd, 0);
    sprintf(ptr, "%s", msg);
    close(shm_fd);
    return 0;
}
```

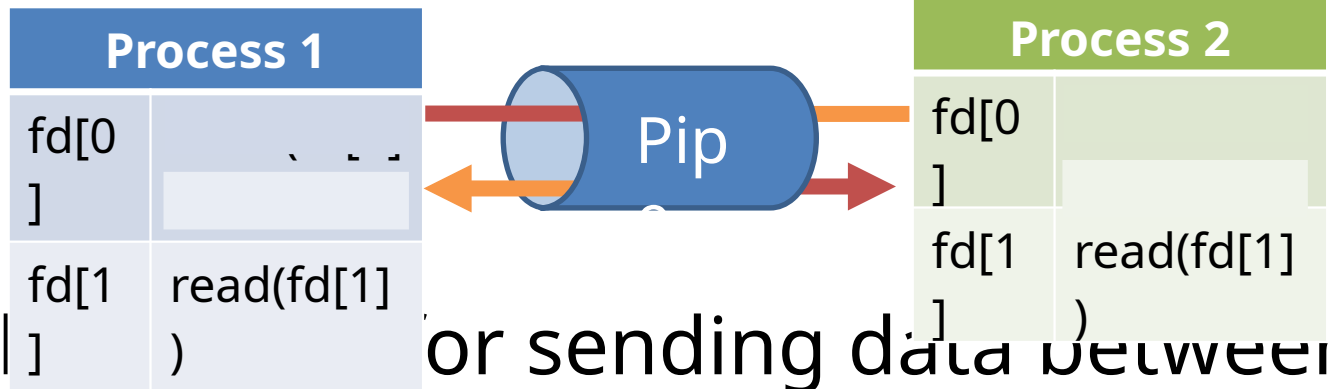


```
/* Program to read some data from shared memory */  
int main() {  
    const int SIZE = 4096; /* size of the shared page */  
        /* name of the shared page */  
    const char * NAME = "MY_PAGE";  
    int shm_fd;  
    char * ptr;  
  
    shm_fd = shm_open(name, O_RDONLY, 0666);  
    ptr = (char *) mmap(0, SIZE, PROT_READ,  
        MAP_SHARED, shm_fd, 0);  
    printf("%s\n", ptr);  
    shm_unlink(shm_fd);  
    return 0;  
}
```

POSIX Message Queues

- Implementation of message passing
 - Producers add messages to a shared FIFO queue
 - Consumer(s) remove messages
 - OS takes care of memory management, synchronization
- Posix API:
 - `msgget()` – creates a new message queue
 - `msgsnd()` – pushes a message onto the queue
 - `msgrcv()` – pops a message from the queue

Pipes



- File-like abstraction for sending data between processes
 - Can be read or written to, just like a file
 - Permissions controlled by the creating process
- Two types of pipes
 - Named pipe: any process can attach as long as it knows the name
 - Typically used for long lived IPC
 - Unnamed/anonymous pipe: only exists between a parent and its children
- Full or half-duplex
 - Can one or both ends of the pipe be read?
 - Can one or both ends of the pipe be written?

You've All Used Pipes

Pipe the output from one process to the input of another process

```
$ ps x | grep ssh  
3299 ?      S      0:00 sshd: cbw@pts/0
```

```

int main() { /* Program that passes a string to a child process through a pipe
*/
    int fd[2], nbytes;
    pid_t childpid;
    char string[] = "Hello, world!\n";
    char readbuffer[80];

    pipe(fd);
    if ((childpid = fork()) == -1) { perror("fork"); exit(1); }
    if (childpid == 0) {
        /* Child process closes up input side of pipe */
        close(fd[0]);
        /* Send "string" through the output side of pipe */
        write(fd[1], string, strlen(string) + 1);
    } else {
        /* Parent process closes up output side of pipe */
        close(fd[1]);
        /* Read in a string from the pipe */
        nbytes = read(fd[0], readbuffer, sizeof(readbuffer));
        printf("Received string: %s", readbuffer);
    }
    return(0);
}

```

Sockets for IPC

- Yes, the same sockets you use for networking
- Server opens a listen socket, as usual
- Clients connect to this socket
 - The server can check the clients IP and drop connections from anyone other than 127.0.0.1
- Send and receive packets as usual

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- What is the capacity of each link?
- Are messages fixed size or variable size?
- Is the link unidirectional or bidirectional?
- Is the link synchronous or asynchronous?
- Does the API guarantee atomicity?
- What is the overhead of the API?

- Programs
- Processes
- Context Switching
- Protected Mode Execution
- Inter-process Communication
- **Threads**

Are Processes Enough?

- At this point, we have the ability to run processes
 - And processes can communicate with each other
- Is this enough functionality?
- Possible scenarios:
 - A large server with many clients
 - A powerful computer with many CPU cores

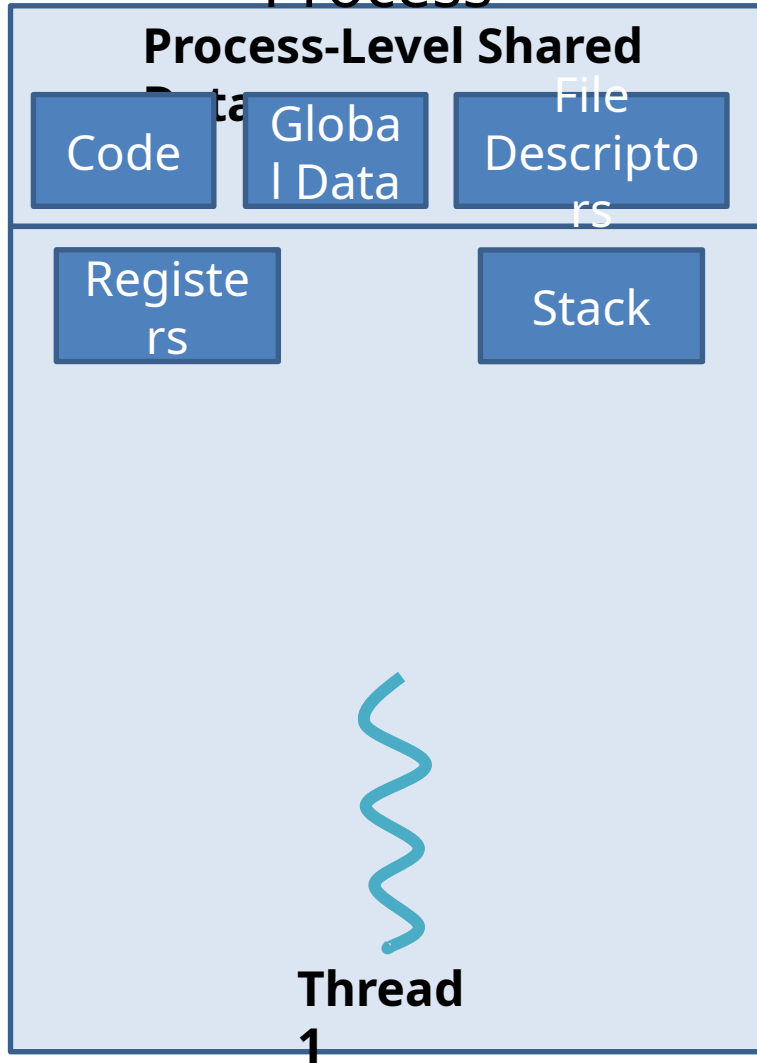
Problems with Processes

- Process creation is heavyweight (i.e. slow)
 - Space must be allocated for the new process
 - `fork()` copies all state of the parent to the child
- IPC mechanisms are cumbersome
 - Difficult to use fine-grained synchronization
 - Message passing is slow
 - Each message may have to go through the kernel

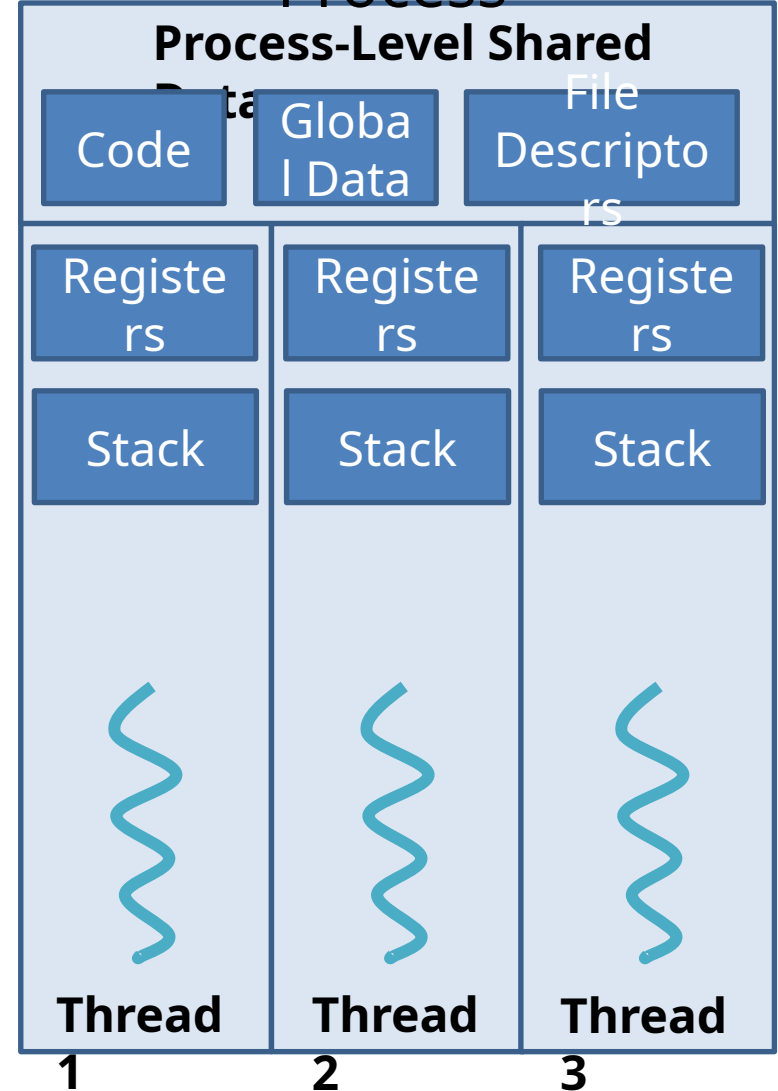
Threads

- Light-weight processes that share the same memory and state space
- Every process has at least one thread
- Benefits:
 - Resource sharing, no need for IPC
 - Economy: faster to create, faster to context switch
 - Scalability: simple to take advantage of multi-core CPUs

Single-Threaded Process



Multi-Threaded Process



Thread Implementations

- Threads can be implemented in two ways:
 1. User threads
 - User-level library manages threads within a single process
 2. Kernel threads
 - Kernel manages threads for all processes

POSIX Pthreads

- POSIX standard API for thread creation
 - IEEE 1003.1c
 - *Specification, not implementation*
 - Defines the API and the expected behavior
 - ... but not how it should be implemented
- Implementation is system dependent
 - On some platforms, user-level threads
 - On others, maps to kernel-level threads

Pthread API

- `pthread_attr_init()` – initialize the threading library
- `pthread_create()` – create a new thread
- `pthread_exit()` – exit the current thread
- `pthread_join()` – wait for another thread to exit
- Pthreads also contains a full range of synchronization primitives

Pthread Example

```
pthread_t tid; // id of the child thread
pthread_attr_t attr; // initialization data
pthread_attr_init(&attr);
pthread_create(&tid, &attr, runner, 0);
pthread_join(tid, 0);
```

```
void * runner(void * params) {
    ...
    pthread_exit(0);
}
```


Linux Threads

- In the kernel, threads are just tasks
 - Remember the `task_struct` from earlier?
- New threads created using the `clone()` API
 - Sort of like `fork()`
 - Creates a new child task that copies the address space of the parent
 - Same code, same environment, etc.
 - New stack is allocated
 - No memory needs to be copied (unlike `fork()`)

Thread Oddities

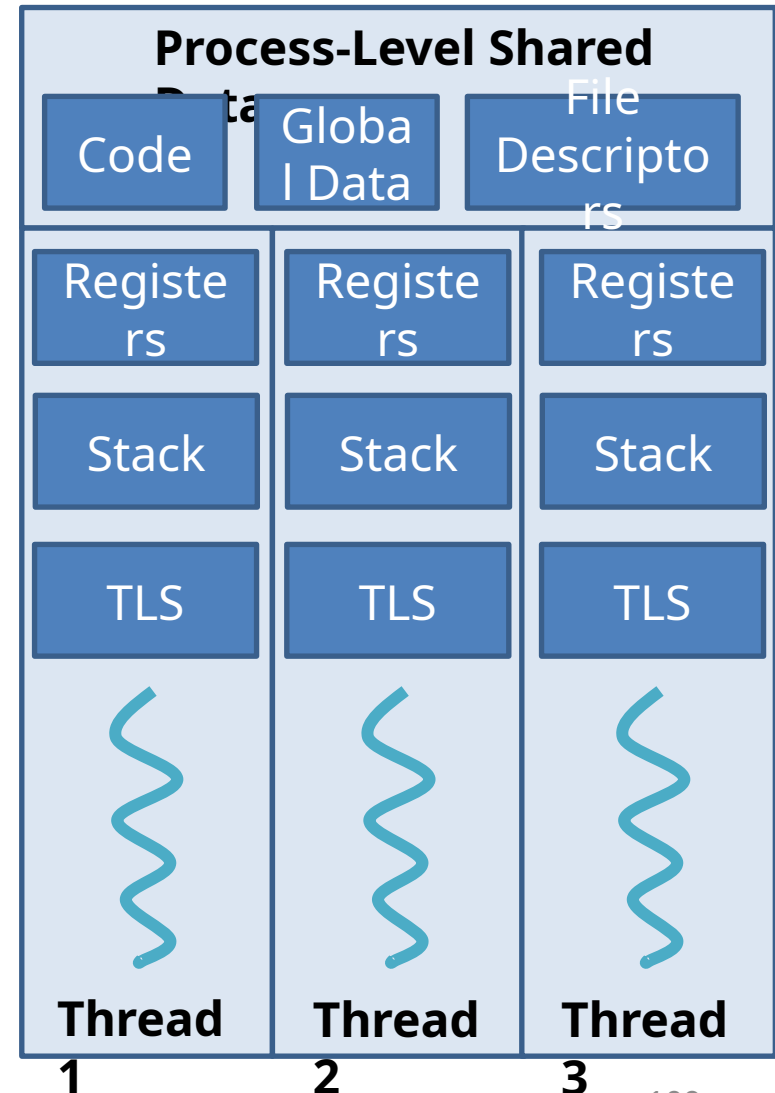
- What happens if you `fork()` a process that has multiple threads?
 - You get a child process with exactly one thread
 - Whichever thread called `fork()` survives
- What happens if you run `exec()` in a multi-threaded process?
 - All but one threads are killed
 - `exec()` gets run normally

Advanced Threading

- Thread pools:
 - Create many threads in advance
 - Dynamically give work to threads from the pool as it becomes available
- Advantages:
 - Cost of creating threads is handled up-front
 - Bounds the maximum number of threads in the process

Thread Local Storage

- Sometimes, you want each thread to have its own “global” data
 - Not global to all threads
 - Not local storage on the stack
- Thread local storage (TLS) allows each thread to have its own space for “global” variables
 - Similar to **static** variables



OpenMP

- Compiler extensions for C, C++ that adds native support for parallel programming
- Controlled with parallel regions
 - Automatically creates as many threads as there

```
#include <omp.h>

int main() {
    int i, N = 20;
    #pragma omp parallel
    {
        printf("I am a parallel region\n");
    }

    # pragma omp parallel for
    for (i = 0; i < N; i++)
        printf("This is a parallel for loop\n");

    return 0;
}
```

Processes vs. Threads

- Threads are better if:
 - You need to create new ones quickly, on-the-fly
 - You need to share lots of state
- Processes are better if:
 - You want protection
 - One process that crashes or freezes doesn't impact the others
 - You need high security
 - Only way to move state is through well-defined, sanitized message passing interface