CS 5600
Computer Systems

Lecture 10: File Systems

What are We Doing Today?

- Last week we talked extensively about hard drives and SSDs
  - How they work
  - Performance characteristics
- This week is all about managing storage
  - Disks/SSDs offer a blank slate of empty blocks
  - How do we store files on these devices, and keep track of them?
  - How do we maintain high performance?
  - How do we maintain consistency in the face of random crashes?

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- Extents and B-Trees (ext4)
- Log-based File Systems

Building the Root File System

- One of the first tasks of an OS during bootup is to build the root file system

1. Locate all bootable media
   - Internal and external hard disks
   - SSDs
   - Floppy disks, CDs, DVDs, USB sticks

2. Locate all the partitions on each media
   - Read MBR(s), extended partition tables, etc.

3. Mount one or more partitions
   - Makes the file system(s) available for access

The Master Boot Record

Address		Description	Size (Bytes)
Hex	**Dec.**		
0x000	0	Bootstrap code area	446
0x1BE	446	Partition Entry #1	16
0x1CE	462	Partition Entry #2	16
0x1DE	478	Partition Entry #3	16
0x1EE	494	Partition Entry #4	16

Includes the starting LBA and length of the partition

Disk 1

Partition 3 (NTFS)

Partition 4 (FAT32)

Disk 2

MBR

Partition 1 (NTFS)

Extended Partitions

- In some cases, you may want >4 partitions
- Modern OSes support extended partitions
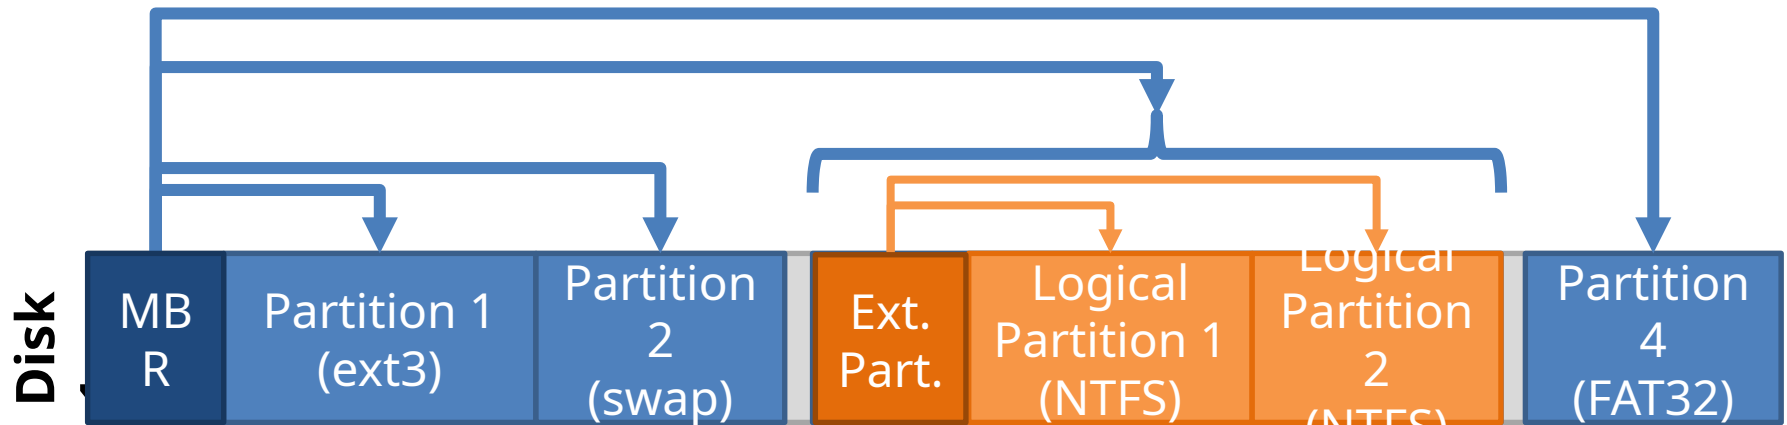
- Extended partitions may use OS-specific partition table formats (meta-data)
  - Thus, other OSes may not be able to read the logical partitions

Types of Root File Systems

```
[cbw@ativ9 ~] df -h
Filesystem  Size Used   Avail Use%   Mounted on
/dev/sda7  39G  14G    23G   38%   /
/dev/sda2296M  48M 249M   16%   /boot/efi
/dev/sda5127G  86G   42G   68%   /media/cbw/Data
/dev/sda4  61G  34G   27G   57%   /media/cbw/Windows
/dev/sdb1 1.9G 352K  1.9G    1%   /media/cbw/NDSS-2019
```

1 drive, 4 partitions

1drive, 1 partition

- Linux has a single root
  - One partition is mounted as /
  - All other partitions are mounted somewhere under /

- Typically, the partition containing the kernel is mounted as / or C:

Mounting a File System

1. Read the super block for the target file system
   – Contains meta-data about the file system
   – Version, size, locations of key structures on disk, etc.

2. Determine the mount point
   – On Windows: pick a drive letter
   – On Linux, mount the new file system under a

```
Filesystem Size Used  Avail Use%  Mounted on
/dev/sda5127G  86G   42G   68%  /media/cbw/Data
/dev/sda4  61G  34G   27G   57%  /media/cbw/Windows
/dev/sdb1 1.9G 352K  1.9G    1%  /media/cbw/NDSS-2013
```
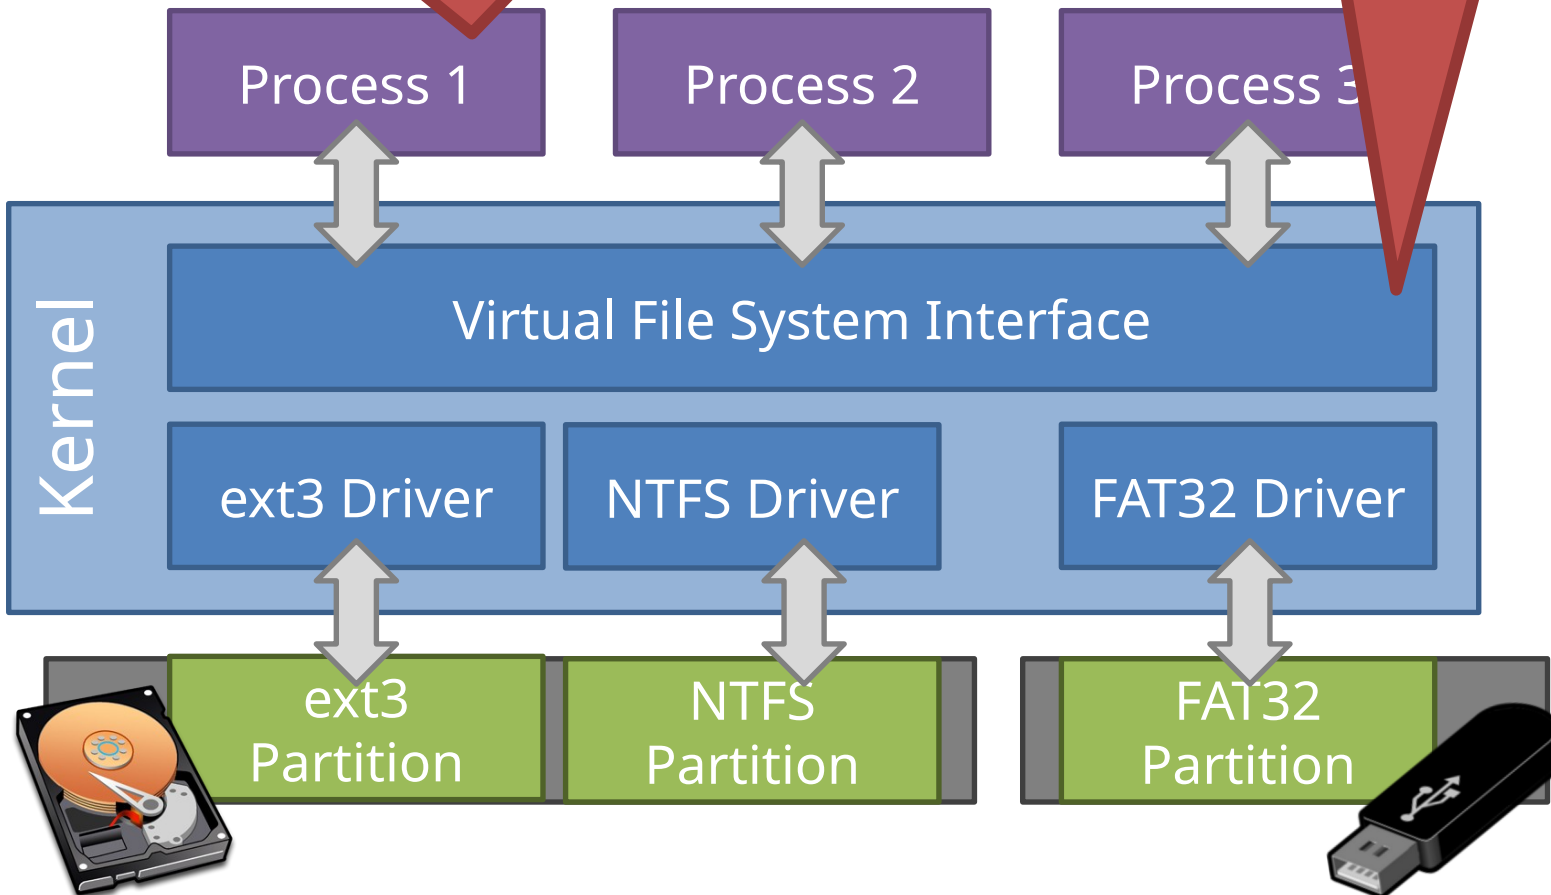
Virtual File System Interface

- Problem: the OS may mount several partitions containing different underlying file systems
  - It would be bad if processes had to use different APIs for different file systems
- Linux uses a Virtual File System interface (VFS)
  - Exposes POSIX APIs to processes
  - Forwards requests to lower-level file system specific drivers
- Windows uses a similar system

9

VFS Flowchart

Processes (usually) don't need to know about low-level file system details

Relatively simple to add additional file system drivers

Process 1

Process 2

Process 3

Kernel

Virtual File System Interface

ext3 Driver

NTFS Driver

FAT32 Driver

ext3 Partition

NTFS Partition

FAT32 Partition

Mount isn't Just for Bootup

- When you plug storage devices into your running system, mount is executed in the background
- Example: plugging in a USB stick
- What does it mean to "safely eject" a device?
  - Flush cached writes to that device
 - Cleanly unmount the file system on that device

11

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- Extents and B-Trees (ext4)
- Log-based File Systems

Status Check

- At this point, the OS can locate and mount partitions
- Next step: what is the on-disk layout of the file system?
  - We expect certain features from a file system
    - Named files
    - Nested hierarchy of directories
    - Meta-data like creation time, file permissions, etc.
  - How do we design on-disk structures that support these features?

The Directory Tree
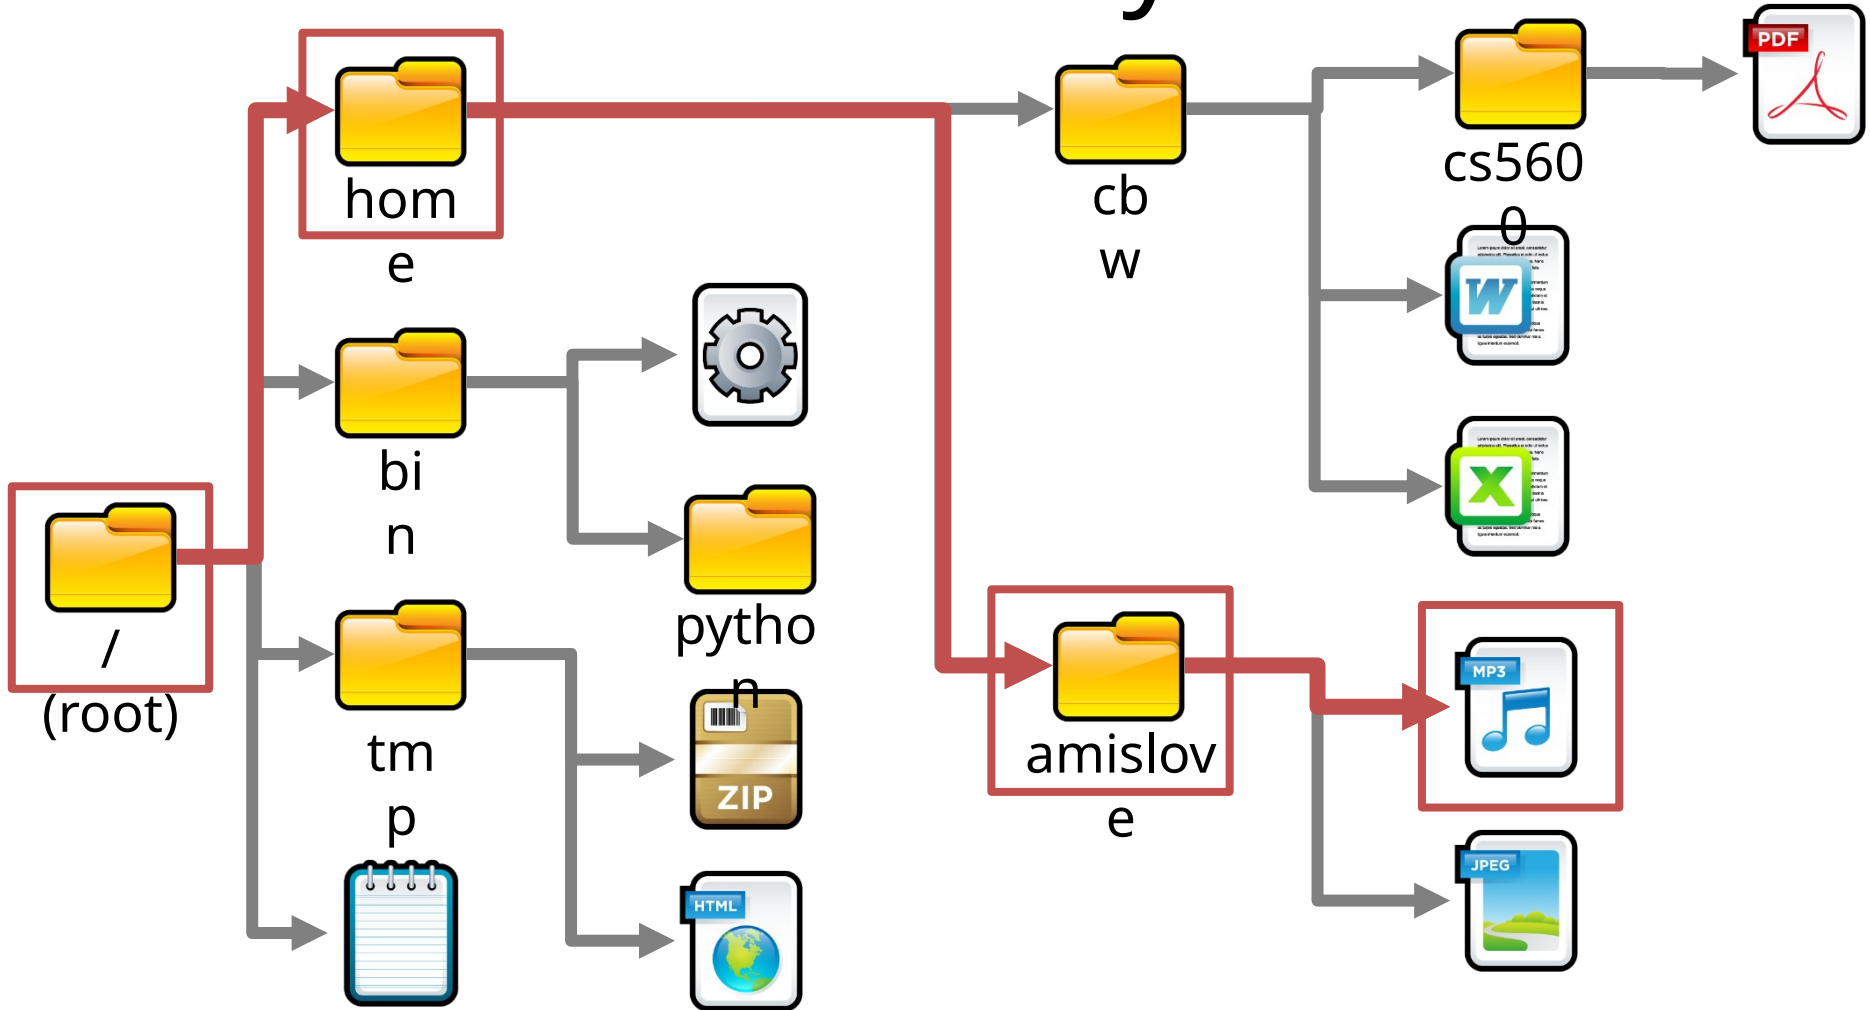
- Navigated using a path
  - E.g. /home/amislove/music.mp3

Absolute and Relative Paths

- Two types of file system paths
  - Absolute
    - Full path from the root to the object
    - Example: /home/cbw/cs5600/hw4.pdf
    - Example: C:\Users\cbw\Documents\
  - Relative
    - OS keeps track of the working directory for each process
    - Path relative to the current working directory
    - Examples [working directory = /home/cbw]:
      - syllabus.docx [→ /home/cbw/syllabus.docx]
      - cs5600/hw4.pdf [→ /home/cbw/cs5600/hw4.pdf]
      - ./cs5600/hw4.pdf [→ /home/cbw/cs5600/hw4.pdf]
      - ../amislove/music.mp3 [→ /home/amislove/music.mp3]

Files

- A file is a composed of two components
  - The file data itself
    - One or more blocks (sectors) of binary data
    - A file can contain anything
  - Meta-data about the file
    - Name, total size
    - What directory is it in?
    - Created time, modified time, access time
    - Hidden or system file?
    - Owner and owner's group
    - Permissions: read/write/execute

16

File Extensions

- File name are often written in dotted notation
  - E.g. program.exe, image.jpg, music.mp3
- A file's extension **does not mean anything**
 - Any file regardless of its name given an extension

  Rename

  Has the data in the file changed from music to an image?

- Graphical shells (like Windows explorer) use extensions to try and match files → programs
  - This mapping may fail for a variety of reasons

More File Meta-Data

- Files have additional meta-data that is not typically shown to users
  - Unique identifier (file names may not be unique)
  - Structure that maps the file to blocks on the disk
- Managing the mapping from files to blocks is one of the key jobs of the file system

Disk

Mapping Files to Blocks

- Every file is composed of >=1 blocks
- Key question: how do we map a file to its blocks?

List of blocks

[1]

[4, 5, 7, 8]

[6]

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |

As (start, length) pairs

(1, 1)

(4, 4)

(9, 1)

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |

- Problem?
  - Really large files

- Problem?
  - Fragmentation
  - E.g. try to add a new file with 3 blocks

Directories

- Traditionally, file systems have used a hierarchical, tree-structured namespace
  - Directories are objects that contain other objects
    - i.e. a directory may (or may not) have children
  - Files are leaves in the tree
- By default, directories contain at least two entries

"..." points the the parents directory

"." self pointer

.

.

/
(root)

bi
n

pytho
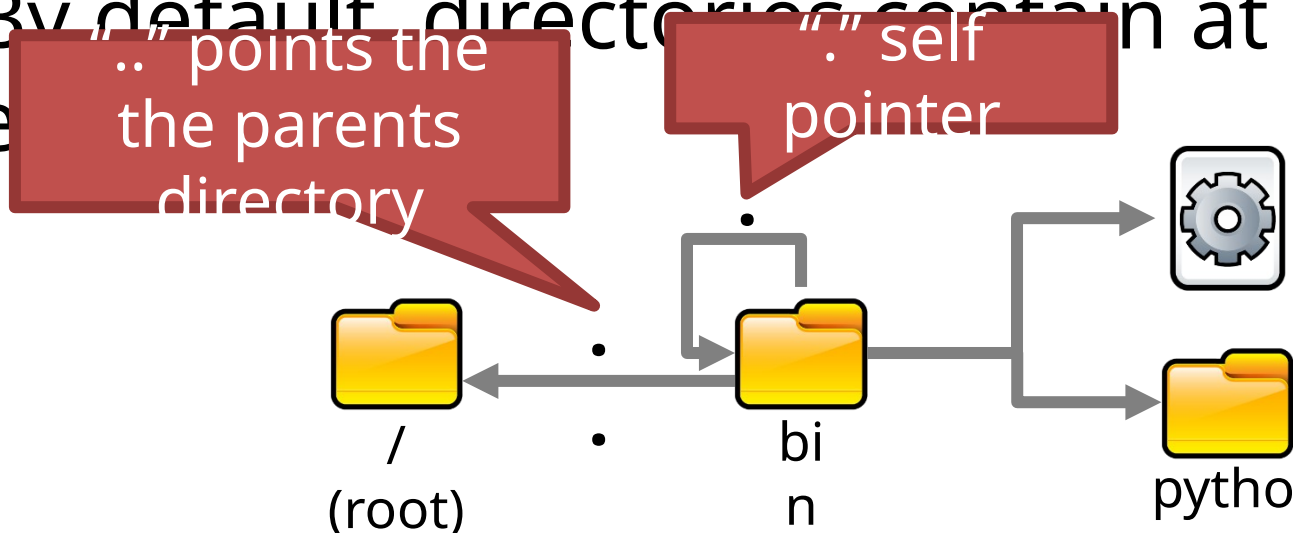
More on Directories

- Directories have associated meta-da
  - Name, number of entries
  - Created time, modified time, access tim
  - Permissions (read/write), owner, and gr
- The file system must encode directories and store them on the disk
  - Typically, directories are stored as a special type of file
  - File contains a list of entries inside the directory, plus some meta-data for each entry

Example Directory File

Name	Index	Dir?	Perms
.	2	Y	rwx
Windows	3	Y	rwx
Users	4	Y	rwx
pagefile.sys	5	N	r

C:\

Windows

Users

pagefile.sys

0 1 2 3 4 5 6 7 8 9

Disk

C:

Directory File Implementation

- Each directory file stores many entries
- Key Question: how do you encode the entries?

N			
.			
w			
U			
pagefile.s ys	5	N	r

| Windows | 3 | Y | rwx |

- Other alternatives: hash tables, B-trees
  - More on B-trees later…
- In practice, implementing directory files is complicated
  - Example: do filenames have a fixed, maximum length or variable length?

- Good: O(1) to add new entries
  - Just append to the file
- Bad: O(n) to search for an entry

- Good: O(log n) to search for an entry
- Bad: O(n) to add new entries
  - Entire file has the be

File Allocation Tables (FAT)

- Simple file system popularized by MS-DOS
  - First introduced in 1977
  - Most devices today use the FAT32 spec from 1996
  - FAT12, FAT16, VFAT, FAT32, etc.
- Still quite popular today
  - Default format for USB sticks and memory cards
  - Used for EFI boot partitions
- Name comes from the index table used to track directories and files

- system
- FAT version, location of boot files
- Total number of blocks
- Index of the root directory in the

- File allocation table (FAT)
- Marks which blocks are free or in-use
- **Linked-list structure** to manage large files

- Store file and directory data
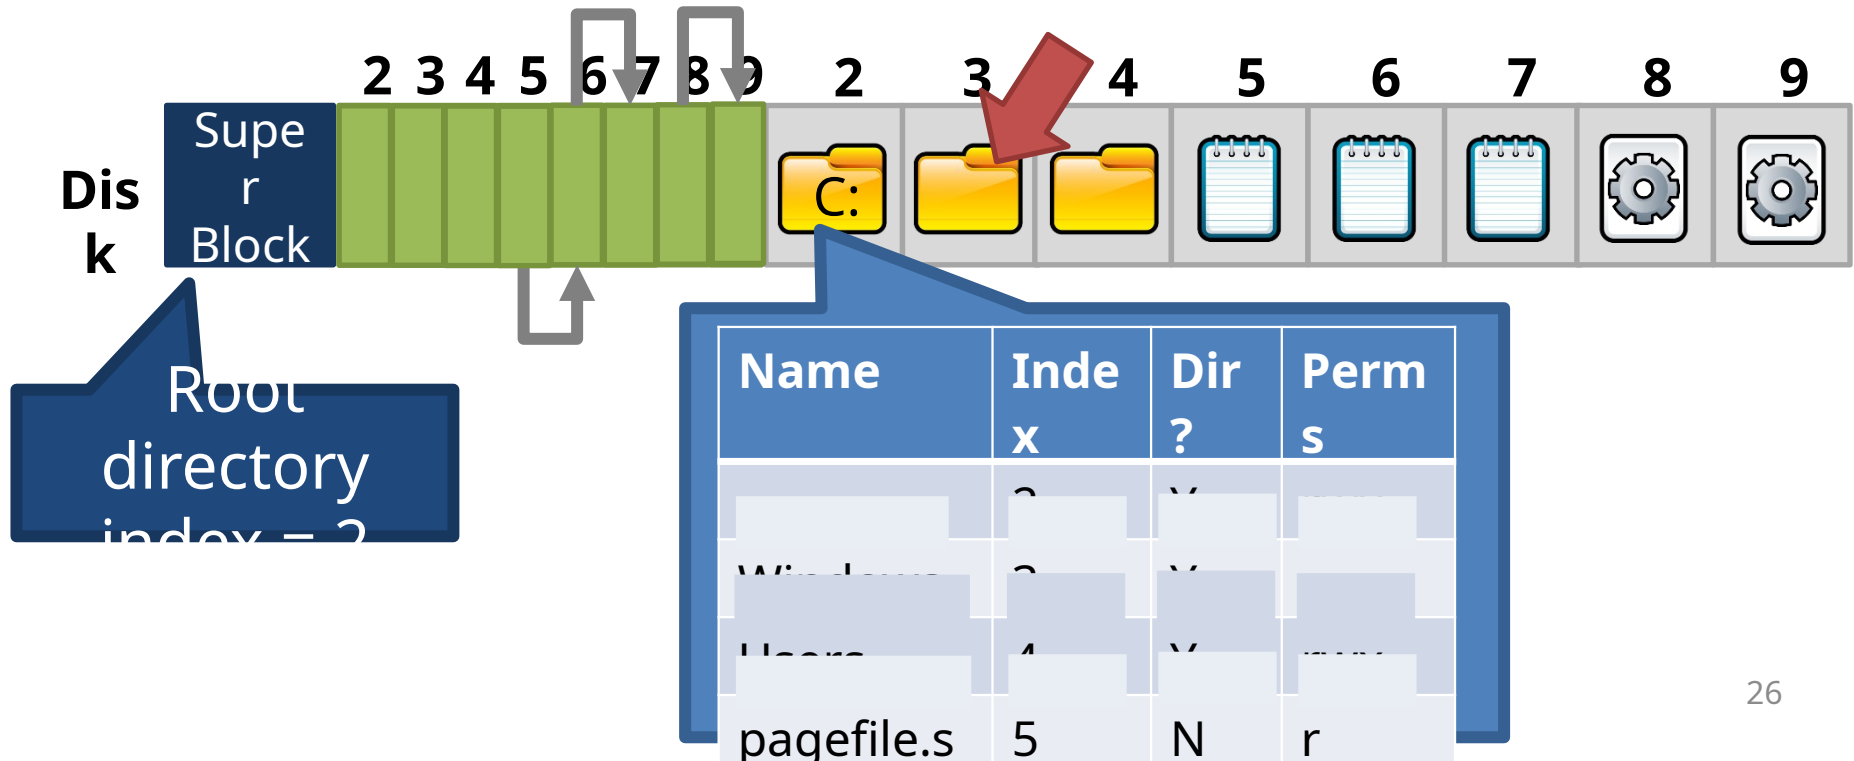- Each block is a fixed size (4KB – 64KB)
- Files may span multiple blocks

Disk

| Super Block | | | | | | | | | | | | | | | | | | | | | |

C:\

Windows

Users

- Directories are special files
  - File contains a list of entries inside the directory
- Possible values for FAT entries:
  - 0 – entry is empty
  - 1 – reserved by the OS
  - 1 < N < 0xFFFF – next block in a chain
  - 0xFFFF – end of a chain

Disk

Super Block

2 3 4 5 6 7 8 9 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9

C:

Root directory index = 2

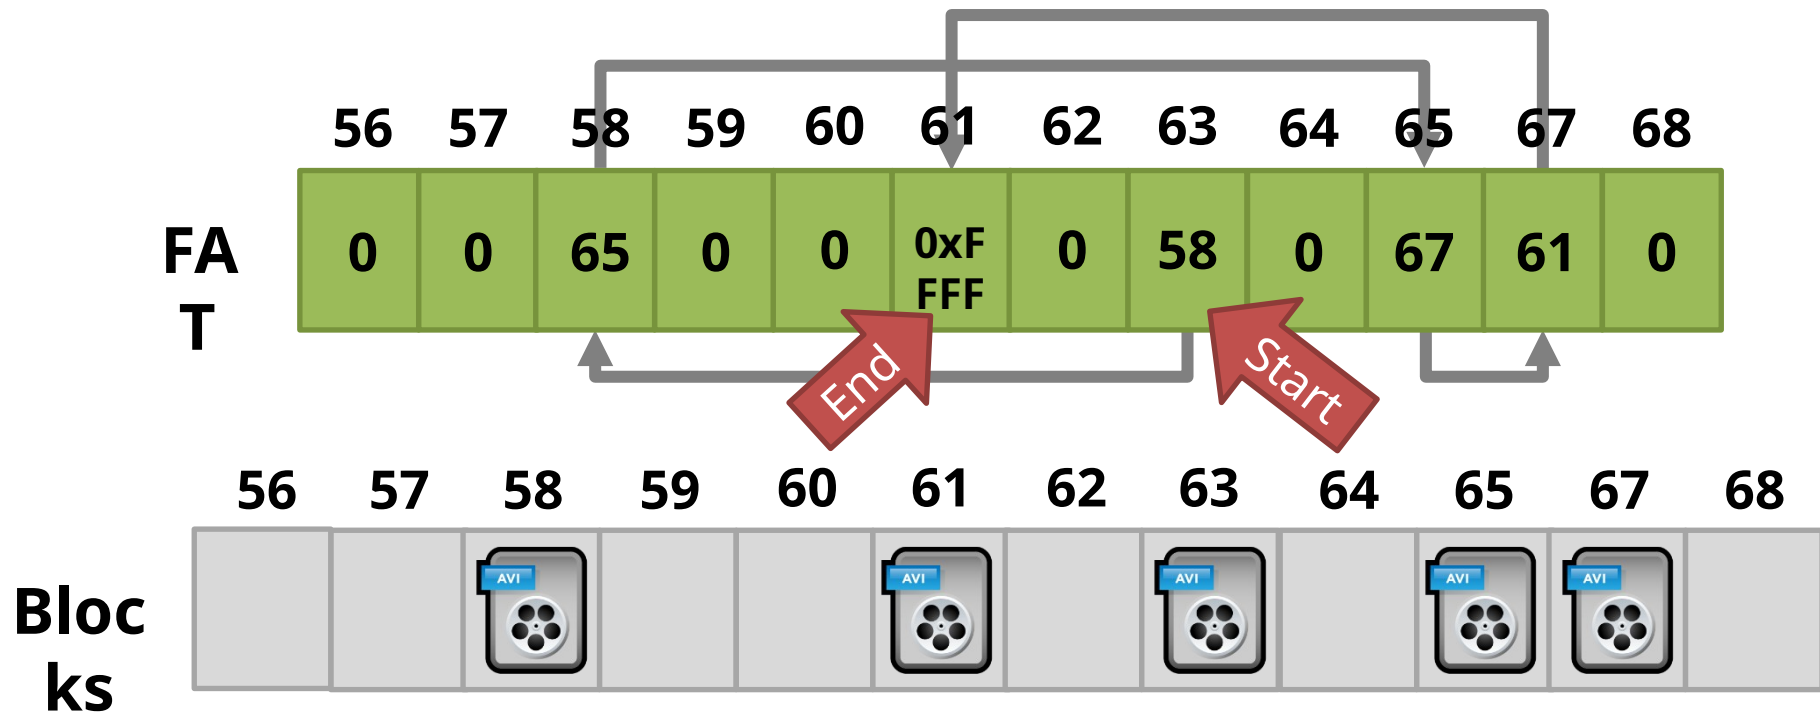| Name	Index	Dir?	Perms
	2	Y	
Windows	3	Y	
Users	4	Y	rwx
pagefile.s	5	N	r

Fat Table Entries

- len(FAT) == Number of clusters on the disk
  - Max number of files/directories is bounded
  - Decided when you format the partition
- The FAT version roughly corresponds to the size in bits of each FAT entry
  - E.g. FAT16 → each FAT entry is 16 bits
  - More bits → larger disks are supported

Fragmentation

- Blocks for a file need not be contiguous

	56	57	58	59	60	61	62	63	64	65	67	68
FAT	0	0	65	0	0	0xFFFF	0	58	0	67	61	0

	56	57	58	59	60	61	62	63	64	65	67	68
Blocks			AVI			AVI		AVI		AVI	AVI	

Possible values for FAT entries:
- 0 – entry is empty
- 1 < N < 0xFFFF – next block in a chain
- 0xFFFF – end of a chain
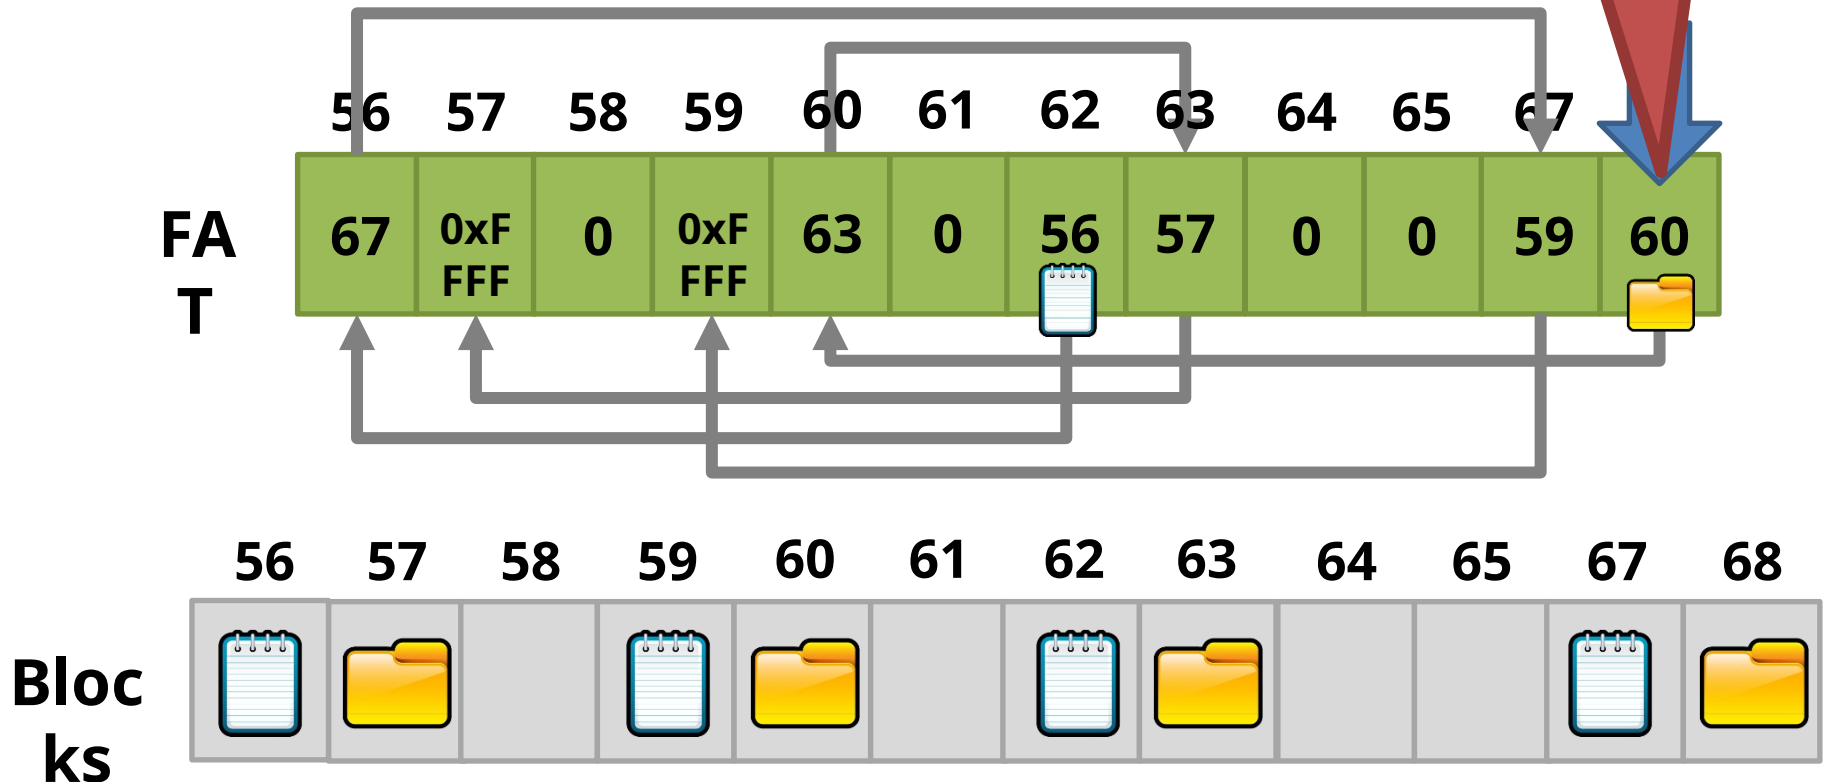
FAT: The Good and the Bad

- The Good – FAT supports:
  - Hierarchical tree of directories and files
  - Variable length files
  - Basic file and directory meta-data
- The Bad
  - At most, FAT32 supports 2TB disks
  - Locating free chunks requires scanning the entire FAT
  - Prone to internal and external fragmentation
    - Large blocks → internal fragmentation
  - **Reads require a lot of random seeking**

Lots of Seeki

- Consider the following cod

int fd = open("my_file.txt", "r");

int r = read(fd, buffer, 1024 * 4 * 4); // 4 4K blocks

FAT may have very low spatial locality, thus a lot of random seeking

FAT

56	57	58	59	60	61	62	63	64	65	67	
67	0xFFFF	0	0xFFFF	63	0	56	57	0	0	59	60

Blocks

56	57	58	59	60	61	62	63	64	65	67	68

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- Extents and B-Trees (ext4)
- Log-based File Systems

Status Check

- At this point, we have on-disk structures for:
  - Building a directory tree
  - Storing variable length files
- But, the efficiency of FAT is very low
  - Lots of seeking over file chains in FAT
  - Only way to identify free space is to scan over the entire FAT
- Linux file system uses more efficient structures
  - Extended File System (ext) uses index nodes (inodes) to track files and directories

Size Distribution of Files

- FAT uses a linked list for all files
  - Simple and uniform mechanism
  - ... but, it is not optimized for short or long files
- Question: are short or long files more common?
  - Studies over the last 30 years show that short files are much more common
  - 2KB is the most common file size
  - Average file size is 200KB (biased upward by a few very large files)
- Key idea: optimize the file system for many small files

33

- Super block, storing:
  - Size and location of bitmaps
  - Number and location of inodes
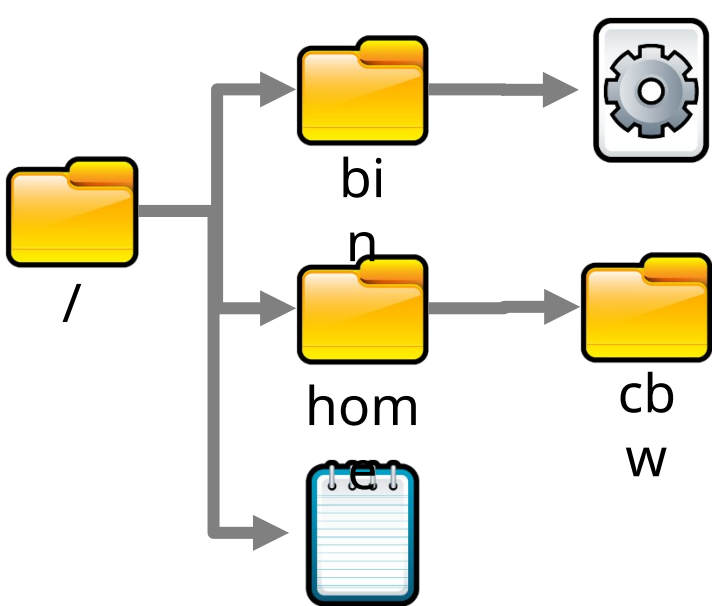 - Number and location of data blocks
  - Index of root inodes

- Table of inodes
- Each inode is a file/directory
- Includes meta-data and lists of associated data blocks

Bitmap of free & used data blocks

Bitmap of free & used inodes

Data blocks (4KB each)

- Directories are files
- Contains the list of entries in the

- Each inode can directly point to 12 blocks
- Can also indirectly point to blocks at 1, 2, and 3 levels of depth

bi n

hom e

cb w

initrd.img   3

Inode Bitma

Data Bitma

Inod es

Data locks

SB

Root inode = 0

35

ext2 inodes

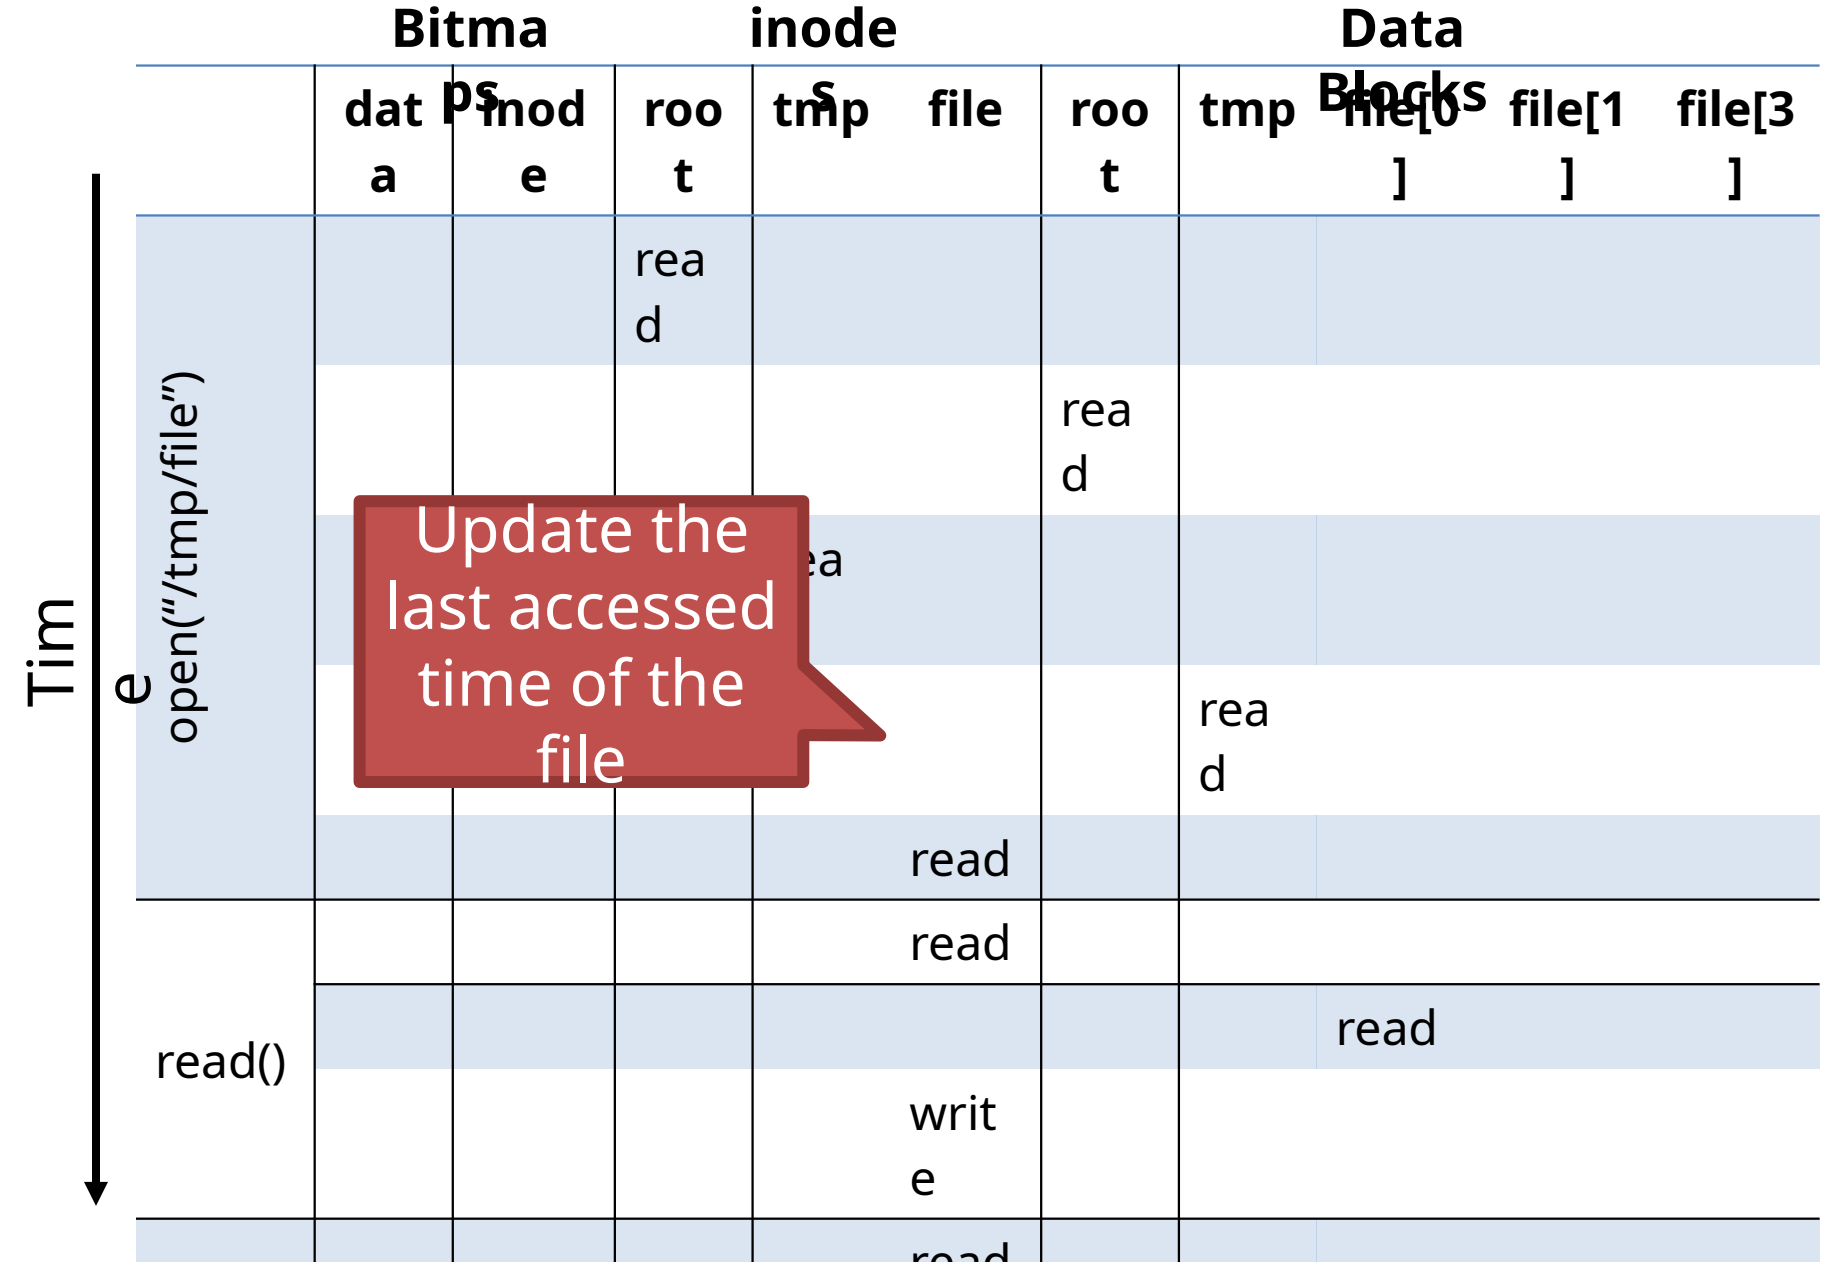| Size (bytes)	Name	What is this field for?
2	mode	Read/write/execute?
2	uid	User ID of the file owner
4	size	Size of the file in bytes
4	time	Last access time
4	ctime	Creation time
4	mtime	Last modification time
4	dtime	Deletion time
2	gid	Group ID of the file
2	links_count	How many hard links point to this file?
4	blocks	How many data blocks are allocated to this file?

inode Block Pointers

- Each inode is the root of an unbalanced tree of data blocks

inode

15 total pointers

Single Indirect

Double Indirect

Triple Indirect

12 blocks * 4KB = 48KB

1024 blocks * 4KB = 4MB

1024 * 1024 blocks * 4KB = 4GB
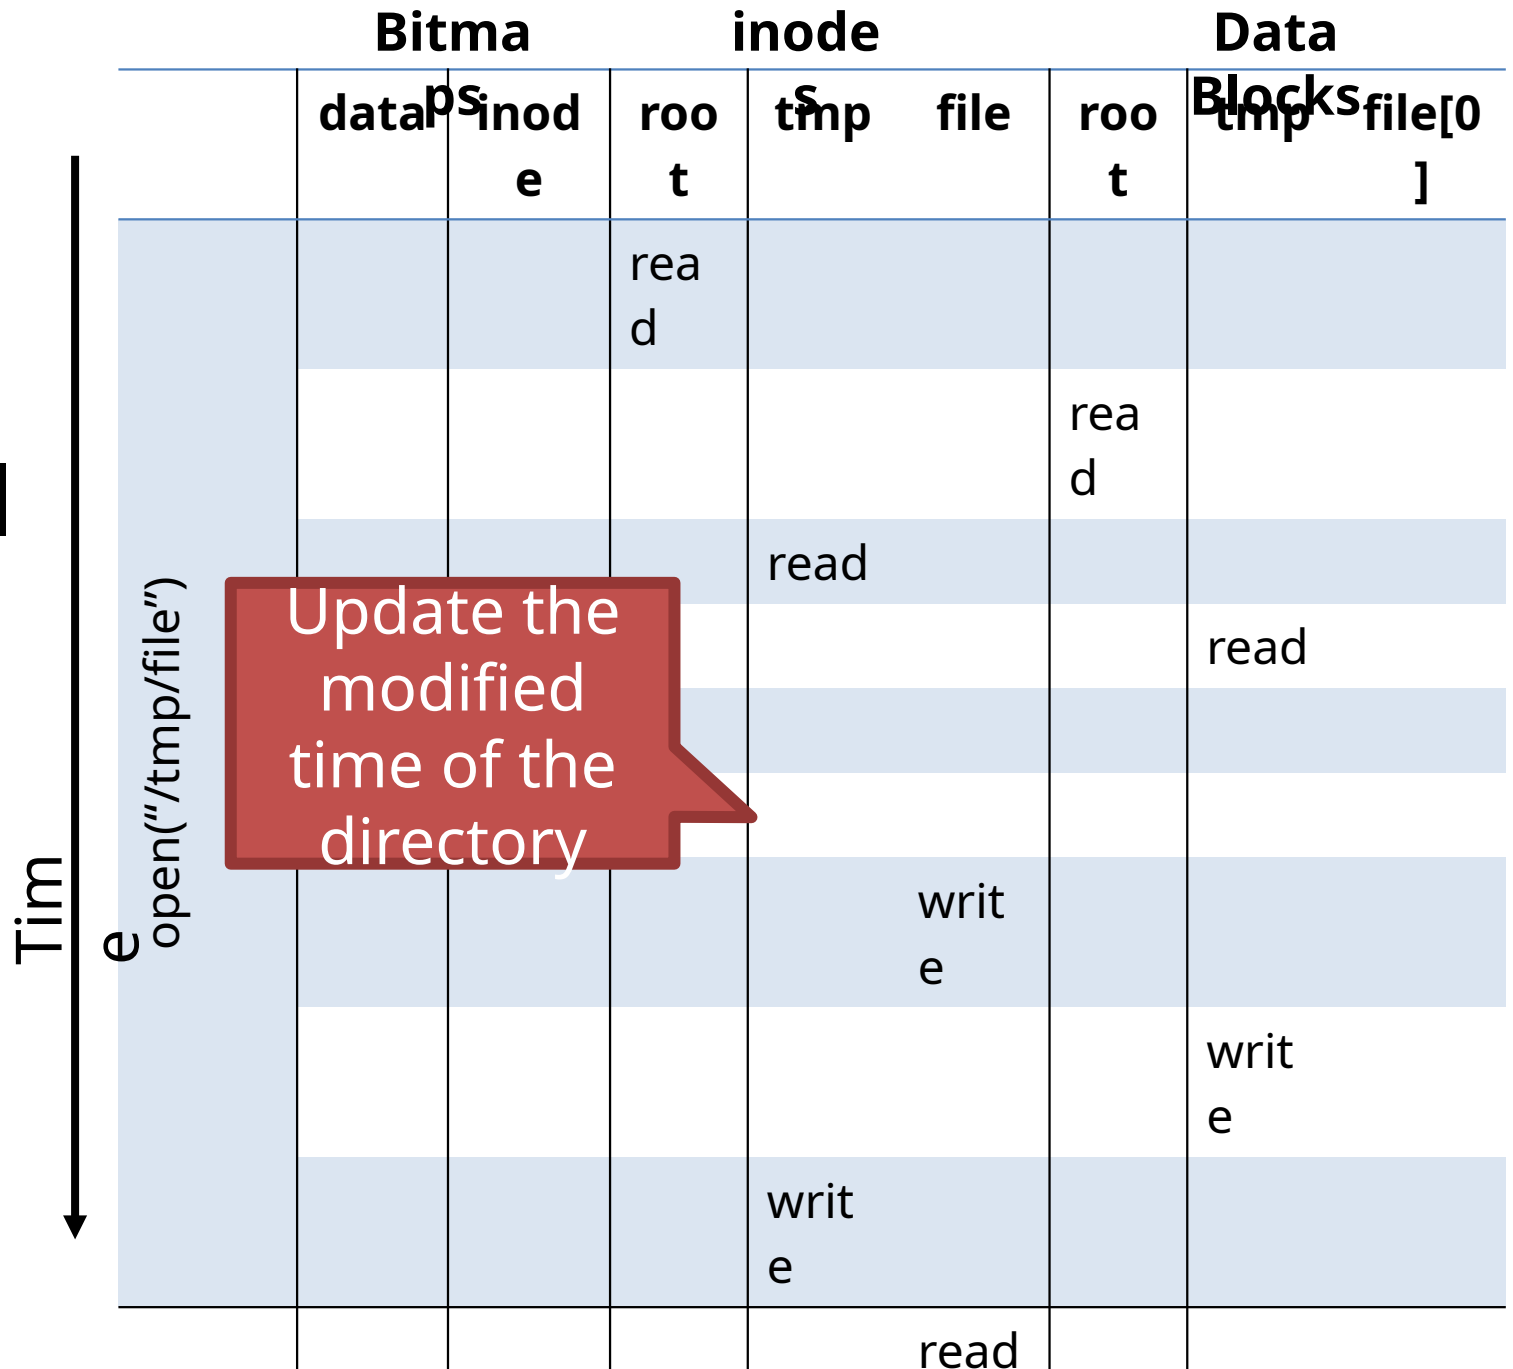
2^{30} blocks * 4KB = 4TB

Advantages of inodes

- Optimized for file systems with many small files
  - Each inode can directly point to 48KB of data
  - Only one layer of indirection needed for 4MB files
- Faster file access
  - Greater meta-data locality → less random seeking
  - No need to traverse long, chained FAT entries
- Easier free space management
  - Bitmaps can be cached in memory for fast

File Reading Example

|	Bitmaps		inodes			Data Blocks				
	data	inode	root	tmp	file	root	tmp	file[0]	file[1]	file[3]
open("/tmp/file")			read							
						read				
				read						
							read			
					read					
					read					
read()								read		
					write					
					read					

Update the last accessed time of the file

File Create and Write Example

	Bitmaps		inodes			Data Blocks		
	data	inode	root	tmp	file	root	tmp	file[0]
			read					
						read		
				read				
							read	
					write			
								write
				write				
				read				

open("/tmp/file")

Update the modified time of the directory

Time

ext2 inodes, Again

Size (bytes)	Name	What is this field for?
2	mode	Read/write/execute?
2	uid	User ID of the file owner
4	size	Size of the file in bytes
4	time	Last access time
4	ctime	Creation time
4	mtime	Last modification time
4	dtime	Deletion time
2	gid	Group ID of the file
2	links_count	How many hard links point to this file?
4	blocks	How many data blocks are allocated to this file?

Hard Link Example

- Multiple directory entries may point to the same inode

[amislove@ativ9 ~] ln –T ../cbw/my_file cbw_file

cb w

my_fil e

hom e

amislov e

cbw_fil e

1. Add an entry to the "amislove" directory
2. Increase the link_count of the "my_file" inode

Inode Bitma

Data Bitma

Inod es

Data Blocks

SB

Hard Link Details

- Hard links give you the ability to create many aliases of the same underlying file
  - Can be in different directories
- Target file will not be marked invalid (deleted) until link_count == 0
  - This is why POSIX "delete" is called *unlink()*
- Disadvantage of hard links
  - Inodes are only unique within a single file system
 - Thus, can only point to files in the same partition

43

Soft Links

- Soft links are special files that include the path to another file
  - Also known as symbolic links
  - On Windows, known as shortcuts
 - File may be on another partition or device

Soft Link Example

1. Create a soft link file
2. Add it to the current directory

home

cbw

my_file

amislove

cbw_file

Inode Bitma

Data Bitma

Inodes

Data Blocks

SB
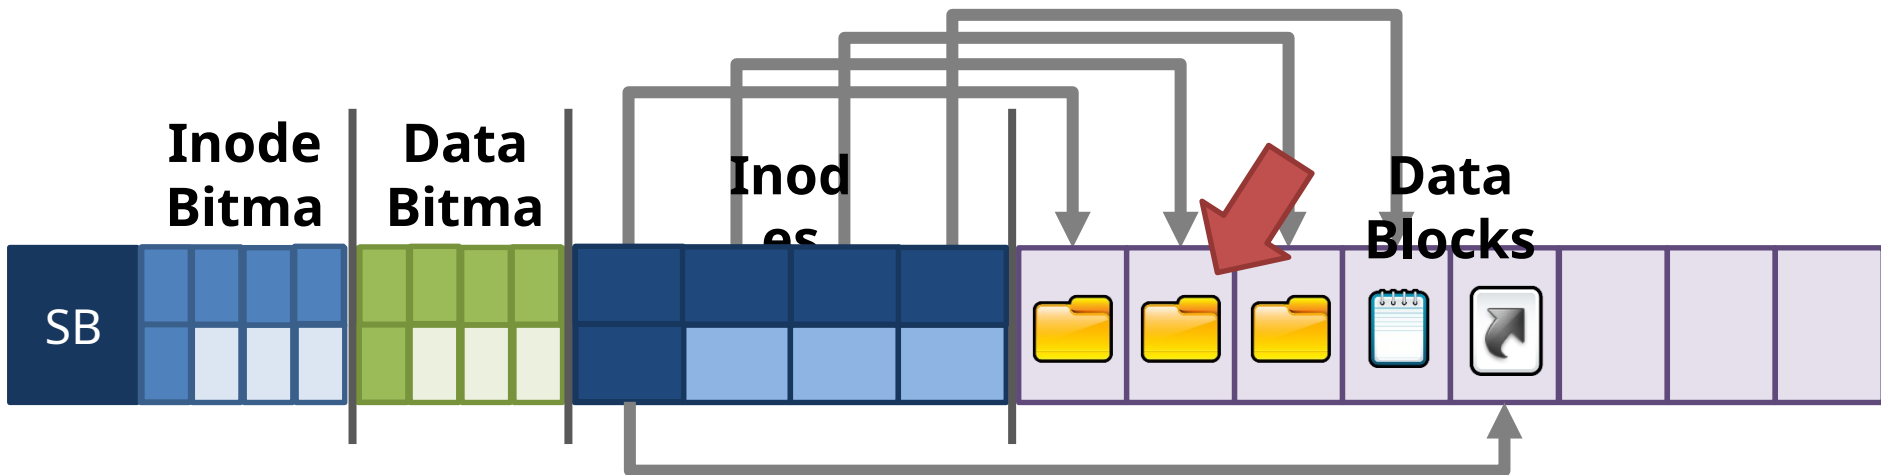
ext: The Good and the Bad

- The Good – ext file system (inodes) support:
  - All the typical file/directory features
  - Hard and soft links
  - More performant (less seeking) than FAT
- The Bad: poor locality
  - ext is optimized for a particular file size distribution
  - However, it is not optimized for spinning disks
  - inodes and associated data are far apart on the disk:

| SB | Inode Bitmap | Data Bitmap | Inodes | Data Blocks |

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- Extents and B-Trees (ext4)
- Log-based File Systems

Status Check

- At this point, we've moved from FAT to ext
  - inodes are imbalanced trees of data blocks
  - Optimized for the common case: small files
- Problem: ext has poor locality
  - inodes are far from their corresponding data
  - This is going to result in long seeks across the disk
- Problem: ext is prone to fragmentation
  - ext chooses the first available blocks for new data
  - No attempt is made to keep the blocks of a file contiguous

Fast File System (FFS)

- FFS developed at Berkeley in 1984
  - First attempt at a disk aware file system
  - i.e. optimized for performance on spinning disks
- Observation: processes tend to access files that are in the same (or close) directories
  - Spatial locality
- Key idea: place groups of directories and their files into cylinder groups
 - Introduced into ext2, called block groups

Block Groups

- In ext, there is a single set of key data structures
  - One data bitmap, one inode bitmap
  - One inode table, one array of data blocks
- In ext2, each block group contains its own key data structures

Inode Bitmap	Data Bitmap	Inodes	Data Blocks

SB	Block Group 1	Block Group 2	Block Group 3	Block Group 4	Block Group 5	Block Group 6

Allocation Policy

- ext2 attempts to keep related files and directories within the same block group

ext2: The Good and the Bad

- The good – ext2 supports:
  - All the features of ext...
  - ... with even better performance (because of increased spatial locality)
- The bad
  - Large files must cross block groups
  - As the file system becomes more complex, the chance of file system corruption grows
    - E.g. invalid inodes, incorrect directory entries, etc.

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- Extents and B-Trees (ext4)
- Log-based File Systems

Status Check

- At this point, we have a full featured file system
  - Directories
  - Fine-grained data allocation
  - Hard/soft links
- File system is optimized for spinning disks
  - inodes are optimized for small files
 - Block groups improve locality
- What's next?

Maintaining Consistency

- Many operations results in multiple, independent writes to the file system
  - Example: append a block to an existing file
  1. Update the free data bitmap
  2. Update the inode
  3. Write the user data
- What happens if the computer crashes in the middle of this process?

File Append Example

owner: christo
permissions: rw
size: 2
pointer: 4
pointer: 5
pointer: null
pointer: null

- These three operations can potentially be done in any order
- ... but the system can crash at any time

Inode Bitmap

Data Bitmap

Inodes

v2

Data Blocks

D_1 D_2

Update the data bitmap

Update the inode

Write the data

| Inode Bitmap	Data Bitmap	Inodes		Data Blocks	

Write the data

Result: file system is consistent, but the data is lost

Update the inode

Result: inode points to garbage data, and file system is inconsistent (data bitmap vs. inode)

Update the data bitmap

Result: space leakage, and file system is inconsistent (data bitmap vs. inode)

Inode Bitmap | **Data Bitmap** | **Inodes** | **Data Blocks**

v2 ... D_1 D_2

Result: inode points to data, but file system is inconsistent

v1 ... D_1 D_2

Result: file system is inconsistent, and the data is useless since it's not associated with an inode

v2 ... D_1

Result: file system is consistent, but the inode points to garbage data

The Crash Consistency Problem

- The disk guarantees that sector writes are atomic
  - No way to make multi-sector writes atomic
- How to ensure consistency after a crash?
  1. Don't bother to ensure consistency
     - Accept that the file system may be inconsistent after a crash
     - Run a program that fixes the file system during bootup
     - File system checker (*fsck*)
  2. Use a transaction log to make multi-writes atomic

Approach 1: File System Checker

- Key idea: fix inconsistent file systems during bootup
  - Unix utility called *fsck* (*chkdsk* on Windows)
  - Scans the entire file system multiple times, identifying and correcting inconsistencies
- Why during bootup?
  - No other file system activity can be going on
  - After fsck runs, bootup/mounting can continue

fsck Tasks

- **Superblock:** validate the superblock, replace it with a backup if it is corrupted
- **Free blocks and inodes:** rebuild the bitmaps by scanning all inodes
- **Reachability:** make sure all inodes are reachable from the root of the file system
- **inodes:** delete all corrupted inodes, and rebuild their link counts by walking the directory tree
- **directories:** verify the integrity of all directories

fsck: the Good and the Bad

- Advantages of *fsck*
  - Doesn't require the file system to do any work to ensure consistency
  - Makes the file system implementation simpler
- Disadvantages of *fsck*
  - Very complicated to implement the *fsck* program
    - Many possible inconsistencies that must be identified
    - Many difficult corner cases to consider and handle
  - *fsck* is **super slow**
    - Scans the entire file system multiple times

Approach 2: Journaling

- Problem: *fsck* is slow because it checks the entire file system after a crash
  - What if we knew where the last writes were before the crash, and just checked those?
- Key idea: make writes transactional by using a write-ahead log
  - Commonly referred to as a journal
- Ext3 and NTFS use journaling

Superblock	Journal	Block Group 1	...	Block Group *N*	

Write-Ahead Log

- Key idea: writes to disk are first written into a log
  - After the log is written, the writes execute normally
  - In essence, the log records transactions
- What happens after a crash...
  - If the writes to the log are interrupted?
    - The transaction is incomplete
    - The user's data is lost, but the file system is consistent
  - If the writes to the log succeed, but the normal writes are interrupted?
    - The file system may be inconsistent, but...
    - The log has exactly the right information to fix the problem

Data Journaling Example

- Assume we are appending to a file
  - Three writes: inode v2, data bitmap v2, data D_2
- Before executing these writes, first log them

Journal	TxB ID=1	I v2	B v2	D_2	TxE ID=1	

1. Begin a new transaction with a unique ID=k
2. Write the updated meta-data block(s)
3. Write the file data block(s)
4. Write an end-of-transaction with ID=k

Commits and Checkpoints

- We say a transaction is committed after all writes to the log are complete

- After a transaction is committed, the OS checkpoints the update

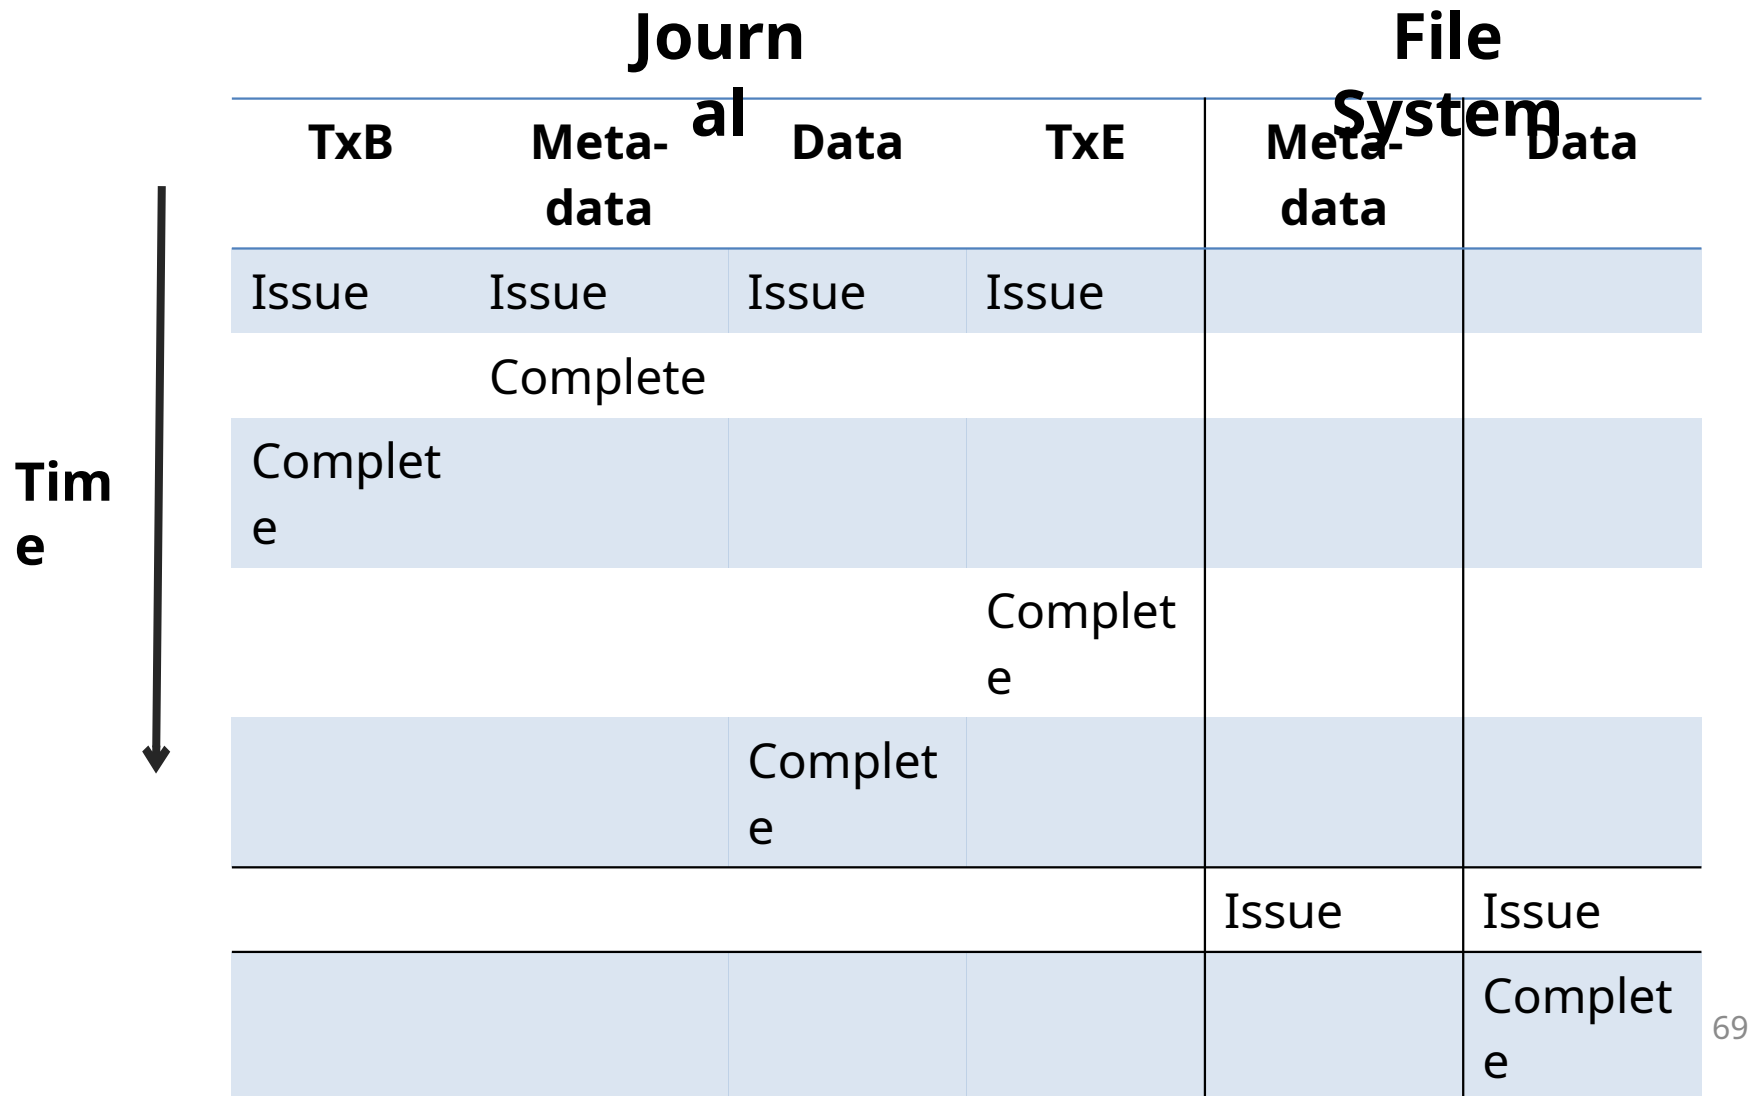Committed!

Checkpointed!

Journal

| TxB | I v2 | B v2 | D_2 | TxE | |

Inode Bitmap | **Data Bitmap** | **Inodes** | | **Data Blocks**

p | p | v2 | | D_1 D_2

- Final step: free the checkpointed transaction

67

Journal Implementation

- Journals are typically implemented as a circular buffer
  - Journal is **append-only**
- OS maintains pointers to the front and back of the transactions in the buffer
  - As transactions are freed, the back is moved up
- Thus, the contents of the journal are never deleted, they are just overwritten over time

Data Journaling Timeline

Time (downward arrow)

|	Journal			File System	
TxB	**Meta-data**	**Data**	**TxE**	**Meta-data**	**Data**
Issue	Issue	Issue	Issue		
	Complete				
Complete					
			Complete		
		Complete			
				Issue	Issue
					Complete

Crash Recovery (1)

- What if the system crashes during logging?
  - If the transaction is not committed, data is lost
  - But, the file system remains consistent

Journal

TxB	I v2	B v2	D$_2$	

Inode Bitmap | **Data Bitmap** | **Inodes** | **Data Blocks**

v1

D$_1$

Crash Recovery (2)

- What if the system crashes during the checkpoint?
  - File system may be inconsistent
  - During reboot, transactions that are committed but not free are replayed in order
  - Thus, no data is lost and consistency is restored

Journal

TxB	I v2	B v2	D_2	TxE

Inode Bitmap

Data Bitmap

Inodes

v2

Data Blocks

D_1 D_2

Corrupted Transactions

- Problem: the disk scheduler may not execute writes in-order
  - Transactions in the log may appear committed, when in fact they are invalid

Journal

TxB	I v2	B v2	D$_2$	TxE	

- Solution: add a checksum to TxB
- During recovery, reject transactions with invalid checksums
- Implemented on Linux in ext4

- Transaction looks valid, but the data is missing!
- During replay, garbage data is written to the file
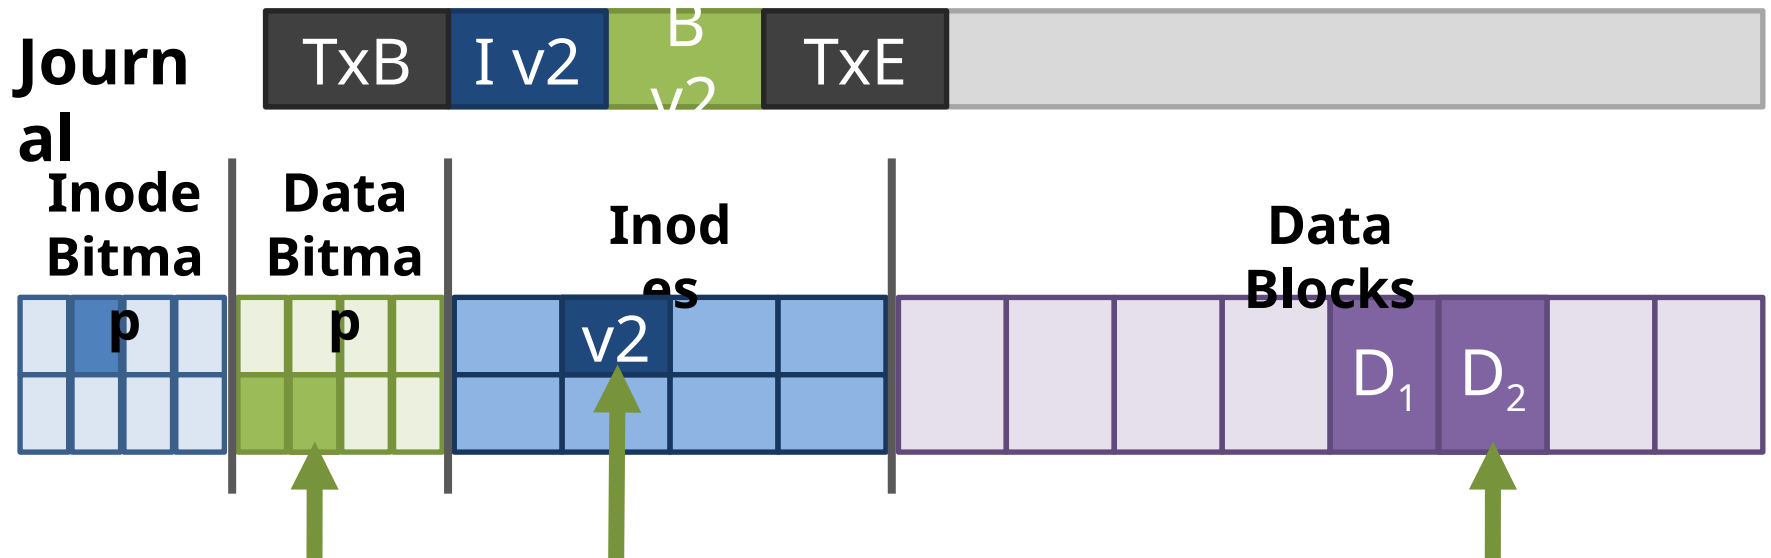
Journaling: The Good and the Bad

- Advantages of journaling
  - Robust, fast file system recovery
    - No need to scan the entire journal or file system
  - Relatively straight forward to implement
- Disadvantages of journaling
  - Write traffic to the disk is doubled
    - Especially the file data, which is probably large
  - Deletes are very hard to correctly log
    - Example in a few slides…
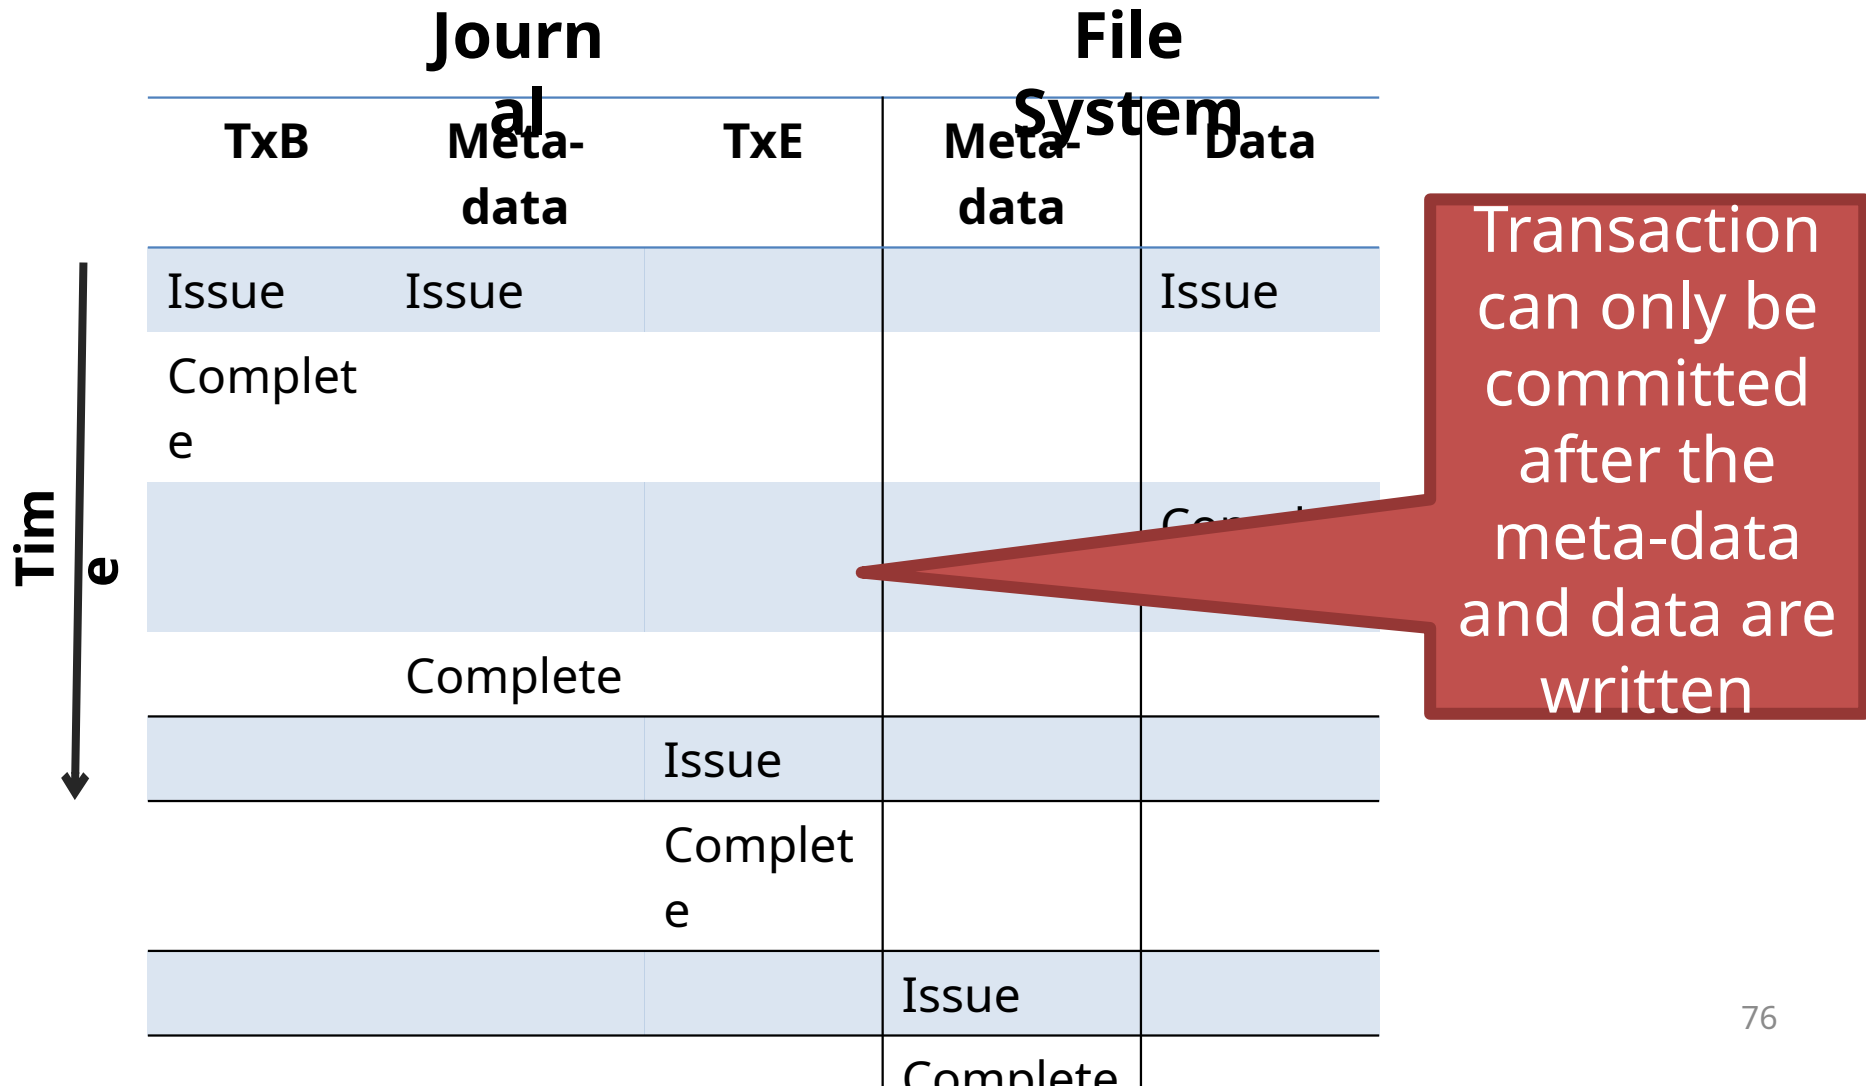
Making Journaling Faster

- Journaling adds a lot of write overhead
- OSes typically batch updates to the journal
  - Buffer sequential writes in memory, then issue one large write to the log
  - Example: ext3 batches updates for 5 seconds
- Tradeoff between performance and persistence
  - Long batch interval = fewer, larger writes to the log
    - Improved performance due to large sequential writes
  - But, if there is a crash, everything in the buffer

Meta-Data Journaling

- The most expensive part of data journaling is writing the file data twice
  - Meta-data is small (~1 sector), file data is large

- ext3 implements meta-data journaling

Meta-Journaling Timeline

Time

Journal			File System	
TxB	**Meta-data**	**TxE**	**Meta-data**	**Data**
Issue	Issue			Issue
Complete				
				Complete
	Complete			
		Issue		
		Complete		
			Issue	
			Complete	

> Transaction can only be committed after the meta-data and data are written
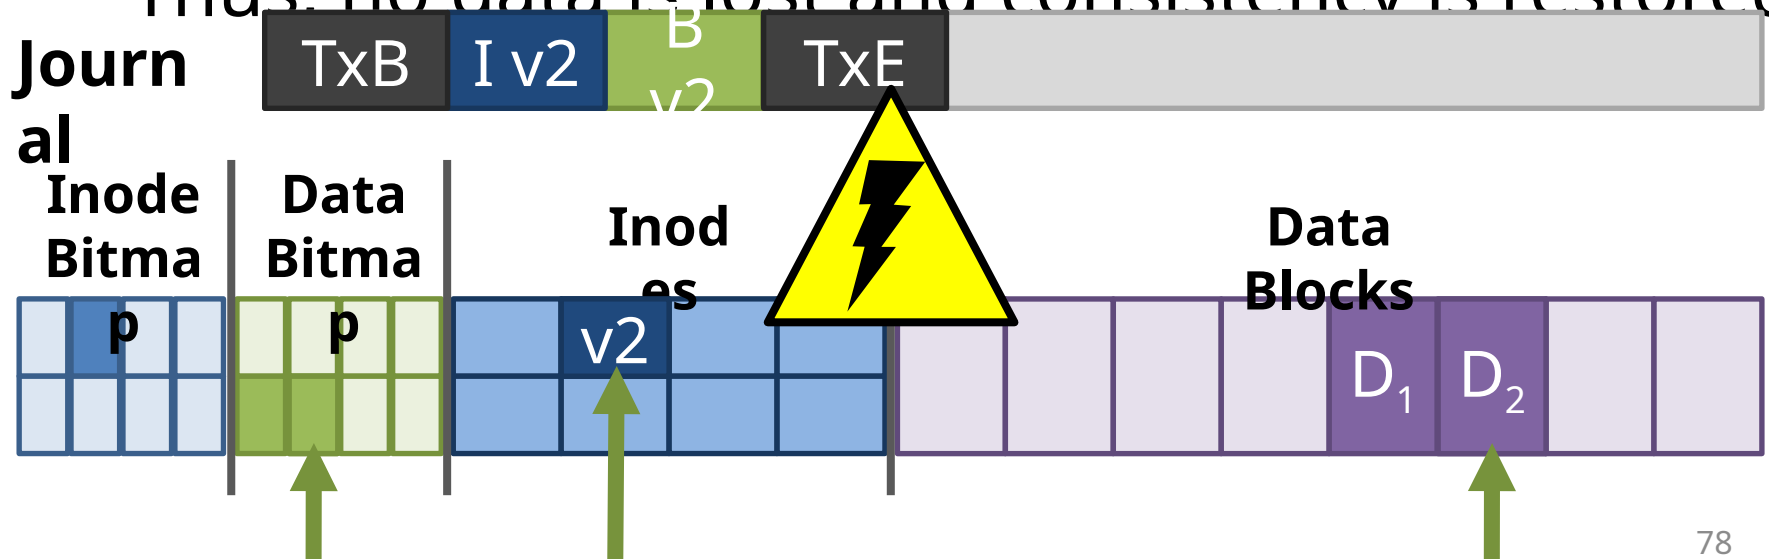
Crash Recovery Redux (1)

- What if the system crashes during logging?
  - If the transaction is not committed, data is lost
  - D_2 will eventually be overwritten
  - The file system remains consistent

Journal

TxB	I v2	B v2

Inode Bitmap

p

Data Bitmap

p

Inodes

v1

Data Blocks
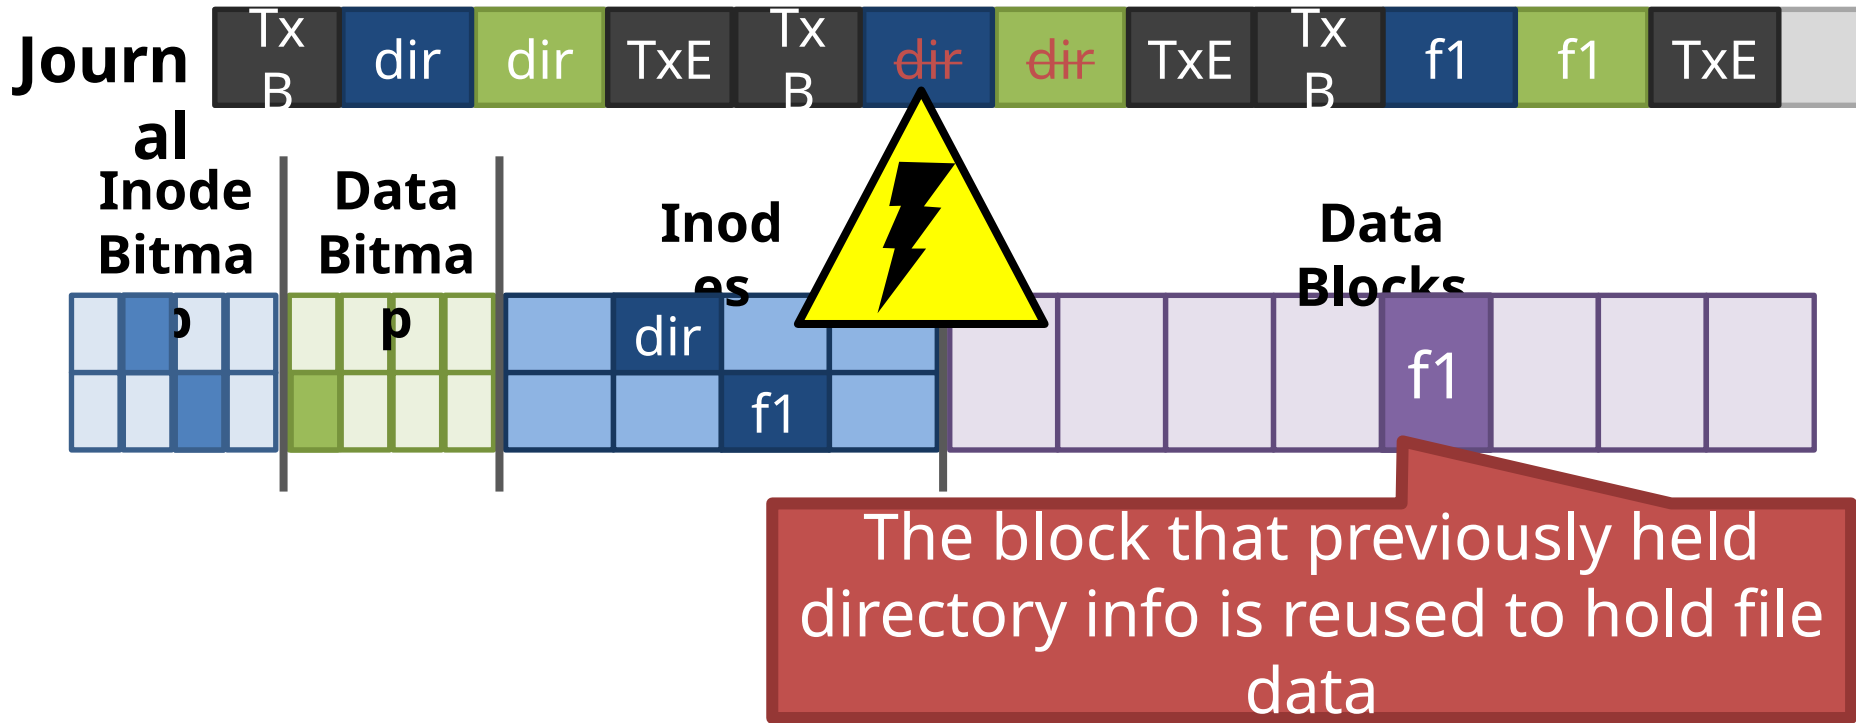
D_1	D_2

Crash Recovery Redux (2)

- What if the system crashes during the checkpoint?
  - File system may be inconsistent
  - During reboot, transactions that are committed but not free are replayed in order
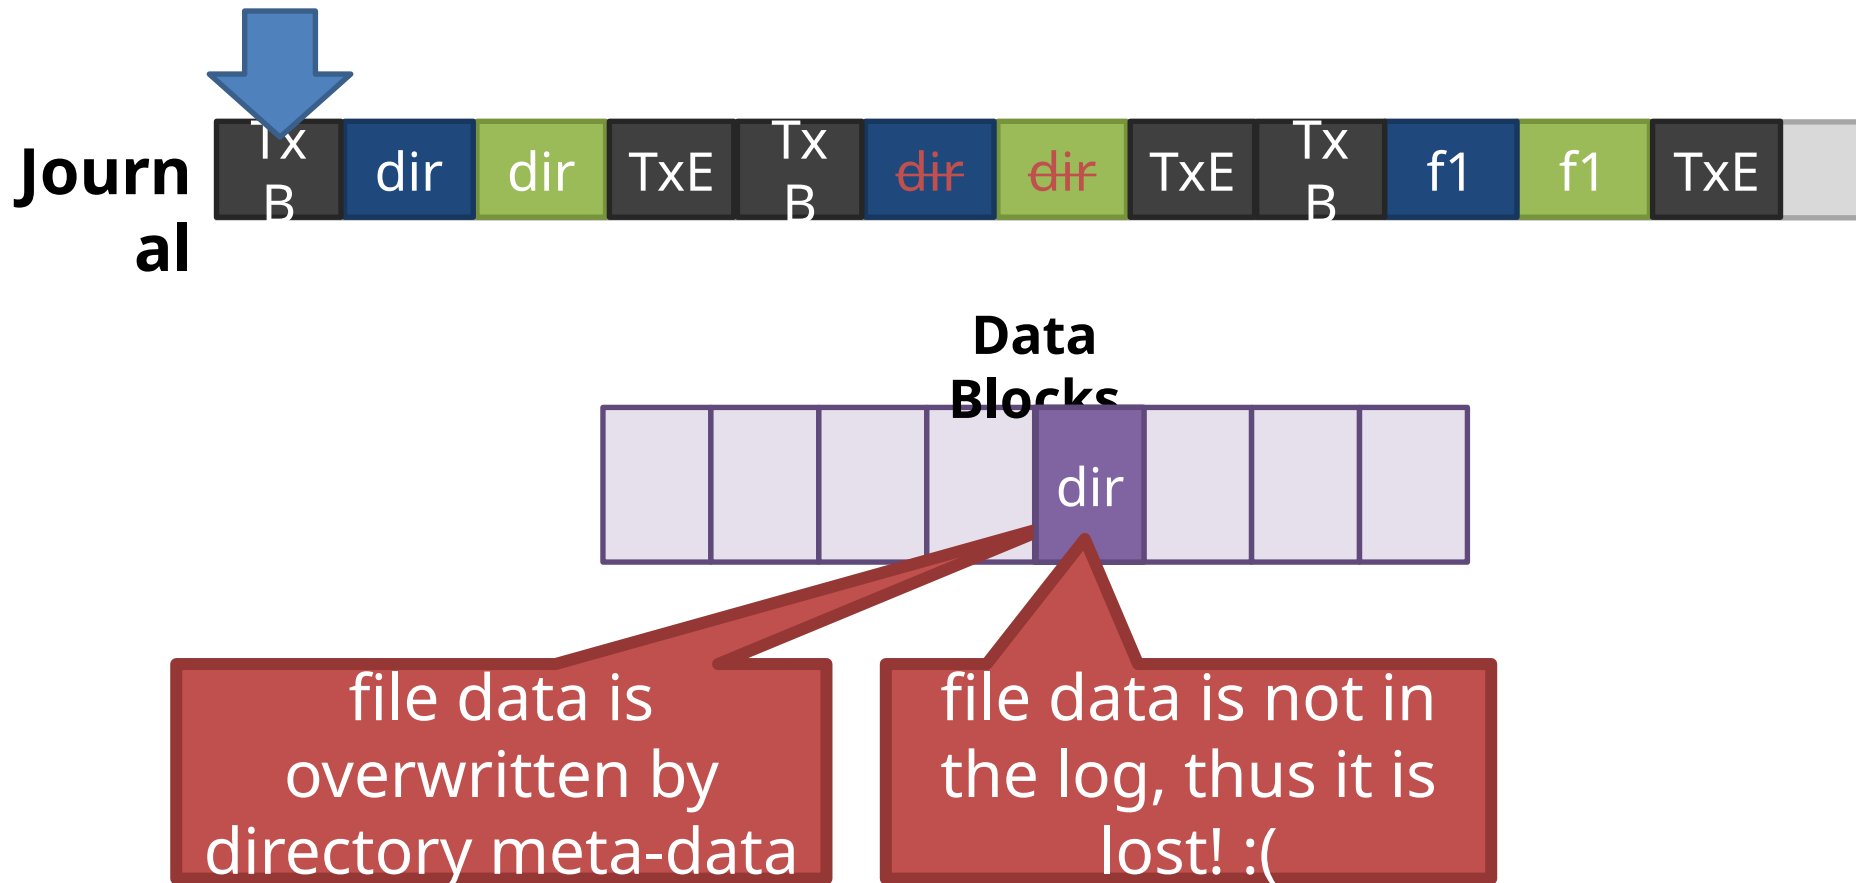 - Thus, no data is lost and consistency is restored

Journal

TxB	I v2	B v2	TxE	

Inode Bitmap

Data Bitmap

Inodes
v2

Data Blocks
D_1 D_2

Delete and Block Reuse

Journal

| Tx B | dir | dir | TxE | Tx B | ~~dir~~ | ~~dir~~ | TxE | Tx B | f1 | f1 | TxE | |

Inode Bitmap | **Data Bitmap** | **Inodes** | **Data Blocks**

dir

f1

f1

The block that previously held directory info is reused to hold file data

1. Create a directory: inode and data are written

2. Delete the directory: inode is removed
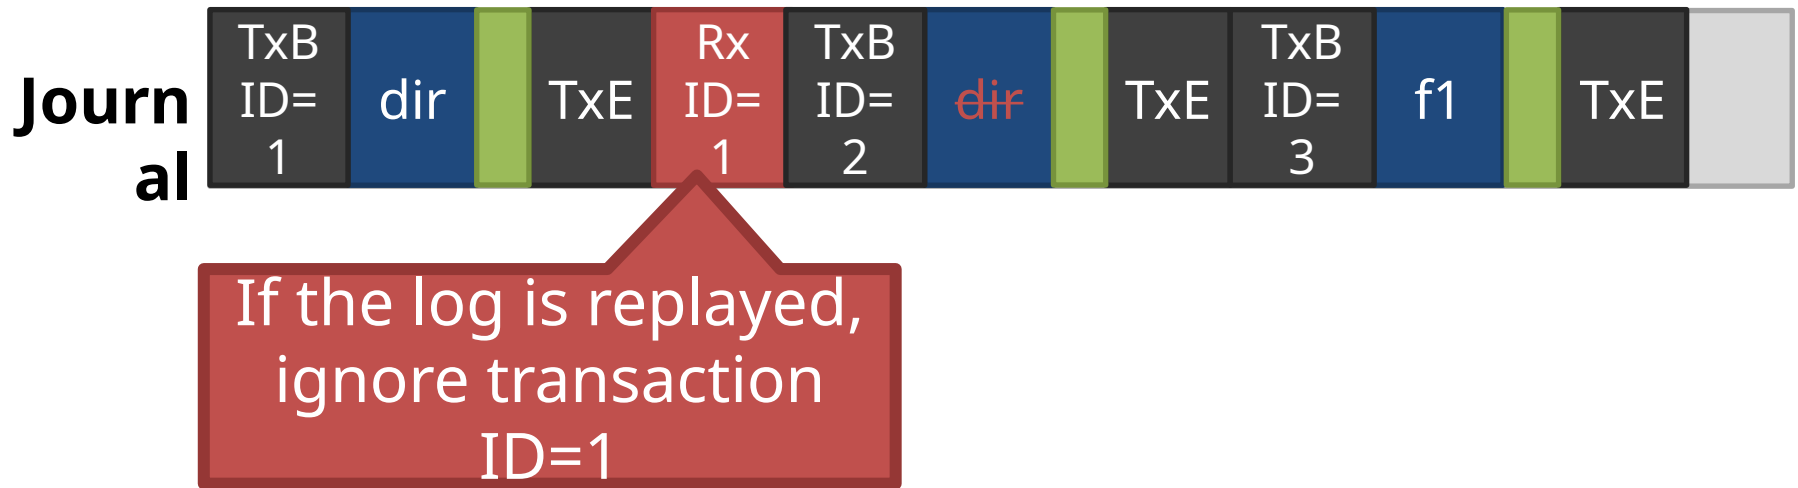
3. Create a file: inode and data are written

The Trouble With Delete

- What happens when the log is replayed?

Journal

| TxB | dir | dir | TxE | TxB | dir | dir | TxE | TxB | f1 | f1 | TxE | |

Data Blocks

| | | | | dir | | | |

file data is overwritten by directory meta-data

file data is not in the log, thus it is lost! :(

Handling Delete

- Strategy 1: don't reuse blocks until the delete is checkpointed and freed
- Strategy 2: add a revoke record to the log
  - ext3 used revoke records

Journal

| TxB ID=1 | dir | | TxE | Rx ID=1 | TxB ID=2 | ~~dir~~ | | TxE | TxB ID=3 | f1 | | TxE | |

If the log is replayed, ignore transaction ID=1

Journaling Wrap-Up

- Today, most OSes use journaling file systems
  - ext3/ext4 on Linux
  - NTFS on Windows
- Provides excellent crash recovery with relatively low space and performance overhead
- Next-gen OSes will likely move to file systems with copy-on-write semantics
  - btrfs and zfs on Linux

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- **Extents and B-Trees (ext4)**
- **Log-based File Systems**

Status Check

- At this point:
  - We not only have a fast file system
  - But it is also resilient against corruption

- What's next?
  - More efficiency improvements!

Revisiting inodes

- Recall: inodes use indirection to acquire additional blocks of pointers
- Problem: inodes are not efficient for large files
  - Example: for a 100MB file, you need 25600 block pointers (assuming 4KB blocks)
- This is unavoidable if the file is 100% fragmented
  - However, what if large groups of blocks are contiguous?

From Pointers to Extents

- Modern file systems try hard to minimize fragmentation
  - Since it results in many seeks, thus low performance
- Extents are better suited for contiguous files

inode
block 1
block 2
block 3
block 4
block 5
block 6

block 1
length
1
block 2
length
2
block 3
length

Each extent includes a block pointer and a length

Implementing Extents

- ext4 and NTFS use extents
- ext4 inodes include 4 extents instead of block pointers
  - Each extent can address at most 128MB of contiguous space (assuming 4KB blocks)
  - If more extents are needed, a data block is allocated
  - Similar to a block of indirect pointers

Revisiting Directories

- In ext, ext2, and ext3, each directory is a file with a list of entries
  - Entries are not stored in sorted order
  - Some entries may be blank, if they have been deleted
- Problem: searching for files in large directories takes O(n) time
  - Practically, you can't store >10K files in a directory
  - It takes way too long to locate and open files
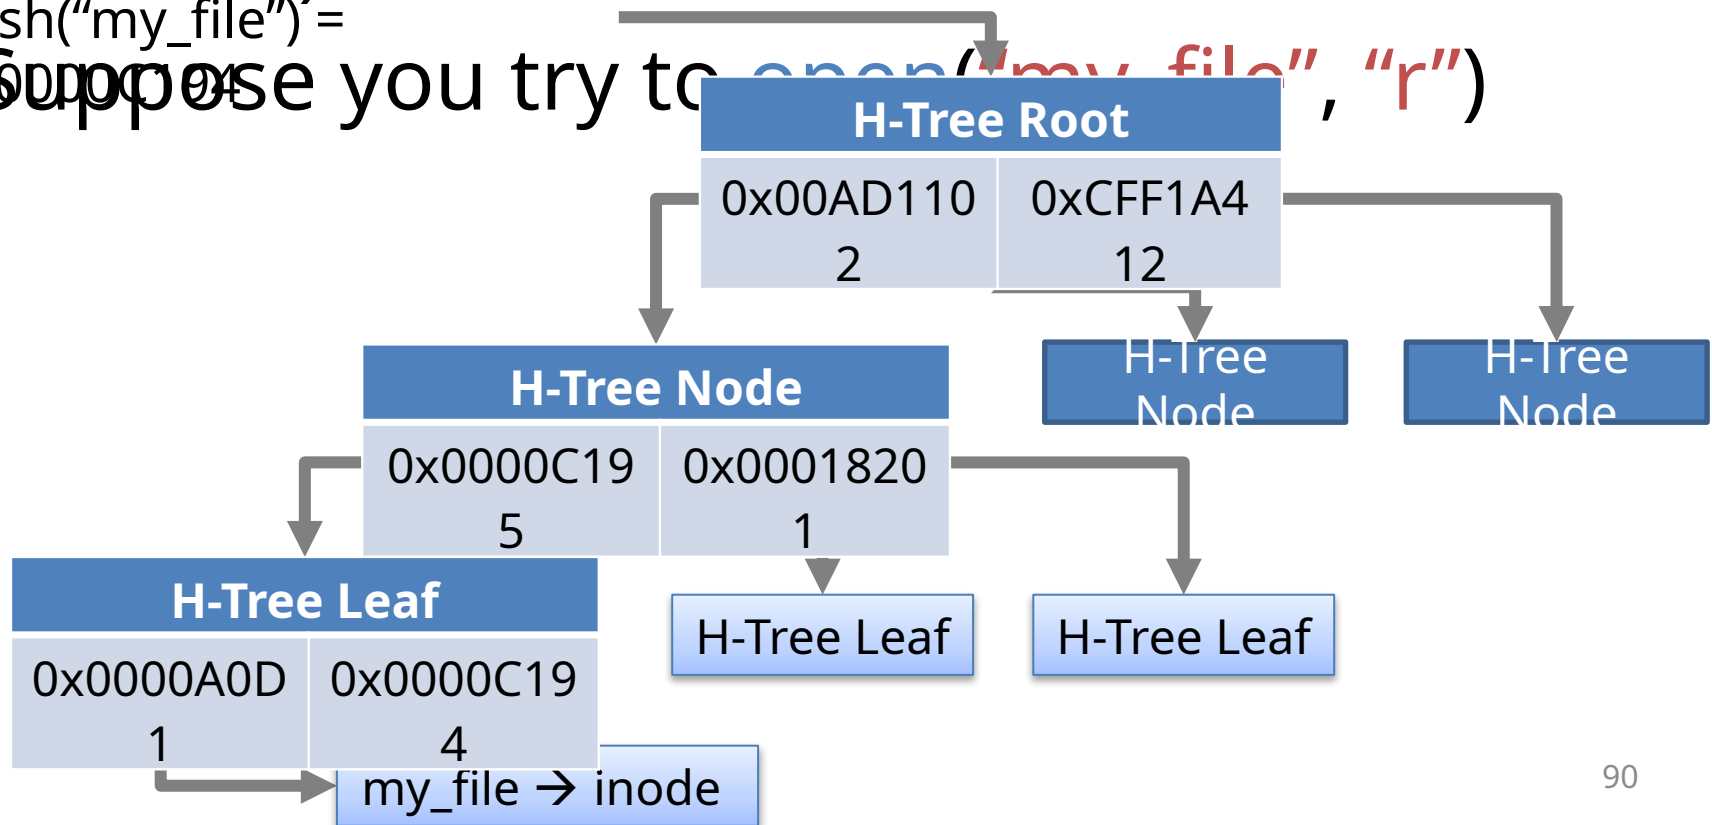
From Lists to B-Trees

- ext4 and NTFS encode directories as B-Trees to improve lookup time to O(log N)
- A B-Tree is a type of balanced tree that is optimized for storage on disk
  - Items are stored in sorted order in blocks
  - Each block stores between m and $2m$ items
- Suppose items i and j are in the root of the tree
  - The root must have 3 children, since it has 2 items
  - The three child groups contain items $a < i$, $i < a < j$, and $a > j$

Example B-Tree

- ext4 uses a B-Tree variant known as a H-Tree
  - The *H* stands for *hash* (sometime called B+Tree)

hash("my_file") = 0x0000A0D1

- Suppose you try to open("my_file", "r")

H-Tree Root	
0x00AD1102	0xCFF1A412

H-Tree Node

H-Tree Node

H-Tree Node	
0x0000C195	0x00018201

H-Tree Leaf

H-Tree Leaf

H-Tree Leaf	
0x0000A0D1	0x0000C194

my_file → inode

ext4: The Good and the Bad

- The good – ext4 (and NTFS) supports:
  - All of the basic file system functionality we require
  - Improved performance from ext3's block groups
  - Additional performance gains from extents and B-Tree directory files
- The bad:
  - ext4 is an incremental improvement over ext3
  - Next-gen file systems have even nicer features
    - Copy-on-write semantics (btrfs and ZFS)

- Partitions and Mounting
- Basics (FAT)
- inodes and Blocks (ext)
- Block Groups (ext2)
- Journaling (ext3)
- Extents and B-Trees (ext4)
- **Log-based File Systems**

Status Check

- At this point:
  - We have arrived at a modern file system like ext4
- What's next?
  - Go back to the drawing board and reevaluate from first-principals

Reevaluating Disk Performance

- How has computer hardware been evolving?
  - RAM has become cheaper and grown larger :)
  - Random access seek times have remained very slow :(
- This changing dynamic alters how disks are used
  - More data can be cached in RAM = less disk reads
  - Thus, writes will dominate disk I/O
- Can we create a file system that is optimized for sequential writes?

Log-structured File System

- Key idea: buffer all writes (including meta-data) in memory
  - Write these long segments to disk sequentially
  - Treat the disk as a circular buffer, i.e. don't overwrite
- Advantages:
  - All writes are large and sequential
- Big question:
  - How do you manage meta-data and maintain structure in this kind of design?
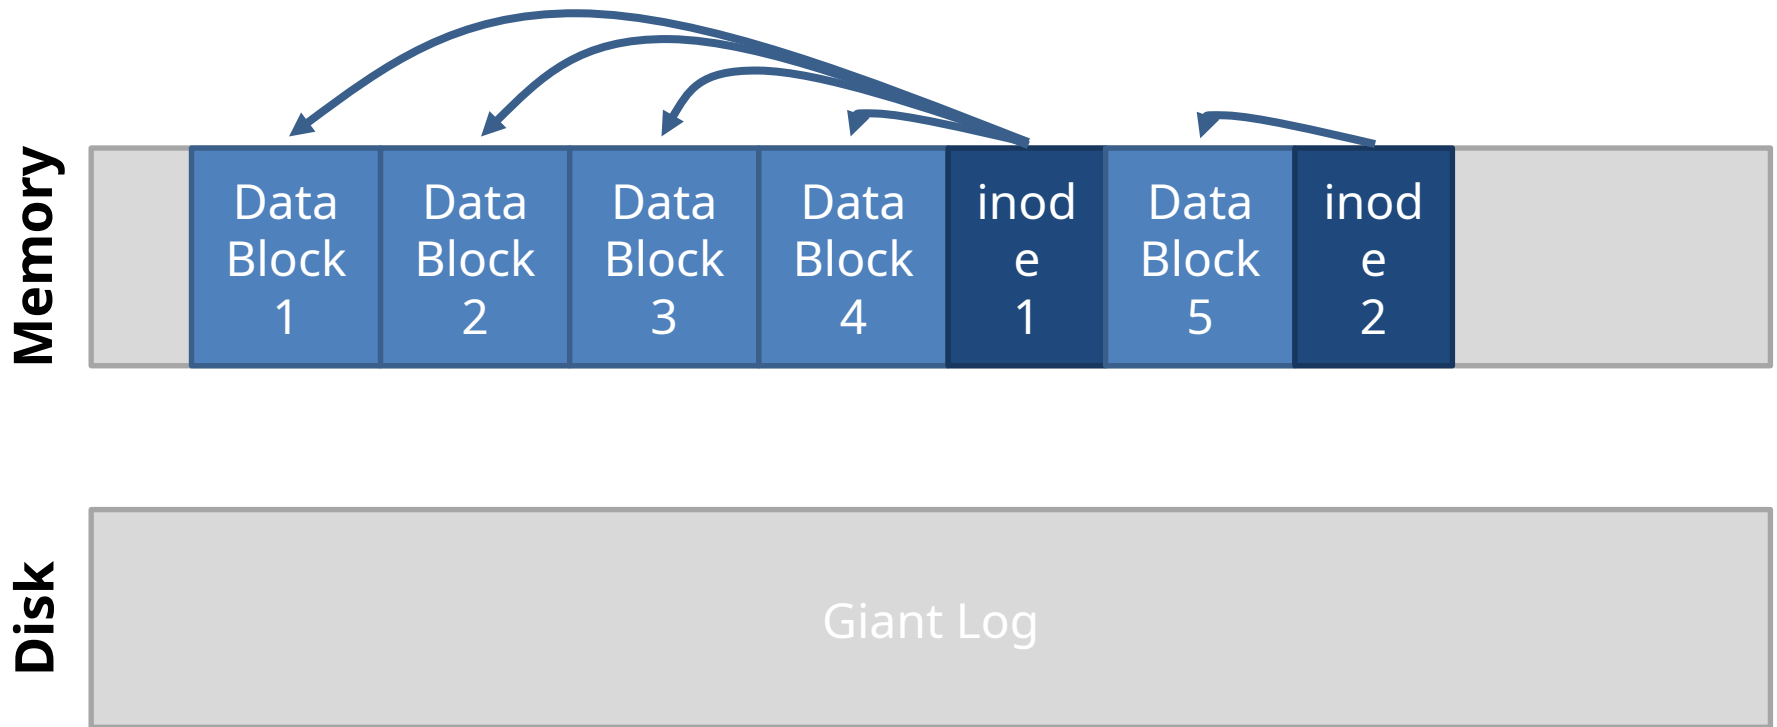
Treating the Disk as a Log

- Same concept as data journaling
  - Data and meta-data get appended to a log
  - Stale data isn't overwritten, its replaced

Disk

| Data Block 1 | Data Block 2 | inode 1 | Data Block 5 | inode 2 | Data Block 1 | inode 1 |

Buffering Writes

- LFS buffers writes in-memory into chunks

Memory

| Data Block 1 | Data Block 2 | Data Block 3 | Data Block 4 | inode 1 | Data Block 5 | inode 2 |

Disk

Giant Log

- Chunks get appended to the log once they are sufficiently large

How to Find inodes

- In a typical file system, the inodes are stored at fixed locations (relatively easy to find)

- How do you find inodes in the log?
  – Remember, there may be multiple copies of a given inode

- Solution: add a level of indirection
  – The traditional inode map can be broken into pieces
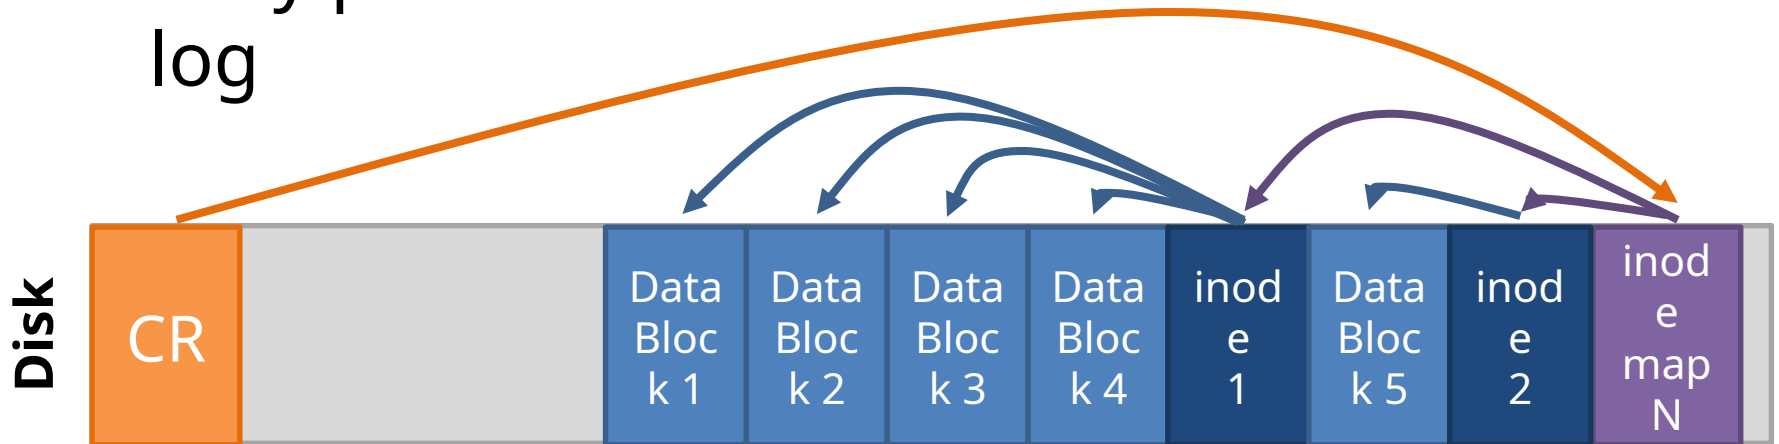 – When a portion of the inode map is updated, write it to the log!

inode Maps
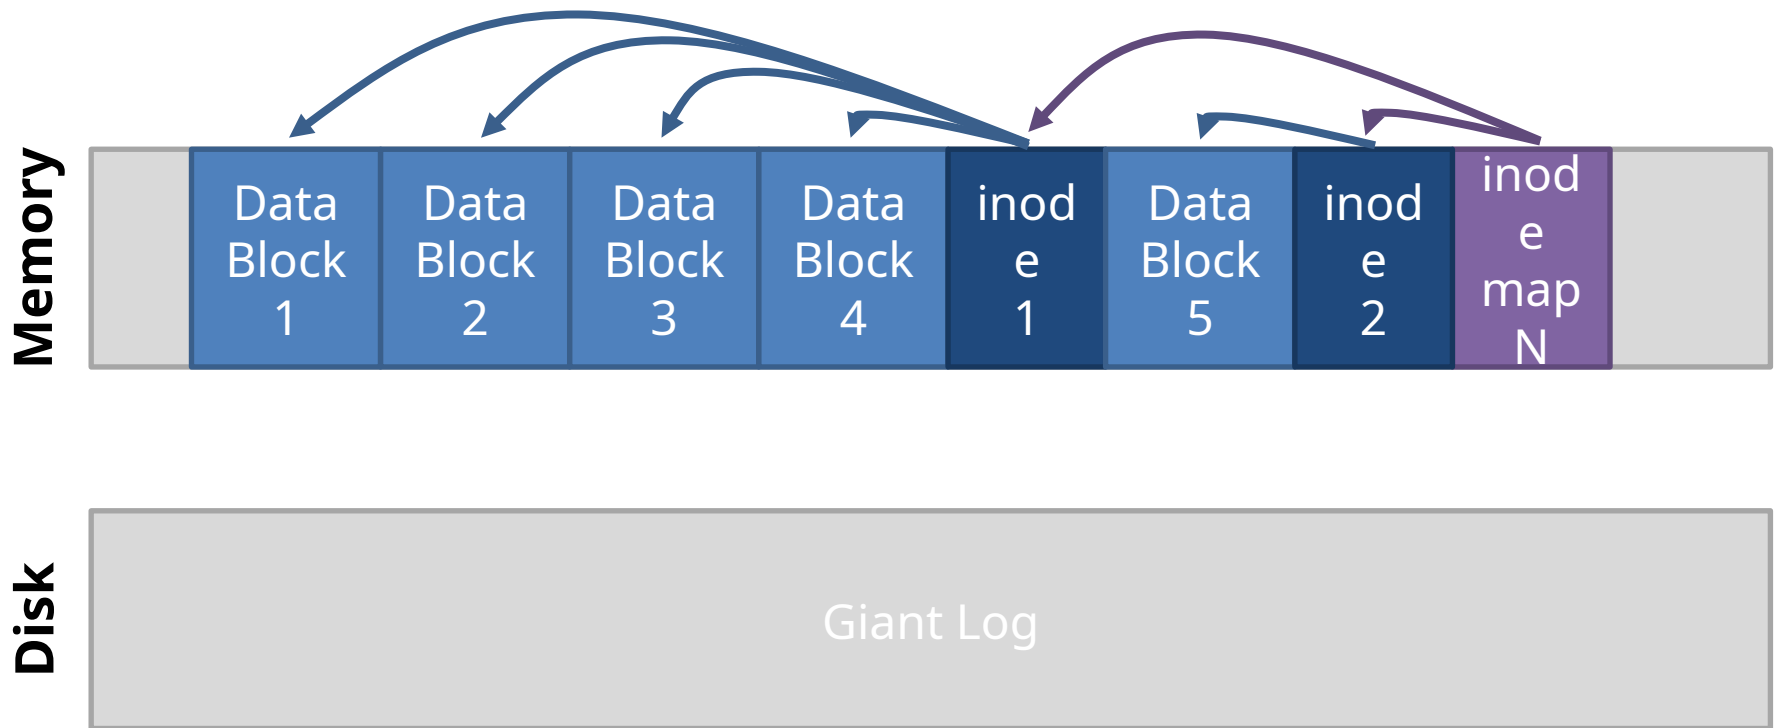
- New problem: the inode map is scattered throughout the log
 - How do we find the most up-to-date pieces?

The Checkpoint Region

- The superblock in LFS contains pointers to all of the up-to-date inode maps
  - The checkpoint region is always cached in memory
  - Written periodically to disk, say ~30 seconds
  - Only part of LFS that isn't maintained in the log

Disk

| CR | | Data Block 1 | Data Block 2 | Data Block 3 | Data Block 4 | inode 1 | Data Block 5 | inode 2 | inode map N | |

How to Read a File in LFS

- Suppose you want to read inode 1
  1. Look up inode 1 in the checkpoint region
     - inode map containing inode 1 is in sector *X*
  2. Read the inode map at sector *X*
     - inode 1 is in sector *Y*
  3. Read inode 1
   - File data is in sectors *A*, *B*, *C*, etc.
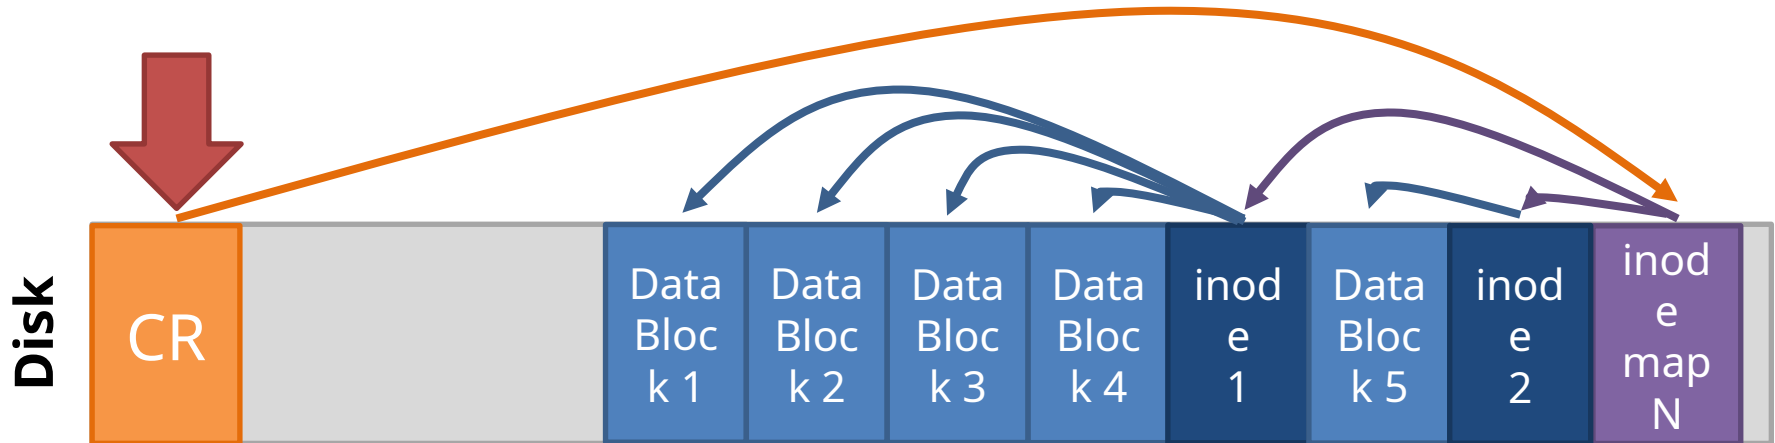
Directories in LFS

- Directories are stored just like in typical file systems
  - Directory data stored in a file
  - inode points to the directory file
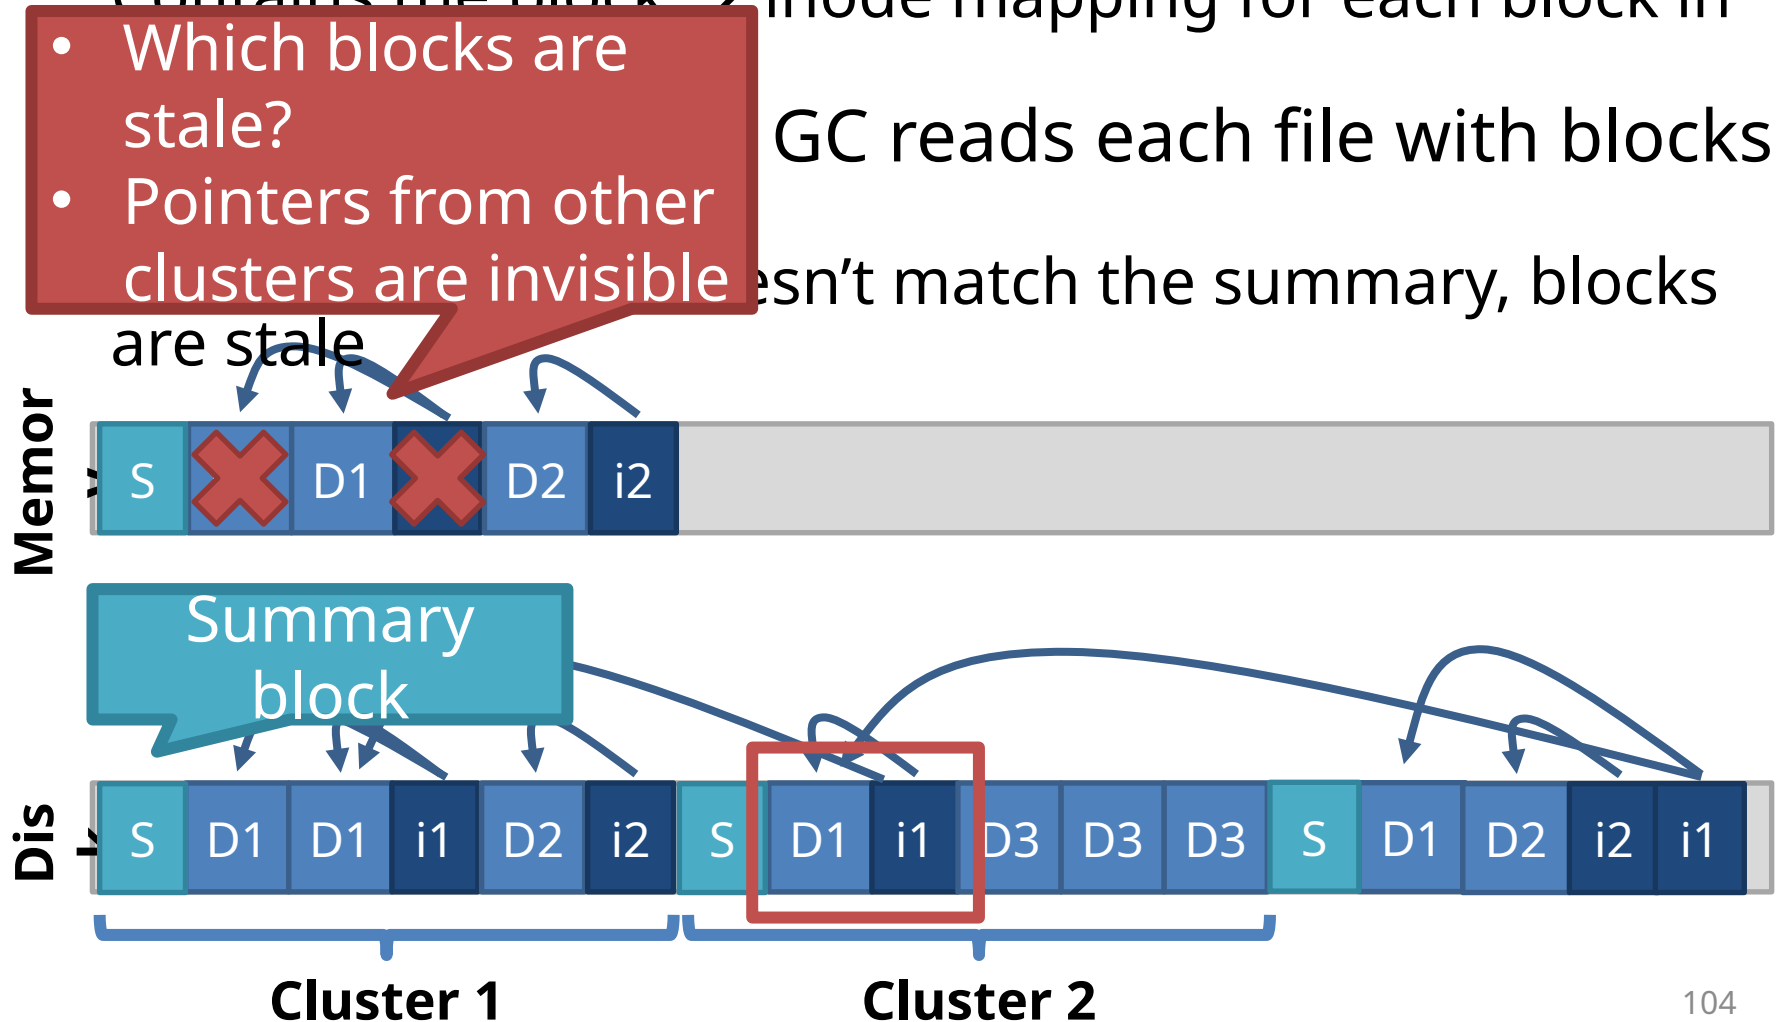 - Directory file contains name → inode mappings
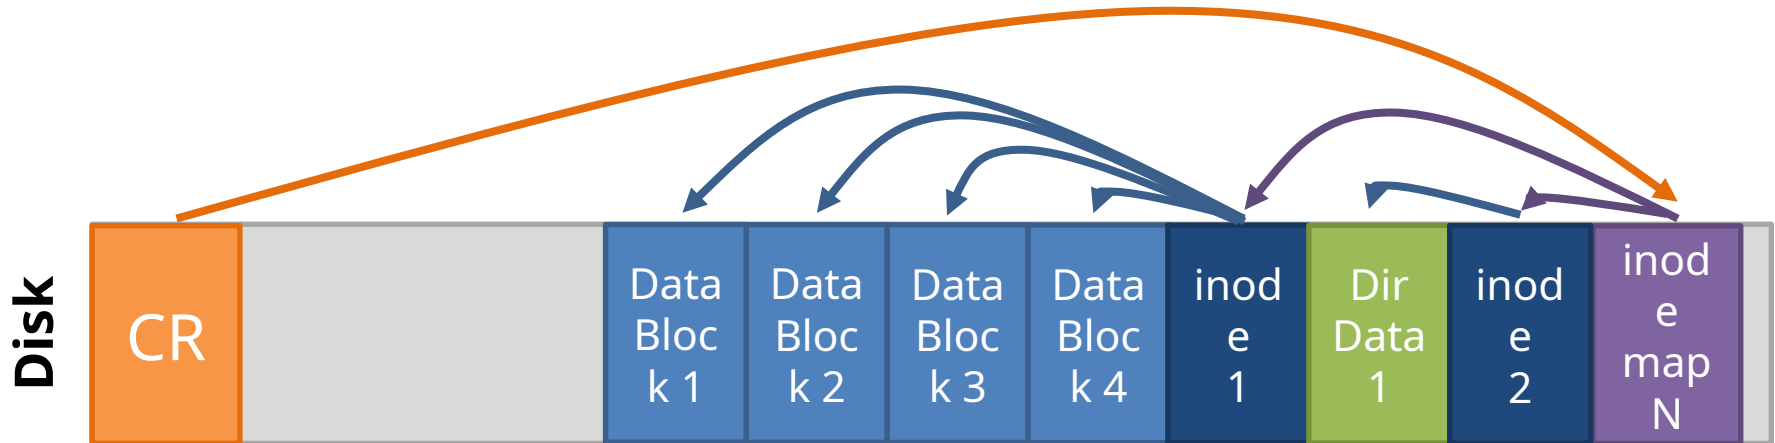
Garbage

- Over time, the log is going to fill up with stale data
  - Highly fragmented: live data mixed with stale data

- Periodically, the log must be garbage collected

Garbage Collection in LFS

- Each cluster has a summary block
  - Contains the block → inode mapping for each block in

- GC reads each file with blocks

  esn't match the summary, blocks
  are stale

Which blocks are stale?

Pointers from other clusters are invisible

Summary block

Memor y

Dis k

Cluster 1   **Cluster 2**

An Idea Whose Time Has Come

- LFS seems like a very strange design
  - Totally unlike traditional file system structures
  - Doesn't map well to our ideas about directory heirarchies
- Initially, people did not like LFS
- However, today it's features are widely used

File Systems for SSDs

- SSD hardware constraints
  - To implement wear leveling, writes must be spread across the blocks of flash
  - Periodically, old blocks need to be garbage collected to prevent write-amplification
- Does this sounds familiar?
- LFS is the ideal file system for SSDs!
- Internally, SSDs manage all files in a LFS
  - This is transparent to the OS and end-users
  - Ideal for wear-leveling and avoiding write-amplification

Copy-on-write

- Modern file systems incorporate ideas from LFS
- Copy-on-write sematics
  - Updated data is written to empty space on disk, rather than overwriting the original data
  - Helps prevent data corruption, improves sequential write performance
- Pioneered by LFS, now used in ZFS and btrfs
  - btrfs will probably be the next default file system in Linux

Versioning File Systems

- LFS keeps old copies of data by default
- Old versions of files may be useful!
  - Example: accidental file deletion
  - Example: accidentally doing *open(file, 'w')* on a file full of data
- Turn LFS flaw into a virtue
- Many modern file systems are versioned
  - Old copies of data are exposed to the user
  - The user may roll-back a file to recover old versions

108